

Thesis Plan

Decentralized Multi-Agent Reinforcement Learning for Power Grid Topology Control

Corentin Plumet

August 18, 2026

Contents

1	Thesis Plan	1
1.1	Working Title	1
1.2	Central Thesis Question	1
1.3	Research Questions	1
1.4	Expected Contributions	2
1.5	Open Decisions	2
2	Introduction	3
2.1	Topology control	3
2.2	Why it has to be automated	3
2.3	Why decentralized agents	4
2.4	Reinforcement learning for topology control	4
2.5	Starting point: MARL2Grid and MAPPO	5
2.6	Outline	6
3	Establishing a Baseline	7
3.1	Evaluation Protocol	7
3.2	IEEE 14-Bus Experiments	7
3.2.1	MLP Baseline	7
3.2.2	Initial Do-Nothing Bias	9
4	Sparse Intervention Control	11
4.1	Evaluation-Time Rho Heuristics	11
4.2	Intervention-Gated Actor	12
4.3	Fixed Sparse-Action Penalty	13
4.4	Adaptive Intervention Budget	13
4.5	Behavioral Cloning from Heuristics	14
4.6	Results	14
4.7	Preliminary WCCI Stress Test	15
5	Methods	17
5.1	Physical Scaling and Voltage-Angle Representation	17
5.1.1	Physical scaling	17
5.1.2	Angle representation	17
5.1.3	Use across the experiments	18
5.2	Graph Representations	18
5.2.1	Busbar Representation	18
5.2.2	Disaggregated Representation	19
5.2.3	Disaggregated Line-Node Representation	21
5.2.4	Local actor information boundary and the global state graph	23
5.2.5	Graph actor and centralized critic	25
5.3	Substation Representation: Direct Edges and Summary Nodes	26
5.3.1	Direct same-substation edges (e)	27
5.3.2	Substation summary nodes (n)	28
5.4	Pooling and Graph Readout	28
5.4.1	Mean and maximum pooling	29
5.4.2	Virtual-node graph readout	30

5.5	Candidate-Action Scoring from the Disaggregated Line-Node Graph	31
5.5.1	From action indices to candidate scoring	31
5.5.2	What each action touches	32
5.5.3	Pooling the touched nodes	33
5.5.4	Scoring and the idle action	34
5.5.5	Cost and current limits	34
5.6	Action-Space Reduction for Large Grids	35
6	Graph-Design Screening	36
6.1	Preliminary GNN Encoder Viability	36
6.2	How the Screens Are Organized	38
6.3	Shared Protocol	40
6.3.1	Endpoint statistics	40
6.3.2	Compute	40
6.4	Screen A: Encoder Depth and Width	41
6.5	Screen B: Graph Readout Scope and Reducer	43
6.6	Screen C: Structural Augmentations of the Busbar Graph	45
6.7	Screen D: Generator and Load Message Direction	46
6.8	Screen E: Candidate-Action Scoring on the Corrected Boundary	48
6.8.1	Replication with Corrected Graph Inputs	48
7	What the Shared Encoder Reads	50
7.1	Why the input distribution matters	50
7.2	Measurement protocol	50
7.3	What the encoder reads	51
7.4	Results	51
7.5	Two failure modes	53
7.6	Correcting the inputs	54
7.6.1	Running standardization makes the mismatch worse	54
7.6.2	Physical scaling for the power channels	55
7.6.3	The angle displacement is a gauge artifact	55
7.6.4	The action descriptor must not be normalized per agent	56
7.6.5	The corrected distribution	56
8	Transfer to the 36-Substation WCCI Grid	58
8.1	The target environment	58
8.1.1	Removing maintenance	59
8.1.2	Action space	59
8.2	Corrected inputs	60
8.3	Transfer experiment design	60
A	Supplementary Results	62
A.1	Sparse Intervention Control	62
A.2	Screen A: Encoder Depth and Width	64
A.3	Screen B: Graph Readout Scope and Reducer	65
A.4	Screen C: Structural Augmentations of the Busbar Graph	66
A.5	Screen D: Generator and Load Message Direction	68
A.6	Screen E: Candidate-Action Scoring	69
A.7	Encoder Input Distributions	70
A.8	Deterministic Feature Scales	71
B	GINE Graph-Actor Architecture	72
B.1	Shared input and encoder pipeline	72
B.2	GINE message passing	72
B.3	Pooling and graph output	73
B.4	Agent-specific action head	73
B.5	Exact node inputs by representation	74
B.6	Exact edge inputs by representation	74

C	Graph Construction Details	75
C.1	Candidate graph sizes	75
C.2	Indexing and masking conventions	75
C.3	Local actor graphs	76
C.4	Boundary verification	76
D	Experiment Configurations	77
E	Full-Test Checkpoint Registry	80

Chapter 1

Thesis Plan

1.1 Working Title

Graph Neural Networks for Decentralized Multi-Agent Reinforcement Learning in Power-Grid Topology Control

1.2 Central Thesis Question

How can graph neural networks help decentralized reinforcement learning agents learn reliable power-grid topology control policies that preserve survival performance while scaling to combinatorial action spaces?

Three terms fix what the thesis has to measure:

- **Reliable** is episodic survival on held-out chronics, always reported from the checkpoint with the best test survival and re-evaluated on the full test split.
- **Decentralized** means each agent owns a fixed group of substations, observes a local subgraph, and acts on its own discrete topology actions without having any information about the other agent actions; only the critic is centralized.
- **Scaling** means larger grids, which a graph encoder addresses because its parameters are independent of grid size.

1.3 Research Questions

1. **Can decentralized MAPPO first learn a strong, stable baseline?** The flat MLP actor must achieve reliable survival on bus14 before either its objective or its representation is changed. (Chapter 3.)
2. **Can the agents intervene sparsely without losing survival?** Evaluation-time rho overrides, an intervention-gated actor, fixed non-idle penalties, and adaptive Lagrangian intervention budgets are evaluated on two axes at once, survival and intervention rate, since improving either alone is not a result. (Chapter 4.)
3. **Can a graph actor replace the flat one without paying for it?** An MLP's input width is tied to one grid, whereas a shared graph encoder is size-independent. That advantage matters only if the GNN first matches the tuned MLP on bus14. (Chapter 6.)
4. **How should the grid be represented as a graph for a local actor?** Node granularity, relations, substation-level augmentations, and graph readout are all design choices. Chapter 5 defines them and Chapter 6 measures which ones move survival.

5. **What is a local agent entitled to observe, and does enforcing that boundary matter?** A local graph carries contextual nodes from neighboring substations. Making context visible only through a live boundary connection, and choosing whether pooling covers the visible graph or only the controlled domain, determine what the policy conditions on. The distinction is between context and leakage. (Sections 5.2.4 and 5.4.)
6. **Does any of this transfer to a larger grid?** Moving from bus14 to the 36-substation WCCI grid tests whether a graph encoder is portable in practice and not only in parameter shape, and what has to be corrected: mainly feature scaling and the size of the action space. (Chapter 8.)

1.4 Expected Contributions

- **A tuned decentralized MAPPO baseline on bus14.** The original framework trained unstably; the training and initialization changes that fix it are isolated one at a time.
- **An ablation of sparse intervention mechanisms** reported jointly on survival and intervention frequency, connecting the learned behavior to what a control room would accept: act rarely, act locally, and act when the grid is unsafe.
- **A shared GNN actor and a screened local graph design.** The GNN first matches the tuned MLP, then three representations are compared using a fixed candidate graph with a dynamic edge mask and an explicit information boundary. Encoder depth and width, message-passing operator, readout, and structural augmentations are screened with full-test evaluation.
- **A size-invariant mechanism for combinatorial action spaces.** A candidate-action head that scores each action from the encoded nodes it modifies, replacing the fixed-width layer over an action index.
- **A measurement of what actually breaks under transfer.** Encoder parameters cross grids unchanged; the input distribution they were calibrated on does not. Two findings generalize beyond power grids: per-feature standardization is harmful when a channel is a mixture dominated by structural zeros, and for gauge-dependent quantities such as voltage angle, changing the representation beats any choice of normalization.

1.5 Open Decisions

- Confirm the final thesis title and whether to emphasize GNNs, sparse-control, or decentralized MARL most strongly.
- Decide which experiment families are final and which are only exploratory.
- Decide whether the main claim is performance improvement, improved sparsity at equal survival, or reproducible ablation insight.
- Choose the final bibliography style required by the university.

Chapter 2

Introduction

2.1 Topology control

A transmission grid is not a fixed graph. Inside a substation, every connected element, each transmission line, generator, and load, is attached to one of several busbars, and an operator can move an element from one busbar to another or split a substation into two electrically separate nodes. Doing so changes which paths the power actually flows through, and therefore redistributes loading across the network. This is topology control, also called busbar switching or network reconfiguration.

Its appeal is economic. Congestion is classically relieved by redispatching generation or by curtailing renewable production, both of which cost money at every intervention, or by building new lines, which costs a great deal once and takes years. Topology control uses assets that already exist and switching equipment that is already installed, so its marginal cost is close to zero [1]. The energy transition makes this attractive rather than merely elegant: variable renewable generation makes flows less predictable and pushes existing corridors closer to their thermal limits, while grid expansion is slow [2], [3].

The catch is that topology is hard to exploit. Operators use it, but mostly through a small repertoire of remedial actions established by experience and offline study, so a large part of the available flexibility is never searched [3].

2.2 Why it has to be automated

Three properties make manual search impractical.

The action space is combinatorial. A substation with d connected elements and two busbars admits 2^d assignments, and the number of joint configurations of a grid is the product over its substations. Even the small IEEE 14-bus benchmark used throughout this thesis exposes 178 distinct unitary topology actions; the 36-substation grid of Chapter 8 exposes more than 66 000, which is why it has to be reduced offline before an agent can be trained on it at all (Section 5.6).

The effect of an action cannot be read off the grid. Power flows are governed by non-linear AC equations, so the consequence of moving one line onto another busbar is a global redistribution, not a local change. There is no reliable shortcut for ranking actions other than simulating them.

Finally, the decision is sequential and time-constrained. An action taken now changes which contingencies are survivable later, cooldown periods forbid acting on the same substation twice in quick succession, and an operator has minutes, not hours. Optimizing a single step greedily is not the same problem as keeping a grid alive for a day.

These are the conditions under which reinforcement learning is a reasonable candidate: a sequential decision problem, a simulator that is expensive but available, no analytic optimum, and a clear success

signal — the grid does not black out. The Learning to Run a Power Network (L2RPN) competition series was created by RTE to make exactly this problem tractable for machine-learning research, and the Grid2Op simulator provides the environment, the chronics of load and generation, and the operational rules that go with it [2], [4].

2.3 Why decentralized agents

Most published work treats the grid as a single agent controlling everything. Real grids are not operated that way. Transmission system operators divide the network into regions, each with its own control centre and its own operators, who act on their own substations with a detailed view of their own area and a much coarser view of the rest [3]. A decentralized formulation is therefore closer to how the job is actually done, and to how any automated assistant would eventually be deployed, as support inside one control room, not as a single controller over an entire interconnection.

It is also the formulation that scales. A monolithic controller has to represent the joint action space, which grows multiplicatively with the number of substations. Partitioning the grid into regions and giving each region its own policy replaces that product by a sum: each agent chooses among its own local actions only. This is the classical argument for decentralized multi-agent reinforcement learning, and it comes with the classical cost [5]. Because every agent learns while the others are also changing, each one faces an environment that is non-stationary from its own point of view, and because each observes only part of the grid, agents can take individually sensible actions that are jointly harmful.

The usual answer is *centralized training with decentralized execution*: a critic is allowed to see the global state during training, while each actor conditions only on what it will have access to at deployment time [6], [7]. This thesis uses that setting throughout.

2.4 Reinforcement learning for topology control

Work on RL for topology control falls into four groups, in roughly historical order. A broader survey of graph-based RL for power systems is given by Hassouna et al. [8].

Single-agent policies. The first approaches train one agent to control the whole grid. Subramanian et al. [9] showed that a plain deep RL agent acting on a curated set of topology actions already outperforms a do-nothing policy on L2RPN benchmarks. The strongest example is SMAAC [10], which won the L2RPN WCCI 2020 challenge with an actor-critic method over *afterstates* — the grid configuration that results from an action, evaluated before the physics resolves it — combined with a hierarchical decomposition of where to act and how. The common obstacle in this group is the action space: it has to be reduced, curated, or factorized before learning becomes feasible at all.

Learned policies wrapped in heuristics. A second group accepts that a learned policy should not be in charge at every step, and surrounds it with rules. The dominant pattern is an activation threshold: the agent is queried only when the grid is stressed, measured by the maximum line loading $\rho_{\max}$, and does nothing otherwise. PowRL added a heuristic overload manager and a reduced action set on top of an RL policy and topped the L2RPN NeurIPS 2020 robustness track [11]. Lehna et al. [12] compared advanced rule-based agents against RL agents directly and found comparable survival, with the RL agent far cheaper at inference — a useful reminder that beating a well-engineered heuristic is not automatic. Heuristics of this kind are now standard practice rather than a baseline, and this thesis uses one.

Hierarchical and centrally coordinated multi-agent methods. A third group decomposes the decision but keeps a coordinator. Manczak et al. [13] split control into choosing whether to act, choosing where, and choosing which local configuration to apply. van der Sar et al. [14] assign a low-level agent to each substation under a higher-level agent that selects which substation acts. de Mol et al. [15] make the coordination explicit: regional agents propose actions and a central agent selects among the proposals.

These methods recover much of the scalability of factorization, but they retain a component that sees the whole grid and decides for everyone, so they do not answer whether purely local policies suffice.

Decentralized multi-agent control. The remaining question — whether agents that act on local observations, without a coordinator and without communication, can operate a grid — is much less explored, and it is where this thesis sits. The gap is recognized: the MARL2Grid-TR benchmark was built specifically because prior work “focused on single-agent settings, neglecting the often decentralized, multi-agent nature of grid control” [16].

Graph representations. Cutting across all four groups is the question of how the grid is given to the network. A flat vector ties the policy to one grid size and discards the topology entirely. Graph neural networks are the natural alternative [17], [18], [19], and their parameters are independent of the number of substations, which is what makes transfer between grids conceivable at all. The choice of representation is not neutral, however: de Jong et al. [20] show that the standard homogeneous graph suffers from *busbar information asymmetry* — it encodes the connections that currently exist but not the ones an action could create — and that a heterogeneous representation resolves it. That observation is close to the subject of Chapters 5 and 6.

2.5 Starting point: MARL2Grid and MAPPO

This thesis builds on MARL2Grid, the multi-agent benchmark and codebase of Marchesini et al. [16], itself the multi-agent successor of RL2Grid [21]. It wraps Grid2Op in a multi-agent interface and supplies the decomposition, the observation and action spaces, the reward, and the training loop that the rest of this work modifies.

Decomposition. The grid is partitioned into disjoint regions, one per agent, fixed in advance. On the IEEE 14-bus environment (`bus14`, 14 substations and 20 lines) three agents own substations $\{0, 1, 2, 4\}$, $\{3, 6, 7, 8\}$ and $\{5, 9, \dots, 13\}$, giving them 52, 50 and 76 non-idle actions respectively — the 178 actions of the full grid, partitioned rather than multiplied. Each agent observes only quantities belonging to its own region and chooses among the topology actions available at its own substations; action 0 is always do-nothing. Agents act simultaneously and are given no information about what the others are doing.

Reward and heuristic. The training reward is dominated by a flat +1 per surviving step, plus an overload term, so the objective is essentially to keep the episode alive; the reported metric is episodic survival, the fraction of evaluation episodes completed without a blackout. Following the practice described above, an idle heuristic fast-forwards the simulation while the grid is safe: agents are queried only once $\rho_{\max} \geq 0.90$.

MAPPO. The learning algorithm is multi-agent PPO [7], the natural multi-agent extension of PPO [22] under centralized training with decentralized execution. Each agent i has its own stochastic policy $\pi_{\theta_i}(a_i | o_i)$ over its discrete local action space, and is updated with the clipped surrogate objective

$$L^{\text{CLIP}}(\theta_i) = \hat{\mathbb{E}}_t \left[\min \left(r_t(\theta_i) \hat{A}_t, \text{clip} \left(r_t(\theta_i), 1 - \epsilon, 1 + \epsilon \right) \hat{A}_t \right) \right], \quad r_t(\theta_i) = \frac{\pi_{\theta_i}(a_{i,t} | o_{i,t})}{\pi_{\theta_i^{\text{old}}}(a_{i,t} | o_{i,t})},$$

where $\hat{A}_t$ is a generalized advantage estimate. The advantages come from a single *centralized* value function trained on the global state, which is available in simulation but never given to the actors. Only the actors are kept at execution time, so the deployed system is decentralized even though training was not.

What this thesis starts from. The baseline is that system with flat multilayer-perceptron actors and critic, trained on `bus14`. It is unstable at low survival out of the box, and it is tied to one grid: an MLP’s input width is a function of the number of substations, so nothing it learns can be reused elsewhere. The chapters that follow address these limitations in the order they arose: first stabilize training, then make interventions more operator-like, and finally replace the fixed-width representation.

2.6 Outline

Chapter 3 tunes the flat baseline until it trains stably. Chapter 4 then changes the objective and asks whether the same agents can keep their survival while intervening rarely; its preliminary WCCI pilot exposes the limitations of retaining a fixed-width MLP actor. Chapter 5 consequently defines size-independent graph inputs: feature scaling, three grid representations, substation-level augmentations, the readout, and the candidate-action head. Chapter 6 first checks that a shared graph encoder remains competitive, then screens those design choices. Chapter 8 moves to the 36-substation WCCI grid and tests what actually transfers.

Chapter 3

Establishing a Baseline

This chapter establishes a robust MLP actor baseline for the IEEE 14-bus task.

3.1 Evaluation Protocol

All experiments use an 80/20 train/test chronic split. The 80% training split is used for policy updates, while the 20% test split estimates generalization. Evaluation rollouts run periodically during training using 10 episodes per split.

The primary performance metric is test episodic survival (expressed as a percentage), representing the fraction of evaluation timestep where the agent survives and successfully prevents a blackout. Results are reported using seed-aggregated mean curves, with shaded regions indicating variability. Final evaluations report the performance of the checkpoint that achieved the highest training survival.

3.2 IEEE 14-Bus Experiments

The initial MAPPO baseline from the MARL2Grid framework lead to unstable policies and low survival rates on the IEEE 14-bus topology-control task. The following experiments test training and architectural modifications to establish a high-performing baseline.

3.2.1 MLP Baseline

Question. Which changes are needed to obtain a stable MLP MAPPO baseline on the 14-bus task?

Baseline and changes. We start from the MLP MAPPO baseline of the MARL2Grid framework and test several main stabilization choices:

- using reward normalization;
- changing the critic update so that the critic is updated once for the whole batch instead of once per agent;
- tuning of optimization hyperparameters to stabilize the training, as summarized in Table [3.1](#).

Table 3.1: New hyperparameters for MAPPO training.

Hyperparameter	New Value	Paper Baseline
ENVIRONMENT		
n_envs (number parallel environment)	20	10
PPO OPTIMIZATION		
Total time step	15,000,000	25,000,000
Actor learning rate	1e-4	3e-5
Critic learning rate	1e-4	3e-5
Gamma	0.9	0.99
Update epochs	10	80
Minibatches	16	4
Maximum gradient norm	1.0	10.0

Evidence. Table 3.3 reports final and peak survival for the core ablations. Figure 3.1 overlays all core MLP baseline curves.

Interpretation. The **Optimized MLP Baseline**, which combines reward normalization and an optimized critic update, achieves the best performance and convergence speed, yielding the highest and most stable final survival ($\sim 93\%$). Removing or reducing any of these three elements, using a smaller architecture, disabling reward normalization, or reverting to a non-optimized critic update, results in a drop in either learning speed or final survival. The original unoptimized baseline fails entirely to learn a competitive policy. Because the **Optimized MLP Baseline** consistently outperforms the ablations, it is established as the standard reference for all subsequent comparisons.

Table 3.2: MLP baseline run configurations. The reference run is **Optimized MLP Baseline**.

Run family	Actor/critic hidden layers	Reward norm.	Initial action-0 prob.	Optimized critic update	Difference from Optimized MLP Baseline
Optimized MLP Baseline	$3 \times 256 / 3 \times 256$	true	0.0	true	Reference baseline
Disable optimized critic update	$3 \times 256 / 3 \times 256$	true	0.0	false	Disables optimized critic updates
Smaller MLP actor/critic	$2 \times 256 / 2 \times 256$	true	0.0	true	Uses a smaller MLP actor and critic
Disable reward normalization	$3 \times 256 / 3 \times 256$	false	0.0	true	Disables reward normalization
Old baseline (no val)	$3 \times 256 / 3 \times 256$	true	0.3	false	Uses action-0 bias 0.3 and non-optimized critic updates
Initial Bias 50%	$3 \times 256 / 3 \times 256$	true	0.5	true	Uses action-0 bias 0.5
Initial Bias 70%	$3 \times 256 / 3 \times 256$	true	0.7	true	Uses action-0 bias 0.7




Figure 3.1: All MLP baseline core rerun survival curves shown on one axis.

Table 3.3: MLP baseline core rerun survival summary.

Run	Intended change	Final survival (%)	Peak survival (%)
Optimized MLP Baseline	Optimized MLP baseline	93.4	98.6
Disable optimized critic update	Non-optimized critic update	93.1	95.8
Smaller MLP actor/critic	Smaller MLP architecture	93.7	100.0
Disable reward normalization	No reward normalization	81.8	97.1

3.2.2 Initial Do-Nothing Bias

Question. Does biasing the initial policy toward action 0, the do-nothing action, help or hurt learning? In theory, a do-nothing bias makes sense because in reality, human operators most often do not operate the grid (i.e., they take no action) under normal conditions.

Baseline and changes. Our baseline uses no initial do-nothing bias in this rerun set. The bias controls keep the same general protocol but change the initial probability mass assigned to action 0.

Implementation detail. The bias is applied only at initialization: it changes the starting action distribution before learning, without changing the reward, architecture, or PPO objective. This is implemented by artificially increasing the initial pre-softmax logits for the do-nothing action, such that the policy initially selects this safe action with a higher probability before any learning updates occur. This makes the ablation a test of exploration bias rather than a test of representational capacity.

Evidence. Table 3.4 and Figure 3.2 compare the zero-bias baseline with Initial Bias 50% and Initial Bias 70%.

Interpretation. The effect of the action-0 bias is large and non-monotonic. The 0.5-bias curve remains unstable and substantially below the zero-bias baseline, whereas the 0.7-bias curve eventually reaches the highest final survival in this comparison. The current conclusion is that adding an initialization bias is generally hurtful or highly unstable for the models. Instead of aiding exploration, a strong initialization prior appears to restrict the agent’s ability to reliably discover diverse beneficial actions and should be treated cautiously as a first-order ablation variable.

Table 3.4: Initial do-nothing bias survival summary.

Run	Final survival (%)	Peak survival (%)
Optimized MLP Baseline	93.4	98.6
Initial Bias 50%	50.8	70.2
Initial Bias 70%	99.3	100.0




Figure 3.2: Optimized MLP baseline compared with initial do-nothing bias controls Initial Bias 50% and Initial Bias 70%.

Chapter Summary

Objective: Establish a stable, high-performing flat MAPPO baseline.

Findings:

- **MLP optimization:** Hyperparameter tuning, reward normalization, and optimized critic updates raise the flat MLP baseline to roughly 93% survival.
- **Do-nothing initialization:** The effect is large but non-monotonic; it is not a reliable substitute for learning when to intervene.

Conclusion: The tuned flat actor is a strong survival baseline. The next step is to ask whether it can behave more like an operator by intervening less often without losing that survival, which is the subject of Chapter 4.

Chapter 4

Sparse Intervention Control

The preceding chapter optimized survival on held-out chronics. This chapter introduces a secondary objective: sparsity. A policy that keeps the grid alive by reconfiguring substations at every opportunity is impractical for a control room: interventions carry operational costs, and an agent acting constantly is neither auditable nor trustworthy. The question is whether decentralized agents can maintain high survival while intervening rarely, locally, and only when necessary.

In the literature, sparse activity is most often enforced by hard-coding a physical threshold to dictate when the agent is permitted to act. Instead of relying on an external override, this chapter attempts to make sparse intervention behavior natively learnable.

Mechanisms are therefore evaluated on two competing axes: survival rate and intervention sparsity. An effective mechanism must improve sparsity without degrading survival.

Question. Can the agents keep high survival while behaving more like operators: acting rarely and mainly when the grid state makes an intervention useful?

Baseline and change. This block starts from the flat decentralized MAPPO topology-control actor. For agent i , the baseline policy is one categorical distribution over the full local action space:

$$\pi_i(a_i \mid o_i), \quad a_i \in \{0, 1, \dots, |\mathcal{A}_i| - 1\}.$$

Action 0 is the no-op action. During training, actions are sampled from the policy, while deterministic evaluation uses the greedy action by default:

$$a_{i,t}^{eval} = \arg \max_a \pi_i(a \mid o_{i,t}).$$

The sparse-control roadmap tests five ways to reduce unnecessary topology changes: evaluation-time rho overrides, an intervention-gated actor, fixed non-idle action penalties, adaptive intervention budgets, and behavioral cloning from heuristics.

The intervention gate separates whether to act from which action to execute, following hierarchical topology-control decompositions [13], [14]. The fixed and adaptive penalties instead treat non-idle interventions as a limited resource, following constrained and budgeted reinforcement learning [23], [24], [25].

4.1 Evaluation-Time Rho Heuristics

Implementation detail. The rho heuristic is an evaluation-only diagnostic. It does not change the training objective and it does not prove that the policy learned sparse behavior. The policy first proposes the deterministic action $\bar{a}_{i,t}$, then the evaluator may replace it with action 0.

The global version uses the pre-action maximum line loading

$$\rho_t^{global} = \max_{\ell \in \mathcal{L}} \rho_{\ell,t},$$

and executes

$$a_{i,t}^{exec} = \begin{cases} 0, & \rho_t^{global} < \rho_{safe}, \\ \bar{a}_{i,t}, & \text{otherwise,} \end{cases} \quad \rho_{safe} = 0.90.$$

The local version is decentralized. Agent i computes the maximum rho on the lines touching its regional substations,

$$\rho_{i,t}^{local} = \max_{\ell \in \mathcal{L}_i} \rho_{\ell,t},$$

and only that agent is forced to no-op when $\rho_{i,t}^{local} < \rho_{safe}$. A boundary line can appear in the local neighborhood of both adjacent agents, but no communication is required: both agents can observe the line through their own physical neighborhood.

Interpretation. The heuristic is useful to ask what happens if safe-state interventions are blocked, but it is not a learned method. Its action-0 rates must therefore be reported as final executed evaluation actions, not as evidence that the policy itself prefers action 0.

4.2 Intervention-Gated Actor

Implementation detail. The gated actor changes the actor factorization. Instead of one categorical head over all actions, each agent has a gate head and a non-idle action head:

$$\pi_i^{gate}(g_i | o_i), \quad g_i \in \{0, 1\},$$

where 0 means do nothing and 1 means intervene, and

$$\pi_i^{nonidle}(b_i | o_i), \quad b_i \in \{1, \dots, |\mathcal{A}_i| - 1\}.$$

During training, the gate is sampled first. If $g_{i,t} = 0$, then $a_{i,t} = 0$. If $g_{i,t} = 1$, the final action is sampled from the non-idle head. The PPO log-probability of the executed action is therefore

$$\log P(a_i = 0) = \log P(g_i = 0),$$

and, for $a_i > 0$,

$$\log P(a_i) = \log P(g_i = 1) + \log P(b_i = a_i | g_i = 1).$$

Two deterministic decoders are used. The `final_action_map` decoder selects the most likely final executed action:

$$P(a_i = 0) = P(g_i = 0), \quad P(a_i = a > 0) = P(g_i = 1)P(b_i = a | g_i = 1).$$

The `hierarchical_greedy` decoder first chooses the most likely gate decision, then chooses the most likely non-idle action only if the gate selects intervention.

The original coupled entropy term is

$$H_i = H(g_i) + P(g_i = 1)H(b_i).$$

This can accidentally reward intervention because larger $P(g_i = 1)$ unlocks more non-idle entropy. The separated entropy variant is

$$H_i = \alpha_{gate}H(g_i) + \alpha_{nonidle}H(b_i),$$

and is used in the gate-plus-budget diagnostic runs.

Interpretation. The gate is learned and decentralized, but it can restrict exploration more strongly than a reward penalty. If it collapses early to no-op, the non-idle head is rarely executed and receives little learning signal. For this reason, the gate is treated as an important diagnostic rather than the main sparse control contribution.

4.3 Fixed Sparse-Action Penalty

Implementation detail. The fixed-penalty sweep adds a reward penalty to the final executed action of each agent:

$$r'_{i,t} = r_{i,t} - \lambda_{fixed} \mathbf{1}[a_{i,t} \neq 0],$$

with

$$\lambda_{fixed} \in \{0.0, 0.001, 0.003, 0.01\}.$$

The grid is

$$\text{actor} \in \{\text{flat}, \text{gated}\}, \quad \lambda_{fixed} \in \{0.0, 0.001, 0.003, 0.01\}, \quad \text{seed} \in \{0, 1, 2\},$$

which gives $2 \times 4 \times 3 = 24$ runs. The state-dependent safety penalty is disabled in every condition (`safe_intervention_penalty = 0.0`), so the penalty is not rho-weighted.

Interpretation. This mechanism changes only the training reward. It does not override actions during evaluation and it does not hard-mask exploration during training. The policy can still sample non-idle actions, but each such action must be useful enough to pay its fixed cost. Comparing the flat and gated actors at $\lambda_{fixed} = 0.010$ tests whether the gate adds value once a simple sparse objective already exists.

4.4 Adaptive Intervention Budget

Implementation detail. The adaptive intervention budget turns sparse topology control into a local constraint. For each agent, the cost is $c_i(s_t, a_{i,t}) = \mathbf{1}[a_{i,t} \neq 0]w_i(s_t)$, and the desired constraint is $E_{\pi_i}[c_i(s_t, a_{i,t})] \leq d$.

The implemented rollout reward removes the constant term that does not affect the PPO action gradient:

$$r'_{i,t} = r_{i,t} - \lambda_i c_i(s_t, a_{i,t}).$$

After each rollout, the local multiplier is updated by

$$\lambda_i \leftarrow \text{clip}\left(\lambda_i + \eta(\hat{C}_i - d), 0, \lambda_{\max}\right), \quad \hat{C}_i = \frac{1}{T} \sum_{t=1}^T c_i(s_t, a_{i,t}).$$

If the observed cost $\hat{C}_i$ is above the budget d , future interventions become more expensive; if it is below the budget, the penalty relaxes. The default settings are `intervention_budget_lr = 0.02`, `intervention_budget_init_lambda = 0.0`, and `intervention_budget_max_lambda = 10.0`.

Two cost modes are used:

- **Non-idle cost:** $w_i(s_t) = 1$, so $c_i(s_t, a_{i,t}) = \mathbf{1}[a_{i,t} \neq 0]$. This is an adaptive version of a plain sparse-action penalty.
- **Local-safe cost:** $w_i^{local}(s_t) = \sigma(k(\rho_{safe} - \rho_{i,t}^{local}))$. This is the decentralized version: safe local states make interventions expensive, while stressed local states make them cheap.

In the AIB setting, $\rho_{safe} = 0.90$ is not a hard action override. It is the center of a smooth sigmoid safety weight.

AIB experiment grid. The main local-budget ablation contains five conditions:

- flat actor, local-safe cost, $d = 0.20$;
- flat actor, stricter local-safe budget, $d = 0.10$;
- flat actor, looser local-safe budget, $d = 0.35$;
- gated actor, hierarchical-greedy evaluation, separated entropy, local-safe cost, $d = 0.20$;
- flat actor with adaptive plain non-idle cost, $d = 0.20$.

Comparing the local-safe and plain non-idle costs at $d = 0.20$ isolates the value of local rho weighting.

4.5 Behavioral Cloning from Heuristics

Implementation detail. The evaluation-time rho heuristics (Section 4.1) produce highly sparse, safe behavior, but they require a manual override. To internalize this behavior into the policy weights, a teacher-student distillation approach is tested. The teacher is a MAPPO model running the local rho heuristic, and the student is trained via Behavioral Cloning (BC) to match the teacher’s executed actions. The student is then evaluated without the heuristic override to see if it learned to voluntarily imitate the sparsity. Two balancing configurations are evaluated to manage the severe class imbalance caused by the teacher’s high action-0 rate.

Interpretation. As reported in Appendix A.1, the student policies reach roughly 91–92% survival and 94–96% sparsity. While this shows that sparse behavior can be partially distilled into the actor weights, the resulting student is weaker than the AIB methods and sacrifices a portion of the teacher’s survival (97.1%). Consequently, the distillation route is less reliable than natively learning the sparse objective through the Adaptive Intervention Budget.

4.6 Results

Figure 4.1 summarizes survival and the mean action-0 fraction across agents. Exact values and per-agent action-0 fractions are reported in Appendix A.1.

Table 4.1: Sparse-control mechanisms and where they affect the policy.

Method	Training effect	Evaluation effect	Main risk
Baseline	Samples from the flat policy	Greedy argmax by default	Standard entropy-controlled exploration
Rho heuristic	No training change	Hard no-op override in safe states	Evaluation is manually changed
Intervention gate	Changes action factorization and sampling	Final-action MAP or hierarchical greedy	Gate collapse can starve non-idle exploration
Fixed penalty	Fixed reward penalty on non-idle actions	No action override	Penalty coefficient must be tuned
AIB non-idle	Adaptive reward penalty on non-idle actions	No action override	Multiplier can grow if target is too strict
AIB local-safe	Adaptive state-dependent reward penalty	No action override	Depends on the local rho
Behavioral cloning	Distills heuristic into actor weights	No action override	safety-weight design Struggles with extreme action-0 class imbalance




Figure 4.1: Joint action-0/survival comparison of the selected sparse-control mechanisms on **bus14**. Large markers denote condition means and faint markers the individual seeds. The plot uses checkpoint-aligned action-0 rates and full-test survival, so it visualizes the chapter’s two objectives.

Interpretation. The current roadmap interpretation is that the gate is not the strongest contribution by itself. Sparse reward shaping already improves action-0 usage, the adaptive budget makes the sparse penalty self-tuning, and the local-safe cost makes the penalty more physically meaningful. In particular, the flat actor with $\lambda_{fixed} = 0.010$ is a strong fixed-penalty baseline, while the flat local-safe AIB policy with $d = 0.20$ is the most principled learned sparse-control candidate: it keeps the original decentralized actor, does not rely on an evaluation-time override, learns the intervention penalty adaptively, and uses a local state-dependent safety weight. Figure 4.1 also makes the main failure mode visible: high action-0 usage is not sufficient when the gate collapses and survival falls.

4.7 Preliminary WCCI Stress Test

As a preliminary scale test, representative mechanisms were retrained from scratch with an MLP actor on the maintenance-enabled WCCI grid; no **bus14** weights were transferred. All runs use the same reduced action space, three seeds, 60M training steps, and learning-rate schedule. The reported values are the final deterministic test evaluation, averaged over seeds and ten held-out episodes per seed; they are not full-test estimates.

Table 4.2: Preliminary WCCI stress test with scheduled maintenance at the final training-time test evaluation.

Method	Survival (%)	Action-0 fraction
Matched MLP baseline	23.75 ± 7.70	0.168
Global rho heuristic, $\rho_{safe} = 0.90$	28.35 ± 11.06	0.974
Local rho heuristic, $\rho_{safe} = 0.90$	9.46 ± 3.54	0.995
Intervention gate, final-action MAP	25.92 ± 10.69	0.400
Flat fixed penalty, $\lambda_{fixed} = 0.010$	27.00 ± 12.13	0.489
Flat local-safe AIB, $d = 0.20$	21.65 ± 2.21	0.927




Figure 4.2: Action-0/survival trade-off for the preliminary maintenance-enabled WCCI runs. Each point is a seed-aggregated run variant. The upper-right region remains preferred, but no method combines high sparsity with strong survival on this larger grid.

The mechanisms change intervention frequency, but none reliably improves survival over the weak MLP baseline: the differences are small relative to the seed variation. The local rho heuristic and AIB are highly sparse but remain weak controllers, suggesting that sparsity transfers more readily than control quality. This failed pilot changes grid size, action space, and maintenance at once, so it cannot identify the cause. It nevertheless exposes the limitation of continuing with a fixed-width MLP actor on the larger grid. The next two chapters therefore introduce a size-independent graph representation and test its design choices. Chapter 8 later returns to WCCI and isolates the confounded factors before studying cross-grid transfer.

Chapter Summary

Objective: Preserve survival while reducing unnecessary topology interventions.

Findings:

- **Evaluation heuristics** reach 97–98% survival with action 0 in 94–97% of decisions, but sparsity is imposed after the policy acts.
- **A fixed penalty** already gives a strong learned controller: $\lambda_{fixed} = 0.010$ reaches 97.44% survival and a 0.905 action-0 fraction without an evaluation override.
- **Local-safe AIB** gives the best learned trade-off: budgets $d = 0.10$ and $d = 0.20$ reach about 98.4% survival and 0.98 action-0 usage. The gated variants are less reliable and can collapse to almost permanent no-op behavior.

Conclusion: Sparse behavior can be learned on `bus14`, but the same MLP mechanisms remain weak on WCCI. This motivates the graph representation introduced next.

Chapter 5

Methods

Chapter 4 made the flat policy more operator-like on bus14, but its preliminary WCCI stress test remained weak. Sparse intervention control changes when the policy acts; it does not remove the fixed-width MLP representation that ties the actor to one grid. This chapter therefore introduces the size-independent graph inputs and graph-actor variants used in the remaining experiments. The presentation separates four design choices: feature scaling, the base graph representation, substation-level augmentations, and graph pooling.

5.1 Physical Scaling and Voltage-Angle Representation

Two input transformations make observations more comparable across grids: a fixed physical scaling of the continuous features, and a reference-independent representation of voltage angles. Both act before the graph encoder and leave it unchanged.

5.1.1 Physical scaling

Each scaled continuous feature is divided by a fixed reference of the same unit, $x'_f = x_f/s_f$. This changes the unit without changing the information and — importantly for the sparse graph features — keeps zero at zero, so a busbar carrying no asset still reads 0 on every grid. Two references carry the physical content: power is divided by the rating of the largest installed generator in the grid, $s_P = \max_g P_g^{\max}$, and angles by a fixed 180 degrees. The remaining discrete channels, the two cooldown counters, the overflow counter, and the maintenance timers, are divided by their Grid2Op parameter bounds, listed in Table A.13. Every reference is fixed at construction rather than estimated during training. The dimensionless loading ratio `rho` and all binary, mask, type, and relation indicators are left unchanged.

5.1.2 Angle representation

An absolute voltage angle needs an arbitrary zero, normally fixed by the slack bus, so adding a constant to every bus angle describes the same physical state. The angle difference across a line does not depend on that choice:

$$\Delta\theta_\ell = \theta_\ell^{\text{or}} - \theta_\ell^{\text{ex}}.$$

This quantity is *gauge-invariant*, and being defined per line it is also dense, whereas an absolute angle exists only where a generator or load attaches. Enabling the drop removes the absolute-angle channels. Chapter 7 measures why this matters: the two grids differ by 5 to 7° in absolute angle, largely because they place their slack bus differently, and no fixed scaling can remove a displacement.

The drop is used only with the disaggregated line-node representation of Section 5.2.3, where each physical line is itself a node and its attachment edges carry no measurements, so the drop is stored on the line node. This gives every local actor the drop for each line it observes, boundary lines included. A disconnected line is assigned a zero drop, because the simulator may retain stale endpoint angles after

an outage. The two schemas whose lines are edges would carry the drop on the edge instead, but no reported run uses that combination.

5.1.3 Use across the experiments

The screening experiments of Chapter 6 and the sparse-control experiments of Chapter 4 use raw features and absolute angles. The cross-grid transfer study of Chapter 8 pairs runs with both transformations enabled against runs with neither, on both grids, so the comparison isolates their effect. Running standardization is also implemented but no reported policy run uses it; Section 7.6.1 analyses it offline and finds that it increases the mismatch between grids.

5.2 Graph Representations

The graph observation can use one of three base representations, named here after how much equipment receives its own node. The *busbar* representation aggregates equipment values onto busbar nodes. The *disaggregated* representation gives each generator and load its own node while keeping physical lines as busbar-to-busbar edges. The *disaggregated line-node* representation goes one step further and turns each transmission line into a node as well.

Table 5.1 states what separates them. The node schema also determines how a boundary transmission line is represented in a local actor graph. Consequently, the three representations are not information-equivalent: they expose different amounts of context about the substation at the far end of a shared line. This distinction is made explicit in Section 5.2.4.

Table 5.1: The three graph schemas. Each row is a consequence of how much equipment receives its own node. In the node counts, S is the number of substations, B the busbars per substation, N_g and N_d the included generators and loads, and L the included transmission lines.

Aspect	busbar	disaggregated	disaggregated line-node
gnn_graph_type	bus	heterogeneous	heterogeneous_line
Node types	Busbar	Busbar, generator, load	Busbar, generator, load, transmission line
Generator and load state	Summed and averaged onto the busbar	One node per asset	One node per asset
Physical line	Direct busbar-to-busbar edge	Typed direct busbar-to-busbar edge	Typed node between its endpoint busbars
Line state	Continuous edge attributes	Continuous edge attributes	Features of the line node
Minimum useful depth	One layer for adjacent busbar exchange	One layer for adjacent busbars; two for asset-to-remote-busbar flow	Two layers for adjacent busbar exchange through a line node
Hidden line state	None: the edge modifies a message	None: the edge modifies a message	Yes: each line obtains and updates an embedding
Far-end context in a local actor graph	Live neighboring busbar, including aggregated power and angle	Live neighboring busbar, but no neighboring generator or load nodes	None; the shared line node replaces the need for a far-end busbar
Node count	SB	$SB + N_g + N_d$	$SB + N_g + N_d + L$

5.2.1 Busbar Representation

The busbar schema is the most compact of the three. Busbars are the only nodes, active transmission lines form the graph edges, and generator and load quantities are aggregated into their currently assigned busbars.

Nodes and features

With S substations and at most B busbars each, one node is created per substation-busbar pair, written $i(s, b)$, including busbars that currently hold no asset.

Table 5.2 lists the six node features. Each generator and load contributes to the busbar it is con-




Figure 5.1: Busbar graph representation. Blue nodes denote busbars and black connections denote active transmission-line relations. Generator and load quantities are aggregated into the features of their associated busbars.

nected to: active powers are summed when several assets share a busbar, and their angles are averaged. Disconnected assets contribute nowhere. The substation cooldown is repeated on every busbar node of that substation.

Table 5.2: Features attached to each busbar node.

Feature	Definition
<code>gen_p</code>	Sum of active generator power connected to the busbar.
<code>gen_theta</code>	Mean voltage angle of generators connected to the busbar.
<code>load_p</code>	Sum of active load power connected to the busbar.
<code>load_theta</code>	Mean voltage angle of loads connected to the busbar.
<code>time_before_cooldown_sub</code>	Remaining substation cooldown.
<code>domain_mask</code>	One for a substation controlled by the agent and zero for a contextual neighbor.

Relations and edge features

The only physical relation in the busbar schema is the busbar-to-busbar transmission line. Table 5.3 gives the dynamic attributes copied onto each active line edge; the two maintenance features are present only for environments with scheduled maintenance.

Table 5.3: Features attached to physical-line edges.

Feature	Definition
<code>line_status</code>	Binary connection status of the physical line.
<code>rho</code>	Line loading relative to its thermal limit.
<code>timestep_overflow</code>	Number of consecutive time steps in overload.
<code>time_before_cooldown_line</code>	Remaining cooldown before another line action is legal.
<code>time_next_maintenance</code>	Time until scheduled maintenance, when available.
<code>duration_next_maintenance</code>	Duration of scheduled maintenance, when available.

5.2.2 Disaggregated Representation

The disaggregated schema keeps busbars as the backbone and adds one node per generator and load, joined to their busbar by typed directed relations. Transmission lines remain busbar-to-busbar edges, and generator and load states are preserved individually.




Figure 5.2: Disaggregated representation. Busbars are blue, generators are red, and loads are green. Physical transmission lines remain direct relations between busbar nodes.

Nodes and features

Busbar nodes are created exactly as in Section 5.2.1. In addition, each generator and load whose substation is present in the graph receives its own node.

Every busbar, generator, and load node is represented by the same ordered seven-dimensional feature vector

$$\mathbf{x}_v = [\text{is_busbar}, \text{is_generator}, \text{is_load}, p, \text{theta}, \text{time_before_cooldown_sub}, \text{domain_mask}]^T \in \mathbb{R}^7. \quad (5.1)$$

Table 5.4 shows the value placed in every shared column; a zero there is an explicitly zero-filled column, not an omitted feature.

Table 5.4: Values placed in the seven feature columns shared by every node in the disaggregated graph. Disconnected assets retain their node row but have zero power and angle.

Shared feature column	Busbar node	Generator node	Load node
is_busbar	1	0	0
is_generator	0	1	0
is_load	0	0	1
p	0	gen_p if connected; 0 otherwise	load_p if connected; 0 otherwise
theta	0	gen_theta if connected; 0 otherwise	load_theta if connected; 0 otherwise
time_before_cooldown_sub	Substation value	0	0
domain_mask	1 controlled; 0 contextual	1 controlled; 0 contextual	1 controlled; 0 contextual

Relations and edge features

Every candidate directed edge is represented by the same ordered feature vector. Its physical block is

$$\mathbf{p}_{uv} = [\text{line_status}, \text{rho}, \text{timestep_overflow}, \text{time_before_cooldown_line}, (\text{time_next_maintenance}, \text{duration_next_maintenance})_{\text{optional}}]^T. \quad (5.2)$$

The two maintenance columns are appended only when scheduled maintenance is available. The relation block is the same six-column vector for every edge:

$$\mathbf{r}_{uv} = [\text{relation_physical_line}, \text{relation_same_substation_busbar}, \text{relation_generator_to_busbar}, \text{relation_busbar_to_generator}, \text{relation_load_to_busbar}, \text{relation_busbar_to_load}]^T. \quad (5.3)$$

The complete common edge vector is therefore

$$\mathbf{e}_{uv} = [\mathbf{p}_{uv}^T \mid \mathbf{r}_{uv}^T]^T \in \begin{cases} \mathbb{R}^{10}, & \text{without maintenance columns,} \\ \mathbb{R}^{12}, & \text{with maintenance columns.} \end{cases} \quad (5.4)$$

An active physical-line edge fills $\mathbf{p}_{uv}$ with its measured line state. Every active structural or asset-attachment edge instead uses

$$\mathbf{p}_{\text{neutral}} = \mathbf{0}, \quad (5.5)$$

so a relation carrying no flow reports none. In every active edge row, exactly one column of $\mathbf{r}_{uv}$ is one and the other five are zero.

Table 5.5 shows the values placed in these common blocks for every relation.

Table 5.5: Values placed in the common edge-feature vector shared by every active directed relation in the disaggregated graph. “Other five” refers to the remaining columns of the six-column relation block.

Directed relation	Active rule	Shared physical block	Shared relation block
$i_o \rightarrow i_e$ and reverse (physical line)	Line online and candidate busbars match the current endpoints	Measured $\mathbf{p}_{uv}$	<code>relation_physical_line= 1</code> ; other five = 0
$i(s, b) \rightarrow i(s, b'), b \neq b'$	Optional; every ordered same-substation busbar pair	$\mathbf{p}_{\text{neutral}}$	<code>relation_same_substation_busbar= 1</code> ; other five = 0
$g \rightarrow i(s(g), b)$	Generator connected to candidate busbar b ; direction enabled	$\mathbf{p}_{\text{neutral}}$	<code>relation_generator_to_busbar= 1</code> ; other five = 0
$i(s(g), b) \rightarrow g$	Same mask; reverse direction enabled	$\mathbf{p}_{\text{neutral}}$	<code>relation_busbar_to_generator= 1</code> ; other five = 0
$d \rightarrow i(s(d), b)$	Load connected to candidate busbar b ; direction enabled	$\mathbf{p}_{\text{neutral}}$	<code>relation_load_to_busbar= 1</code> ; other five = 0
$i(s(d), b) \rightarrow d$	Same mask; reverse direction enabled	$\mathbf{p}_{\text{neutral}}$	<code>relation_busbar_to_load= 1</code> ; other five = 0

Generator and load directions are selected independently as `bidirectional`, `asset_to_busbar`, or `busbar_to_asset`. The default is `bidirectional`. Physical-line and same-substation relations remain `bidirectional`.

5.2.3 Disaggregated Line-Node Representation

The disaggregated line-node schema replaces each direct physical-line edge with a transmission-line node between its endpoint busbars, so busbars, generators, loads, and lines are all nodes.




Figure 5.3: Disaggregated line-node representation. Yellow nodes represent physical lines and connect the blue busbars at their endpoints; generators and loads remain separate terminal nodes.

Nodes and features

Every node type uses the same ordered feature vector:

$$\begin{aligned} \mathbf{x}_v = & [\text{is_busbar}, \text{is_generator}, \text{is_load}, \text{is_transmission_line}, \\ & \text{p}, \text{theta}, \text{line_status}, \text{rho}, \text{timestep_overflow}, \\ & \text{time_before_cooldown_line}, \\ & (\text{time_next_maintenance}, \text{duration_next_maintenance})_{\text{optional}}, \\ & \text{time_before_cooldown_sub}, \text{domain_mask}]^T, \\ \mathbf{x}_v \in & \begin{cases} \mathbb{R}^{12}, & \text{without maintenance columns,} \\ \mathbb{R}^{14}, & \text{with maintenance columns.} \end{cases} \end{aligned} \quad (5.6)$$

Table 5.6 specifies every entry; 0 denotes an explicitly zero-filled column, not an omitted feature. For generators and loads, power and angle are set to zero when the asset is disconnected.

Table 5.6: Values placed in the common node-feature vector of the disaggregated line-node graph, when voltage angles are carried as absolute per-asset values. The two maintenance rows are included only when maintenance observations are enabled. When the angle is instead carried as the drop along each line, the angle row changes owner: the generator and load entries become 0 and the transmission line holds the drop across it, set to 0 while the line is out (Section 5.1.2).

Shared feature column	Busbar	Generator	Load	Transmission line
is_busbar	1	0	0	0
is_generator	0	1	0	0
is_load	0	0	1	0
is_transmission_line	0	0	0	1
p	0	gen_p if connected; 0 otherwise	load_p if connected; 0 otherwise	0
theta	0	gen_theta if connected; 0 otherwise	load_theta if connected; 0 otherwise	0
line_status	0	0	0	Observed line value
rho	0	0	0	Observed ρ_ℓ
timestep_overflow	0	0	0	Observed line value
time_before_cooldown_line	0	0	0	Observed line value
time_next_maintenance	0	0	0	Observed line value
duration_next_maintenance	0	0	0	Observed line value
time_before_cooldown_sub	Substation value	0	0	0
domain_mask	1 controlled; 0 contextual	1 controlled; 0 contextual	1 controlled; 0 contextual	1 if either endpoint is controlled; 0 otherwise

Relations and edge features

Every candidate directed edge uses the same ordered nine-dimensional feature vector:

$$\begin{aligned} \mathbf{e}_{uv} = & [\text{relation_same_substation_busbar}, \\ & \text{relation_busbar_to_line_origin}, \text{relation_line_origin_to_busbar}, \\ & \text{relation_busbar_to_line_extremity}, \text{relation_line_extremity_to_busbar}, \\ & \text{relation_generator_to_busbar}, \text{relation_busbar_to_generator}, \\ & \text{relation_load_to_busbar}, \text{relation_busbar_to_load}]^T \in \mathbb{R}^9. \end{aligned} \quad (5.7)$$

These edges contain no continuous line-state attributes because the transmission-line node already contains rho, status, overflow, cooldown, and maintenance. For an active edge, exactly one relation column is one and the other eight are zero. An inactive candidate has all nine entries set to zero and is removed by edge_mask=0 before message passing. Table 5.7 gives the active rule and the nonzero entries for every relation.

Table 5.7: Nonzero entries in the common 9-column edge-feature vector of the disaggregated line-node graph. All relation columns not named in a row are zero.

Directed relation	Active rule	Relation column set to 1
$i(s, b) \rightarrow i(s, b'), b \neq b'$	Optional; every ordered same-substation pair	<code>relation_same_substation_busbar</code>
$i_o \rightarrow \ell$	Online line, current origin busbar, direction enabled	<code>relation_busbar_to_line_origin</code>
$\ell \rightarrow i_o$	Same mask; reverse direction enabled	<code>relation_line_origin_to_busbar</code>
$i_e \rightarrow \ell$	Online line, current extremity busbar, direction enabled	<code>relation_busbar_to_line_extremity</code>
$\ell \rightarrow i_e$	Same mask; reverse direction enabled	<code>relation_line_extremity_to_busbar</code>
$g \rightarrow i(s(g), b)$	Generator attached to b ; direction enabled	<code>relation_generator_to_busbar</code>
$i(s(g), b) \rightarrow g$	Same mask; reverse direction enabled	<code>relation_busbar_to_generator</code>
$d \rightarrow i(s(d), b)$	Load attached to b ; direction enabled	<code>relation_load_to_busbar</code>
$i(s(d), b) \rightarrow d$	Same mask; reverse direction enabled	<code>relation_busbar_to_load</code>

5.2.4 Local actor information boundary and the global state graph

Each agent a controls a set of substations $\mathcal{C}_a$. A boundary line connects $\mathcal{C}_a$ to another agent’s domain, and its far-end busbar is the only possible neighboring busbar. The local graph exposes only the boundary information required by its representation.

Why one hop, and why not more. The one-hop boundary is a methodological choice: it keeps each actor regional while the centralized critic still observes the global state. Power flow has no natural one-hop cutoff, but the loading and angle drop of boundary lines provide the immediate interaction with the exterior without exposing neighboring asset states. One hop is therefore not claimed to be optimal.

Visibility and masking

A contextual busbar is visible only while a live boundary line connects it to the controlled region. When that connection is lost its feature row is zeroed, its incident relations go inactive, and it is left out of the readout — including the denominator when the readout is a mean. It therefore reaches neither message passing nor the graph embedding, and a same-substation relation cannot restore it.

Controlled nodes stay visible when disconnected, and in the disaggregated line-node schema a shared line node stays visible too, since it belongs to the agent’s own graph, though its attachment relations are inactive while the line is offline. Relations are active only when the line is online and the current busbar assignments match their endpoints. Neighboring generators and loads are never constructed.

Representation-dependent information

The representations are not information-equivalent: Table 5.8 states what each one exposes of a neighboring domain.




Figure 5.4: One live boundary line, seen under each representation. In every panel the blue dashed region is the agent’s controlled domain and the orange dashed region is what lies beyond it. The busbar graph (a) reaches a far-end busbar carrying the neighbor’s generation and load aggregated onto it. The disaggregated graph (b) reaches the same busbar, but no neighboring generator or load node is ever constructed, so that busbar carries no measurement of its own. The disaggregated line-node graph (c) reaches only the shared line node, and no far-end busbar exists at all.

Table 5.8: Information from another agent’s domain visible to a local actor.

Representation	Neighbor information available
busbar	Shared-line edge and attached far-end busbar, including aggregated generator/load power and angles
disaggregated	Shared-line edge and attached far-end busbar; no neighboring generator/load nodes or private-asset state
disaggregated line-node	Shared-line node only; no neighboring busbar, equipment, or far-end attachment

Three consequences follow. In the busbar and disaggregated schemas the far-end busbar has to exist at all, because an edge needs both endpoints for the boundary line’s features to reach the controlled busbar through message passing; the remaining busbar candidates at that substation stay masked, and if the line disconnects none of them remains visible. In the disaggregated schema that contextual busbar carries only type and role indicators, since power and angle sit on asset nodes and neighboring assets are never constructed — all neighboring measurement therefore arrives through the boundary-line edge. The disaggregated line-node schema drops the contextual busbar entirely and puts the shared measurements on the line node, which stays visible with `line_status=0` when the line disconnects while its attachment edges go inactive.

These boundaries constrain the actors only. The centralized critic reads the normalized flat state of the whole grid during training, while execution uses each actor’s local graph, as required by MAPPO’s centralized-training, decentralized-execution setup [7].

5.2.5 Graph actor and centralized critic




Figure 5.5: End-to-end graph-actor architecture. Candidate edges masked by the current topology are removed before message passing. The graph encoder and its output projection are shared across agents, whereas the final MLP has one output per action available to agent a and is therefore agent-specific.

In the reported configuration, the input node projection and the two message-passing layers have 128 features. LayerNorm is applied after the input projection and after each message-passing layer. The graph readout produces a 128-dimensional embedding, followed by an agent-specific action head with two 128-unit hidden layers. The graph encoder and output projection are shared across agents, whereas the action heads are separate because their action-set sizes may differ.

This is the default configuration. The candidate-scoring alternative of Section 5.5 replaces this head.

The centralized critic is a separate MLP over the normalized flat state. It does not reuse the actor graph embedding and is used only during centralized training. Its flat-state normalization is independent of the graph-input normalization. The exact GINE update, graph readout, output projection, and action-head equations are given in Appendix B.

Purpose: Define how the electrical state becomes the input of each actor and how that input is encoded.

- **Input processing:** Continuous physical quantities can be divided by fixed physical scales and optionally standardized from training statistics. The loading ratio `rho`, binary features, and masks remain unchanged. The graph-screening experiments use raw features; the cross-grid transfer study uses deterministic physical scaling.
- **Dynamic topology:** Nodes and candidate edges remain fixed, while time-dependent edge and node masks retain only the relations and contextual nodes that are electrically connected to the controlled region. An inactive contextual candidate has a zero feature row and participates in neither message passing nor readout.
- **Graph representations:** The busbar schema aggregates equipment on busbars; disaggregated gives generators and loads their own nodes; and disaggregated line-node also represents each transmission line as a node. This changes both the feature location and the number of message-passing steps between physical objects.
- **Actor and critic:** Each actor receives its controlled region and only the schema-specific boundary context of Table 5.8. The reference architecture uses a two-layer, 128-feature GINE encoder shared across agents, followed by an agent-specific action head. The centralized critic remains a separate MLP over the normalized full-grid flat state.

Takeaway: The graph variants change both how state is represented and how much neighboring information is available. The MAPPO training structure and agent partition remain unchanged, but the three representations must not be described as information-equivalent.

5.3 Substation Representation: Direct Edges and Summary Nodes

Two optional factors change how information is exchanged inside a substation. The factor e adds direct same-substation busbar edges. The factor n adds one learned substation-summary node per included substation. Both augmentations are part of the constructed graph; their nodes and edges therefore enter the same node projection, edge filtering, and message-passing pipeline as the base graph.




Figure 5.6: The four combinations of the independent substation factors e and n . The hollow node b_2 is controlled but non-energized. Factor e adds a bidirectional relation between the controlled busbars, while factor n connects a summary node u_s to both. Neither augmentation extends to the contextual busbar b_N , which retains only its physical relation.




Figure 5.7: Placement of the two augmentations when $e = n = 1$. In every graph representation, the direct relation and summary node attach only to controlled busbars. Equipment and line nodes retain their ordinary busbar relations, and no augmentation is added in the contextual domain.

5.3.1 Direct same-substation edges (e)

In addition to physical-line relations, the augmentation inserts $i(s, b) \rightarrow i(s, b')$ for every ordered pair $b \neq b'$ of busbars in the same *controlled* substation. The edges are therefore bidirectional between each pair of busbars, remain active throughout the episode, and do not represent an electrical connection. Consequently, a non-energized controlled busbar remains connected to the other busbar of its substation through this structural relation, without being considered electrically energized. Figure 5.7 shows this relation across the three graph representations.

Table 5.9 gives the exact feature values of the added edges. Every physical attribute is zero: a relation that carries no flow reports none.

Table 5.9: Node and edge features introduced by the direct same-substation augmentation. “None” means that the augmentation neither adds nodes nor changes the features of existing nodes.

Graph schema	Added node features	Feature values on each added directed edge
Busbar	None	Every entry is 0: this schema has no relation columns. An active physical line always carries <code>line_status= 1</code> , so an all-zero row is itself the signature of a structural relation.
Disaggregated	None	<code>relation_same_substation_busbar= 1</code> ; every other entry is 0.
Disaggregated line-node	None	<code>relation_same_substation_busbar= 1</code> ; every other entry is 0.

Why controlled substations only. The relation represents the actor’s ability to move elements between busbars, not an electrical connection. Adding it inside a neighboring substation would encode another agent’s action space and create a nonphysical message path.

5.3.2 Substation summary nodes (n)

Augmentation n adds one learned summary node u_s per controlled substation and connects it in both directions to every visible busbar of that substation, adding one node and $2B$ directed edges each. Two busbars of the same substation then communicate in two steps, $i(s, b) \rightarrow u_s \rightarrow i(s, b')$: the same intra-substation role as the direct relation of Section 5.3.1, one message-passing step slower, and restricted to the controlled domain for the same reason. Only busbars attach to u_s , and only visible ones, so the summary carries no contextual state of its own. Figure 5.7 shows both augmentations across the three representations.

The model initializes u_s from a *structural descriptor* rather than from the substation’s identity. Let $\phi_s \in \mathbb{R}^4$ hold its numbers of busbars, incident lines, attached generators, and attached loads. The last three are divided by a fixed 4, the same on every grid, so a count means the same thing whatever network it comes from; the busbar count is divided by the grid’s busbars per substation and is therefore 1 throughout both grids studied here, leaving the descriptor to vary through its other three entries:

$$u_s^{(0)} = \mathbf{W}_{\text{sub}} \phi_s + \mathbf{b}_{\text{sub}}, \quad \mathbf{W}_{\text{sub}} \in \mathbb{R}^{d_h \times 4}, \quad \mathbf{b}_{\text{sub}} \in \mathbb{R}^{d_h}. \quad (5.8)$$

The summary node therefore holds no grid measurement of its own; messages from its busbars are what give it an observation-dependent state.

An embedding table indexed by substation identity would have shape $S \times d_h$ and could not transfer between grids with different S . Equation 5.8 applies one shared projection to four fixed-scale counts instead, so the parameter shapes are independent of grid size while substations of different structure still start from different states. This keeps the augmentation compatible with the transfer study of Chapter 8 if necessary.

The augmentation behaves identically in all three schemas: every $i \leftrightarrow u_s$ edge carries an all-zero attribute vector, since no schema defines a summary relation column. Each u_s counts as one node in a mean readout. The two augmentations may be enabled together, and either may be combined with the virtual-node readout of Section 5.4.2.

5.4 Pooling and Graph Readout

The final graph embedding is obtained either by pooling node embeddings directly or by reading out a learned virtual node. The output projection and action head are the same in both cases.

5.4.1 Mean and maximum pooling

A terminal readout is two independent choices: *which* nodes are aggregated, and *how*. Write $w_v \in \{0, 1\}$ for the resulting weight of node v . It is the visibility mask m_v of Section 5.2.4, optionally intersected with the static controlled-node mask c_v , and optionally with an energized mask e_v that requires the node to be incident to an active electrical relation. Self and same-substation relations do not count as electrical. The reducer is an average or a per-channel maximum,

$$\bar{\mathbf{h}}_G^{\text{mean}} = \frac{\sum_v w_v \mathbf{h}_v^{(K)}}{\max(1, \sum_v w_v)}, \quad \left(\bar{\mathbf{h}}_G^{\text{max}}\right)_c = \max_{v: w_v=1} \left(\mathbf{h}_v^{(K)}\right)_c. \quad (5.9)$$

An unnormalized sum is not a candidate: its magnitude grows with the number of pooled nodes, which differs across agents, representations, and grids, so the scale of the embedding would change with the graph it came from and again under the transfer of Chapter 8. A single readout is taken over all node types at once; there is no separate aggregation per type. A node with $w_v = 0$ contributes neither to a numerator nor to a denominator, so masking is an enforcement of the construction boundary rather than an averaging of zero placeholders. The two restrictions combine freely, which gives the eight names of Table 5.10.

Table 5.10: The eight terminal readouts, named `[controlled_][energized_]{mean,max}`. Rows choose the pooled population, columns the reducer. The virtual-node readout of Section 5.4.2 replaces terminal pooling entirely.

Pooled nodes (w_v)	mean	max
All visible, m_v	mean	max
Controlled, $m_v c_v$	controlled_mean	controlled_max
Energized, $m_v e_v$	energized_mean	energized_max
Controlled and energized, $m_v c_v e_v$	controlled_energized_mean	controlled_energized_max

The restrictions act on the readout only, never on message passing, so a contextual node excluded from w_v still reaches the controlled nodes through active relations. They are therefore pooling ablations, not information-boundary fixes.

What each axis trades. The two axes answer different questions, so their trade-offs are separate (Table 5.11). Pending the ablation of Section 6.5, the most defensible default is `controlled_mean`: stable in population, aligned with the action domain, and retaining the empty busbars that actions may target. `controlled_max` is the risk-sensitive alternative when a decision is expected to hinge on one critical node. Without one-hop context the controlled and visible sets coincide, and in the disaggregated line-node schema the context option changes little because boundary line nodes are already present.

Table 5.11: Trade-offs along each axis of Table 5.10.

Choice	Main advantage	Main limitation
<i>Reducer</i>		
mean	Smooth regional summary; gradients reach every pooled node.	A rare but critical activation is diluted.
max	Keeps the strongest activation and ignores how many nodes there are.	Sparse gradients, and channel maxima may come from unrelated nodes.
<i>Pooled population</i>		
All visible	Boundary conditions enter the readout directly.	Context can dilute a mean or dominate a max, and the population moves with boundary topology.
Controlled	Aligned with the action domain and fixed for an episode; context still arrives through message passing.	Relies on message passing to carry every useful boundary signal.
Energized	Focuses on the network currently carrying power.	Drops available action targets, and the population changes after each switching action.

5.4.2 Virtual-node graph readout

A virtual-node readout replaces terminal pooling with one learned graph-summary node v_* , and works with any of the three schemas. The graph gains that node plus one directed edge into it from every member of a summary set $\mathcal{V}_{\text{sum}}$: the controlled substation nodes when the augmentation of Section 5.3.2 is enabled, and the visible controlled busbars otherwise. Every added edge points *toward* v_* in all reported runs, so it aggregates node states and never influences them. It has no direct edge to contextual busbars, generators, loads, or line nodes; those reach it only through ordinary message-passing paths into the controlled region.

Its initial state $\mathbf{h}_*^{(0)}$ is learned, and the same message-passing layers that update the physical nodes update it. After K layers, $\mathbf{h}_*^{(K)}$ enters the usual output projection in place of the pooled vector, so the embedding dimension is unchanged. The added edges are structural. Without substation nodes the augmentation adds one edge per visible controlled busbar; with them, one edge per controlled substation.




Figure 5.8: Virtual-node readout in the reported one-way setting. Without substation nodes, every controlled busbar—including non-energized b_{32} —sends directly to v_* . With augmentation n , the same busbars communicate through u_s , and only the substation nodes send to v_* . The final state $\mathbf{h}_*^{(K)}$ replaces terminal pooling.

Chapter Summary (2/3): Augmentations and Graph Readout

Purpose: Separate changes to information exchange inside the graph from changes to the final graph-level representation.

- **Direct substation edges (e):** Add structural relations between busbars of the same substation, without adding nodes or representing a physical transmission line.
- **Substation nodes (n):** Add one learned node per *controlled* substation. It has no raw electrical measurement, is initialized from a structural descriptor, and is updated from its busbars.
- **Terminal pooling:** A readout is a pooled population — all visible nodes, or the controlled and/or energized subsets — reduced by a mean or a per-channel maximum. Restricting the population changes the readout only: excluded context still reaches the controlled nodes through message passing.
- **Virtual readout (v):** Replace terminal pooling with one learned grid-summary node. Messages travel only toward this node, from the busbars or from the added substation nodes, and its final state becomes the graph embedding.

Takeaway: Factors e and n modify message passing, whereas pooling and factor v determine how the encoded local graph is reduced to one vector for action selection.

5.5 Candidate-Action Scoring from the Disaggregated Line-Node Graph

The default action head produces one output column per discrete action. That column is tied to an action index and knows nothing about the physical objects the action modifies. This section describes an alternative head that scores each action from the graph nodes it touches. It is available only for the disaggregated line-node schema of Section 5.2.3; the default remains the action-index head of Appendix B.4. The PPO objective, the categorical policy, sampling, and deterministic evaluation are unchanged.

5.5.1 From action indices to candidate scoring

Throughout this section, $\mathbf{z}_a$ denotes the graph embedding of agent a : the node states of its local graph pooled as in Section 5.4 and passed through the output projection (Equation B.7). It is one 128-dimensional vector per agent per step, and it is the only thing the actor has ever seen of the graph.

With the default head, agent a produces its logit vector in one step from $\mathbf{z}_a$:

$$\ell_a = \text{MLP}_a(\mathbf{z}_a), \quad \ell_a \in \mathbb{R}^{|\mathcal{A}_a|}. \quad (5.10)$$

The last layer has one row per action index, so what makes an action good must be learned separately for every index, and the head grows with $|\mathcal{A}_a|$.

Candidate scoring replaces that by one scalar-valued network shared across all non-idle actions. The idle action may use the same network or a dedicated one:

$$\ell_{a,j} = \text{score} \left(\underbrace{\mathbf{z}_a}_{\text{whole graph}}, \underbrace{\mathbf{c}_j}_{\text{nodes action } j \text{ touches}}, \underbrace{\phi_j}_{\text{optional action descriptor}} \right), \quad j \geq 1. \quad (5.11)$$

The same parameters score all non-idle actions, so the head no longer grows with $|\mathcal{A}_a|$, and two actions are ranked by the difference between their touched-node contexts rather than by two independent output rows. This is what makes the line nodes usable: an action next to a loaded line is scored from that line’s own embedding.

Three choices combine independently: how the touched nodes are pooled into $\mathbf{c}_j$, whether the static descriptor ϕ_j is appended, and whether the idle action uses the shared scorer or its own. The two uniform pooling rules therefore give eight variants, and learned attention adds a third pooling family. Figure 5.9 shows how the three enter one candidate score; the following subsections define them.




Figure 5.9: Candidate-action scoring for an endpoint-move candidate. Stage 1 is the shared GNN encoder, given in full in Figure 5.5: it maps the local graph to node states $\mathbf{H}$, and the same stack yields the global embedding $\mathbf{z}_a$ by masked readout and projection, so the two inputs to the scorer come from one set of weights rather than two. In stage 2 the violet outlines mark the only rows gathered from $\mathbf{H}$ for the local context: both busbars of the modified substation and the touched line node. One pooling construction then produces $\mathbf{c}_j$. The static descriptor ϕ_j may be appended before the shared non-idle scorer. The idle action uses either zero local context in that scorer or a dedicated scorer.

5.5.2 What each action touches

The first ingredient is one lookup table per agent, with one row per action, answering a single question: which nodes of the local graph does this action physically touch? It is built once, when the environment is created, and never recomputed during training.

A Grid2Op topology action states which entries of the topology vector it changes, and each entry belongs to a known object: a line endpoint, a generator, or a load. Reading those entries gives the objects the action modifies, and the graph builder’s row maps turn each object into the row of the graph that holds its embedding. Table 5.12 follows one action through this.

Table 5.12: One action, from Grid2Op to node rows: moving the origin end of line 4 onto the other busbar of substation 6. Row indices are illustrative.

Step	Result for this action
Objects the action changes	Substation 6, line 4
Their rows in the graph	Both busbars of substation 6, here 12 and 13, and the line node of line 4, here 22
Stored for scoring	Busbar [12, 13], line [22], generator and load empty

Two decoding rules are worth stating. A touched substation contributes *all* of its busbar rows, not only the busbar in use, because the action changes which of them the equipment sits on. A boundary line contributes its line-node row and the controlled endpoint rows that are present in the local graph; scoring does not add the far-end busbar or any neighboring private asset. The rows are stored as fixed-size arrays with a mask, padded to the widest action, so that a single gather serves the whole action space.

Write $\mathcal{N}_{\text{bus}}(j)$, $\mathcal{N}_{\text{line}}(j)$, $\mathcal{N}_{\text{load}}(j)$, and $\mathcal{N}_{\text{gen}}(j)$ for the four sets of node rows that action j touches. Because a topology change moves line endpoints, the line node must be present whenever the line touches the agent’s controlled domain. The leakage-free construction guarantees that membership without constructing the non-controlled endpoint.

Each action also carries a static feature vector $\phi_j \in \mathbb{R}^{12}$, listed in Table 5.13. It describes what the action does independently of the current grid state, so the scorer can tell a single endpoint move from an action that reconfigures a whole substation even when both touch the same nodes.

Table 5.13: The static action-feature vector ϕ_j . The five indicators are binary; the seven counts are divided by their largest value over that agent’s action space, so every entry lies in $[0, 1]$ and is independent of the grid size.

Entry	Meaning
<i>Indicators</i>	
<code>is_do_nothing</code>	The action is the idle action
<code>has_busbar_context</code>	It modifies at least one substation
<code>has_line</code>	It touches at least one line
<code>has_load</code>	It touches at least one load
<code>has_generator</code>	It touches at least one generator
<i>Scaled counts</i>	
<code>n_substations</code>	Substations modified
<code>n_topology_changes</code>	Topology-vector entries changed
<code>n_line_status_changes</code>	Lines connected or disconnected
<code>n_line_endpoint_changes</code>	Line endpoints moved to another busbar
<code>n_generator_changes</code>	Generators moved to another busbar
<code>n_load_changes</code>	Loads moved to another busbar
<code>n_other_changes</code>	Changed objects of any other type

The descriptor is optional in the final scorer. Including it tells the model what operation is proposed and can distinguish actions that touch similar nodes; omitting it isolates what can be learned from the graph state and the touched-node context alone. This choice is independent of the pooling rule and of the idle-action treatment.

5.5.3 Pooling the touched nodes

The second ingredient turns the touched rows into $\mathbf{c}_j$. In $\mathbf{z}_a$ the nodes have already been averaged away, so the encoder also returns the node states themselves:

$$(\mathbf{z}_a, \mathbf{H}) = \text{encoder}_{\text{nodes}}(\mathcal{G}_t^{(a)}), \quad \mathbf{H} \in \mathbb{R}^{|\mathcal{V}| \times d_h}, \quad (5.12)$$

where row v of $\mathbf{H}$ is the state $\mathbf{h}_v^{(K)}$ of node v after the last message-passing layer. These are exactly the vectors that mean pooling averages into $\mathbf{z}_a$; nothing new is computed, and a run with the default head is unaffected. Only the physical nodes appear in $\mathbf{H}$: substation-summary and virtual nodes are excluded, since an action touches equipment and not a learned hub.

Write $\mathcal{N}_t(j)$ for the rows of type $t \in \{\text{bus, line, load, gen}\}$ that action j touches. The masked mean over any such set is

$$\mathbf{c}_j^{(t)} = \frac{1}{\max(1, |\mathcal{N}_t(j)|)} \sum_{v \in \mathcal{N}_t(j)} \mathbf{h}_v^{(K)}, \quad (5.13)$$

where the clamped denominator makes an untouched type give the zero vector instead of a division by zero. The joint rule applies this once to all touched rows regardless of type, giving width d_h ; the type-separated rule applies it per type and concatenates the four results, giving width $4d_h$ and preserving node roles (Table 5.14).

Table 5.14: Implemented constructions of the action-specific node context.

Construction	Eligible rows	Aggregation	Width
Joint uniform mean	All nodes touched by the action	One equal-weight mean; node roles are mixed	d_h
Type-separated uniform mean	All nodes touched by the action	One equal-weight mean per type, then concatenation	$4d_h$
Type-separated learned attention	Determined by the attention scope	One action-conditioned weighted sum per type, then concatenation	$4d_h$

With learned attention, an action-specific query assigns normalized weights within each node type:

$$\mathbf{c}_j^{(t)} = \sum_{v \in \mathcal{E}_t(j)} \alpha_{jv}^{(t)} \mathbf{v} \mathbf{h}_v^{(K)}, \quad \sum_{v \in \mathcal{E}_t(j)} \alpha_{jv}^{(t)} = \mathbb{I}[|\mathcal{E}_t(j)| > 0]. \quad (5.14)$$

Here $\mathcal{E}_t(j)$ is the eligible set; an empty set returns the zero vector. The implemented choices are shown in Table 5.15. The action query is formed from the global graph state and the static action descriptor in the reported attention comparisons. This is separate from whether the descriptor is also appended to the final scorer.

Table 5.15: Possible eligibility rules for learned action-specific attention.

Eligibility rule	Nodes allowed to receive weight	Purpose
Fixed physical support	Only nodes identified as touched by the action	Learn their relative importance without changing the physical mapping.
Physical support as a prior	Every local node of the same type, with a score bonus for touched nodes	Allow attention to extend the physical mapping while favoring it initially.
Fully learned support	Every local node of the same type, without a score bonus	Test whether the model can discover the relevant nodes without the mapping as a prior.

The two uniform means are parameter-free controls. Learned attention can retain a critical node that a mean would dilute, but adds parameters and must also learn a sensible weighting.

5.5.4 Scoring and the idle action

All non-idle actions are scored by one network with a single output unit. Its input depends on whether the static action descriptor is included:

$$\mathbf{x}_{a,j} = \begin{cases} [\mathbf{z}_a; \mathbf{c}_j; \phi_j], & \text{with the descriptor,} \\ [\mathbf{z}_a; \mathbf{c}_j], & \text{without it,} \end{cases} \quad \ell_{a,j} = \text{MLP}_{\text{score}}(\mathbf{x}_{a,j}), \quad j \geq 1. \quad (5.15)$$

The idle action touches no node, so its local context is the zero vector. It either reuses the same scorer or gets a dedicated one:

$$\ell_{a,0} = \begin{cases} \text{MLP}_{\text{score}}(\mathbf{x}_{a,0}) + b_{\emptyset}, & \text{shared, with } \mathbf{c}_0 = \mathbf{0}, \\ \text{MLP}_{\emptyset}(\mathbf{z}_a) + b_{\emptyset}, & \text{dedicated.} \end{cases} \quad (5.16)$$

In both cases $b_{\emptyset}$ is a learned scalar applied only to the idle action. Sharing the scorer preserves parameter sharing across the full action set, but presents it with an empty context not seen for non-idle actions. A dedicated scorer instead treats “do nothing” as a global safety decision, at the cost of extra parameters and a different scoring mechanism.

5.5.5 Cost and current limits

The candidate contexts form a tensor of shape $[\text{batch}, |\mathcal{A}_a|, d_c]$, so memory grows with the size of the action space. This is the reason to start on **bus14** and on reduced action spaces before the full WCCI one.

The compared variants do not have identical parameter counts: separating node types widens the scorer input, static descriptors add input columns, learned attention adds projections, and a dedicated idle scorer adds another network. These counts should therefore be reported. When the dedicated idle scorer is used, its different mechanism also makes the action-0 frequency an explicit comparison metric rather than something assumed comparable.

The head requires the disaggregated line-node graph and does not combine with the intervention gate, which is refused at construction rather than silently ignored. One-hop context makes no difference here: the neighbor-inclusion setting adds no far-end equipment to this schema, because the line node already supplies the shared-line state.

5.6 Action-Space Reduction for Large Grids

The representations above change what an agent sees. This section changes what it can do, and is what makes the larger grids tractable at all: the WCCI bus36 setting has a far larger discrete topology-action space than `bus14`, so before training MAPPO on it the action space is reduced offline from simulated action outcomes. During rollouts, states are collected when the grid is stressed,

$$\rho_{\max} > 0.90.$$

For each collected state s , candidate actions are simulated and compared against do nothing:

$$\Delta(s, a) = \rho_{\text{after}}(s, a) - \rho_{\text{after}}(s, 0).$$

An action is counted as useful when it lowers the post-action loading relative to do nothing, i.e. $\Delta(s, a) < -\epsilon$. Each action is then scored by its empirical improvement rate,

$$\text{score}(a) = \frac{\#\{\text{tested states where } a \text{ improves the grid}\}}{\#\{\text{tested states where } a \text{ is evaluated}\}},$$

with a minimum-count filter to avoid selecting rarely sampled actions. The final reduced action space keeps the k best-scoring actions per agent, while always preserving action 0; an agent with fewer than k actions keeps all of them.

Choosing k is an empirical question rather than a design one, so it is settled where the reduced space is used. The candidate sets are sanity-checked with a one-step simulator-greedy controller that, at each risky state, simulates the available reduced actions and executes the best improving one. That controller also sizes the reduction, since it shows how much survival the action set can still reach once cut. Section 8.1 reports the resulting sizes and the value of k used for the transfer experiments.

Chapter Summary (3/3): Action Scoring and Large-Grid Reduction

Purpose: Define how graph embeddings become action logits and how the large WCCI action space is made tractable.

- **Default action head:** An agent-specific MLP maps the graph embedding directly to one logit per discrete action.
- **Candidate-action head:** The optional disaggregated line-node variant scores each non-idle action from the encoded nodes it modifies and from static action descriptors. The idle action has a separate score. The PPO objective and categorical policy are unchanged.
- **Current scope:** Candidate scoring requires the disaggregated line-node graph, and its memory cost still grows with the number of candidate actions. One-hop context is optional, since the line node already carries the shared measurements.
- **WCCI reduction:** Actions are ranked offline by how often they reduce loading in simulated stressed states. The best k per agent are kept and action 0 is always retained; a one-step simulator-greedy controller checks the reduced sets and fixes k before policy training.

Takeaway: Graph representation, action scoring, and action-space size are separate design choices. This chapter defines them; the following chapters evaluate the resulting policies.

Chapter 6

Graph-Design Screening

This chapter begins by verifying that a shared GNN actor can match the tuned MLP baseline, then reports the graph-design screens. Screens A and B are the completed seed-0 campaigns on the corrected information boundary; Screens C and D retain the earlier three-seed legacy campaigns; Screen E is a seed-0 factorial on the corrected boundary. All ask whether a specific choice in the actor of Chapter 5 changes survival on held-out chronics.

Information-boundary status of the reported screens

The preliminary viability runs and Screens C–D predate the corrected local-graph construction of Section 5.2.4. Their fixed candidate graphs marked all rows of a neighboring substation as visible, including unattached busbars and, in disaggregated graphs, neighboring private assets. Those isolated rows could enter graph readout even when no live boundary line connected them to the agent. Screens C–D therefore characterize the *legacy, information-leaking observation* and must not be presented as strictly comparable to checkpoints trained with the corrected boundary. Screens A, B, and E use connected context and controlled-only structural relations, and are the leakage-free results reported in this chapter. The corrected Screens A and B supersede the earlier campaigns under those names.

- **Screen A** varies the depth and width of the actor encoder, which applies to every graph schema: a 3×4 grid, 12 seed-0 runs. *Complete*.
- **Screen B** varies which nodes enter graph readout and whether they are reduced by a mean or a maximum: a 4×2 factorial, seven new seed-0 runs plus one reused Screen-A cell. *Complete*.
- **Screen C** varies the three structural augmentations of the busbar graph – same-substation edges, substation-summary nodes, and virtual-node readout (Sections 5.3 and 5.4.2): a 2^3 factorial, 24 runs. *Complete*.
- **Screen D** varies the direction of the generator and load attachment relations in the disaggregated graph (Section 5.2.2): a 3×3 factorial, 27 runs. *Complete*.
- **Screen E** varies candidate-node pooling, the action descriptor, and the idle-action head on the leakage-free disaggregated line-node graph: a 2^3 factorial, eight runs at seed 0, repeated with corrected graph inputs. *Complete*.

Screens A and B form one corrected-boundary sequence on the busbar graph and share the nominal mp2_h128 GINE/mean reference. Screen E is the other corrected-boundary screen, on the disaggregated line-node graph. Screens C and D are kept as separate legacy evidence and are not compared numerically with the corrected campaigns.

6.1 Preliminary GNN Encoder Viability

Before screening individual graph-design choices, an exploratory study checks that a graph actor can match the tuned MLP baseline. This is a prerequisite for transfer: the agents must be able to share one size-independent encoder without sacrificing survival.

Question. Can a shared graph encoder perform as well as the MLP? If each agent requires a separate encoder, or if the GNN cannot train stably, its structural transferability is lost.

Baseline and changes. The study compares several GINE variants. It varies encoder size (light or heavy), actor sharing (shared or unshared encoder), critic architecture (MLP or GNN), and the training optimizations established for the MLP baseline, including optimized critic updates and reward normalization.

Protocol. All GINE runs use `gnn_type = gine`, 15,000,000 total timesteps, evaluation every 80,000 timesteps, 2,000 rollout steps, deterministic evaluation, and the same 80/20 chronic split used by the MLP experiments. GINE is used because it incorporates edge features and is expressive over graph structures [19]. The complete graph-actor pipeline is described in Appendix B; Appendix D gives the full configurations in Tables D.1 and D.2.

Evidence. Table 6.1 and Figure 6.1 compare the main architectural improvements with the MLP Baseline. Figure 6.2 shows the broader search over isolated changes.

Table 6.1: GINE rerun survival summary.

Run	Intended change	Final survival (%)	Peak survival (%)
Heavy GINE Baseline	Historical GINE baseline	59.2	93.9
Light GINE (Concat, MLP Critic)	Light GINE with MLP critic	78.9	90.4
No Bias, Opt Critic	Optimized critic and bias 0.0	82.5	85.7
Optimized Light GINE	Light optimized GINE with MLP critic	95.7	97.4




Figure 6.1: Seed-aggregated test episodic survival comparing the MLP Baseline with high-performing GINE architectural variants.

Interpretation. The broad search gives three conclusions:

- **Shared encoder viability.** The unshared-actor variant does not outperform the shared variants, so one encoder can serve all agents without becoming the performance bottleneck.
- **Critic architecture.** A GNN centralized critic is unstable and degrades performance. The strongest variants pair the shared GNN actor with the standard MLP critic.
- **Optimization dependence.** Optimized critic updates, reward normalization, and removal of the initial do-nothing bias are also important for stable GNN training.

The strongest result is the **Optimized Light GINE**, which combines the shared light GINE actor, the MLP critic, and the baseline optimizations to reach 95.7% final survival. A shared GNN is therefore competitive on `bus14`, justifying the representation screens that follow.

6.2 How the Screens Are Organized

Each screen is presented as one block containing its design and its outcome together:

1. **Question.** Which design choice is uncertain?
2. **Design.** Which factors are varied, at which levels, and what is held fixed.




Figure 6.2: GINE survival comparisons. Each subplot compares the MLP Baseline with one variant from `configs/gine_s0_s1_s2`. Full hyperparameters and architectural details are given in Appendices B and D.

3. **Result.** The measured tables.
4. **Reading.** What follows, and what does not.

6.3 Shared Protocol

Table 6.2 separates the common optimization protocol from the boundary and seed settings that differ by campaign. Within a screen, every setting not named as a screened factor is fixed, so each screen is a comparison against its own baseline cell.

Table 6.2: Common optimization settings and campaign-specific boundary and seed settings.

Component	Fixed setting
Environment	bus14 , difficulty 0, topology actions, decentralized agents, normalized observations and rewards, 80/20 train/test chronic split with split seed 0
Actor	One GNN encoder shared by all agents; LayerNorm; node pre-encoder; two 128-unit action-head layers
Information boundary	Screens A, B, E: connected context and controlled-only structural relations (corrected); Screens C, D: legacy whole-substation one-hop context
Critic	Centralized MLP with three 256-unit hidden layers
Graph preprocessing	None: neither deterministic physical scaling nor running standardization
Training budget	15,000,000 environment steps, 72 parallel environments, 576 rollout steps
MAPPO update	10 epochs, 16 minibatches, actor and critic learning rates 10^{-4} with linear annealing, $\gamma = 0.9$, $\lambda = 0.95$, clip 0.2, target KL 0.02, entropy coefficient 0.01
Exploration	No initial do-nothing prior and no action-0 logit bonus
Sparse control	Disabled: no intervention gate, no intervention penalty, no evaluation-time rho heuristic
Evaluation	Deterministic evaluation every 82,944 environment steps, 10 episodes on each split
Seeds	Screens A, B, E: 0; Screens C, D: 0, 1, 2

6.3.1 Endpoint statistics

Survival and the train/test split are as defined in Section 3.1: survival is the percentage of an episode’s maximum length that the agent keeps the grid alive, reported on the test split. What differs here is how a run is reduced to one number. Each run saves the checkpoint with the highest test survival seen during training. That checkpoint is then re-evaluated deterministically over all 201 held-out test chronics.

Every run reaches the 15M-step budget, to within a negligible margin, so the screens do not report how far their runs got. Screens A and B report one seed per cell; Screens C and D aggregate three seeds per cell. Every seed-level endpoint is evaluated over the same 201 held-out chronics.

6.3.2 Compute

The campaigns were executed on JED (CPU) and IZAR (GPU). Within each factorial, the execution setup was held fixed across the compared configurations; the three-seed legacy campaigns also distribute seeds across both systems.

Evaluation is deterministic at a fixed checkpoint and seed. The corrected Screens A and B nevertheless have only one training seed per cell, so their differences are descriptive rather than estimates of between-seed effects.

What limits the campaign is simulation throughput rather than model size. A training step is dominated by stepping 72 Grid2Op environments, so the actor forward and backward passes are a small part of the wall clock: across the depth and width grid of Section 6.4 the actor varies by a factor of three in

parameters, 83.6k to 244.9k, without a comparable effect on run time. Parameter count is therefore not the cost to minimize when choosing an architecture here.

6.4 Screen A: Encoder Depth and Width

Question

This screen measures how message-passing depth K and encoder width d_h behave after the information-boundary correction. On the busbar graph, depth controls reach: one layer crosses one transmission-line edge, while two layers let each busbar combine information across paths of two physical edges. The graph-level readout then pools all valid node embeddings, so K need not equal the diameter of the local graph. Width controls how much information the node and pooled embeddings can retain.

The result applies directly only to the busbar schema; using the same pair for graphs with explicit asset or line nodes remains an assumption.

Design

A complete 3×4 grid at seed 0, 12 runs. Table 6.3 lists the screened parameters and their levels. The graph output dimension is matched to the hidden dimension, so a cell varies one width rather than two. Every other setting is inherited from the corrected-boundary reference configuration of Section 6.3.

Table 6.3: Screen A. Screened parameters and the values tested. All other settings are held at Table 6.2.

Parameter	Setting	Values tested
Message-passing layers K	<code>gnn_layers</code>	1, 2, 3
Encoder width d_h	<code>gnn_hidden_dim</code>	16, 32, 64, 128

Result

The complete endpoint and marginal tables, the training curves, and the parameter-count diagnostic are reported in Appendix A.2.




Figure 6.3: Screen A. The same twelve cells as a heatmap of the full-test survival of each run’s best checkpoint. All cells use corrected graph inputs and seed 0.

Reading

The endpoint leader is mp3_h32. Its best checkpoint reaches 97.80% full-test survival, 3.08 points above mp2_h128 at 94.72%.

Depth does not separate cleanly. The row means are 83.80%, 82.90%, and 84.74% for $K = 1, 2, 3$. Their ranges overlap heavily, and the best and second-worst cells both occur at $K = 3$. These four single-seed cells per row do not establish a reliable depth ordering.

Width is strongly non-monotonic. Averaged across depths, $d_h = 32$ leads at 93.97%, followed by 128 at 91.39% and 16 at 88.42%; every $d_h = 64$ cell is weak, yielding a 61.49% marginal mean.

The factorial does not yield a simple architecture rule. The depth marginals differ by less than two points, width is non-monotonic, and the ranking changes across depth–width combinations. The grid is therefore better used to select a robust operating point than to claim a monotonic depth or capacity effect.

Each endpoint covers all 201 test chronics, but it still reflects a checkpoint selected from roughly 180 ten-episode evaluations and only one training seed. Table 6.4 therefore complements it with metrics computed from the complete training curves.

Table 6.4: Screen A. Convergence measured from the training curves rather than from a single checkpoint. Steps to 90% is the first evaluation at which the smoothed test survival reaches 90%; the curve mean is the time-average of the smoothed curve over training. The final-window mean averages the raw last 20 periodic evaluations, spanning approximately the final 1.58M environment steps. This is the campaign’s late-stability window: it reduces sensitivity to the last few evaluations while remaining local to the converged regime. Dashes mark architectures that never reached the threshold.

Architecture	Steps to 90% (M)	Curve mean (%)	Final-window mean (%)
mp2_h128	6.22	71.40	94.04
mp1_h16	6.88	70.92	84.93
mp3_h32	7.88	69.63	91.28
mp2_h16	8.79	69.16	88.21
mp1_h128	11.03	68.50	91.11
mp1_h32	6.64	63.13	93.77
mp3_h128	—	58.58	80.20
mp1_h64	—	51.60	63.37
mp3_h64	—	49.80	67.58
mp3_h16	—	48.82	60.00
mp2_h32	13.69	46.51	91.08
mp2_h64	—	45.00	53.41

The curve summaries separate a best-checkpoint endpoint from sustained late training behavior. **mp2_h128** leads all three curve criteria: it has the earliest 90% crossing, the highest time-averaged curve, and the highest 20-evaluation final-window mean. Over that same final window its standard deviation is 5.29 points, compared with 10.45 for **mp3_h32**. The latter still has the best full-test checkpoint, but its 3.08-point endpoint advantage comes from one seed and does not persist across the late trajectory.

Why carry mp2_h128 forward. The selection combines the training evidence with an architectural prior rather than treating the noisy factor marginals as a law. In the busbar schema, $K = 2$ is the smallest depth that represents interactions across two consecutive transmission-line edges; the global readout already combines all node embeddings, while a third layer adds another round of mixing. This does *not* mean that every pair of nodes is at most two hops apart. Conditional on this deliberately chosen depth, $d_h = 128$ is the strongest width in the $K = 2$ row both at full test and over the late window, and its additional parameter cost is negligible beside Grid2Op simulation. We therefore select $(K, d_h) = (2, 128)$ as the conservative default for the subsequent screens, while retaining **mp3_h32** as the seed-0 endpoint challenger that would need a multi-seed confirmation.

6.5 Screen B: Graph Readout Scope and Reducer

Question

The corrected local graph may contain nodes that provide message-passing context without being equally useful in the final graph summary. Should readout pool every visible node, only controlled nodes, only energized nodes, or the intersection of controlled and energized nodes? For each scope, should the valid embeddings be averaged or reduced by a per-channel maximum?

The scope mask is applied only at graph readout. Context nodes remain in the message-passing graph, so this screen changes which embeddings form the policy summary without changing what information can propagate to a controlled node.

Design

A 4×2 factorial crosses four readout scopes—all visible, controlled, energized, and controlled $\cap$ energized—with mean and max reduction at seed 0. All cells use GINE with $K = 2$ and $d_h = 128$ on the corrected busbar graph. Seven cells are new; the all-visible/mean cell reuses **mp2_h128** from Screen A because it is

exactly the same configuration.

Result

Table 6.5: Screen B. Full-test survival by readout scope and reducer at seed 0. The all-visible/mean cell is inherited from Screen A.

Readout scope	Reducer	
	Mean	Max
All visible	94.72 [†]	96.31
Controlled	94.60	94.79
Energized	98.93	99.39
Controlled $\cap$ energized	91.45	96.47

[†]Screen-A mp2_h128 reference; not a new run.

Table 6.6: Screen B training-curve diagnostics at seed 0, ranked by curve mean. “To 90%” is the first crossing of 90% by the five-evaluation trailing mean; curve mean is its time-weighted mean over training. Final-20 statistics use the raw final 20 periodic evaluations (approximately 1.58M environment steps).

Scope/reducer	To 90% (M)	Curve mean (%)	Final-20 mean (%)	Final-20 std (pp)
Energized/max	3.40	79.26	97.79	4.12
Energized/mean	8.71	73.59	98.01	2.90
Controlled $\cap$ energized/mean	5.97	72.97	91.18	8.59
All visible/mean [†]	6.22	71.40	94.04	5.29
Controlled/mean	6.47	70.94	91.21	6.08
Controlled $\cap$ energized/max	8.88	68.65	86.31	8.80
All visible/max	4.73	68.38	88.57	8.08
Controlled/max	10.87	68.27	90.67	6.81

The corresponding training curves are reported in Appendix A.3.

Reading

Energized-node readout leads with both reducers. Its mean-pooling cell reaches 98.93% and its max-pooling cell reaches 99.39%, only 0.46 points apart. Relative to pooling all visible nodes, energized-only pooling improves survival by 4.21 points under mean and 3.08 points under max.

Max is directionally consistent across scopes. For every matched scope, max pooling exceeds mean pooling on the full test. The gain is 1.59 points for all visible nodes, 0.19 for controlled nodes, 0.46 for energized nodes, and 5.01 for the controlled-and-energized intersection. The magnitude therefore depends strongly on scope, but the direction does not.

Why carry energized/max forward. It combines the best scope with the reducer that wins every within-scope full-test comparison. Table 6.6 also identifies it as the strongest overall training curve: it is the earliest to cross 90% (3.40M steps) and has the highest time-weighted curve mean (79.26%). Energized/mean has the better final-20 mean and standard deviation, so the late-window evidence alone does not select max, and this remains a seed-0 screening choice rather than a general claim that max pooling must always dominate. Taken together with the 99.39% full-test endpoint, **energized_max** is the most consistently supported operating point and is fixed for the corrected reruns of Screens C and D, preventing readout from becoming an additional factor in those experiments.

6.6 Screen C: Structural Augmentations of the Busbar Graph

Question

Three optional augmentations add computational structure to the busbar graph without adding any physical measurement: same-substation edges (e), which let the busbars of one substation exchange messages directly; substation-summary nodes (n), which insert one learned hub per substation; and virtual-node readout (v), which replaces mean pooling with a learned graph-summary node. Each is motivated by the same intuition, shorten the path between related busbars, so the question is whether any of them earns its cost, and whether they compose.

Design

A complete 2^3 factorial with three seeds per cell, 24 runs. The graph type is busbar throughout and the baseline cell is the unaugmented $e0n0v0$. It has the same nominal $K = 2$, width-128, GINE/mean design as the corrected reference, but predates the boundary fix and is not numerically comparable to Screens A and B. Cells are named by their three switches, so $e1n0v1$ means same-substation edges on, summary nodes off, virtual-node readout on.

Result

The detailed result tables, training curves, and per-seed heatmap are reported in Appendix A.4.




Figure 6.4: Screen C. The factorial as a cube: each vertex is one of the eight cells and each edge flips a single factor with the other two held fixed. The baseline $e0n0v0$ is circled in red.

Reading

No augmentation improves the result. The unaugmented `e0n0v0` baseline is the best cell at 97.31%, and all twelve single-factor flips in Table A.10 are negative. It is also the most stable cell, with a seed standard deviation of 1.22 points against 3.9–33.9 for the augmented cells. Figure A.5 shows that the cells separate well before the training budget ends.

The hierarchy factors have the largest effects: -15.13 points for summary nodes and -14.11 for virtual-node readout. Each is relatively mild alone but severe when combined, with both `n1v1` cells between 66% and 69%. Their four conditional flips span about 23 points, so the main effects hide a strong interaction visible on the upper face of Figure 6.4.

The likely explanation is redundancy: both mechanisms add a learned aggregation node, and together they stack two hubs between the busbars and the readout (Figure 5.8). With only two message-passing layers, most physical nodes cannot reach that readout; Screen A’s $K = 3$ row directly tests this interpretation.

Same-substation edges are milder and more consistent: their mean effect is -3.28 points and the four flips range from -1.11 to -5.71 . The best edge cell is only 1.11 points below the baseline, within the observed seed variability (Figure A.6); they therefore show no reliable benefit and add candidate edges without measured return.

The subsequent stages therefore keep the plain busbar graph with mean readout. Any future hierarchy should be tested with a deeper encoder.

6.7 Screen D: Generator and Load Message Direction

Question

In the disaggregated graph, each generator and load is a node attached to its busbar. Section 5.2.2 leaves the direction of that attachment open: the asset may send messages to the busbar, receive them from it, or both. The two directions are not symmetric in effect, because the mean readout of Section 5.4.1 pools asset nodes as well as busbars. Does the choice matter, and do the generator and load relations behave independently?

Design

A complete 3×3 factorial over the generator and load directions with three seeds per cell: 27 runs. The graph type is disaggregated throughout; the line and summary directions stay **bidirectional** and are inert for this schema. The encoder matches the reference configuration: GINE, two layers, 128 features, mean readout, no augmentations, so the only screened factors are the two attachment directions. The baseline cell is bidirectional on both relations. Directions are abbreviated **bi** (bidirectional), **a2b** (asset to busbar), and **b2a** (busbar to asset).

Result

The training curves and complete endpoint ranking are reported in Appendix A.5.

Generator × load direction — cell mean with the three seed values




Figure 6.5: Screen D. Full-test survival of the best checkpoint by generator direction (rows) and load direction (columns), with the three individual seed values printed under each cell mean. The baseline cell is bi/bi.

Reading

Busbar-to-asset is the better direction for both relations. Averaging over the other relation, the generator marginal is 93.43% for b2a, 90.58% for bi, and 78.01% for a2b; the load marginal follows the same order at 93.39%, 86.91%, and 81.73%. The strongest part of the result is not the mean but the spread: the two best cells, b2a/b2a and b2a/a2b, have seed standard deviations of 2.20 and 1.42 percentage points against 8.24 for the bidirectional baseline, and their worst seeds (95.40% and 96.32%) are above the baseline’s mean. The seed values printed in Figure 6.5 show this at a glance: the bottom row holds three tightly grouped seeds per cell, while the baseline cell spans 83.2% to 98.4%.

The failure case is equally clear. Making both relations asset-to-busbar collapses the policy to 56.41% with a 22.12-point spread and a worst seed at 30.88%, which is 36 points below the baseline and the largest single effect measured anywhere in this chapter. Its panel in Figure A.7 separates from the baseline within the first two million steps and never recovers, so this is a training failure rather than a late collapse; one of its seeds never improved on a checkpoint from 4.31M steps.

A plausible mechanism, consistent with the readout defined in Section 5.4.1 but not directly measured here, is that the mean pooling includes generator and load nodes. Under b2a an asset node receives busbar context and its updated embedding enters the readout, so the asset row carries information about its local neighborhood. Under a2b the asset node never receives a message, so after K layers its embedding is still a projection of its own raw features, and it dilutes the pooled vector with context-free rows; with both relations set that way, every generator and load row is inert.

6.8 Screen E: Candidate-Action Scoring on the Corrected Boundary

Question. The candidate-action head of Section 5.5 scores every action from the encoded nodes it modifies rather than from a fixed-width layer over an action index. Three choices in that head are unresolved: how the touched nodes are pooled, whether the static action descriptor is appended to the score, and whether the idle action gets its own head. This screen varies all three.

Design. A 2^3 factorial on `bus14`, seed 0, eight runs. Pooling is `mean` or `typed_mean`; the action descriptor is appended (`f1`) or withheld (`f0`); the idle action is scored by a dedicated head (`a0h1`) or by the shared scorer (`a0h0`). Every run uses the disaggregated line-node graph without one-hop context, a two-layer GINE encoder of width 128, and the 15M-step budget of Section 6.3. All eight reached 15M steps.

Result. Table 6.7 reports the full-test survival of each best checkpoint over the 201 held-out chronics. The corresponding training curves and endpoint ranking are reported in Appendix A.6.

Table 6.7: Screen E, full-test survival of the best checkpoint on 201 held-out chronics. One seed per cell.

Pooling	Descriptor	Idle head	Survival (%)	Best checkpoint
<code>mean</code>	<code>f0</code>	<code>a0h1</code>	99.52	14,929,920
<code>typed_mean</code>	<code>f1</code>	<code>a0h0</code>	99.50	14,929,920
<code>mean</code>	<code>f0</code>	<code>a0h0</code>	99.18	14,929,920
<code>mean</code>	<code>f1</code>	<code>a0h1</code>	98.66	14,929,920
<code>mean</code>	<code>f1</code>	<code>a0h0</code>	98.32	14,846,976
<code>typed_mean</code>	<code>f0</code>	<code>a0h1</code>	98.17	14,764,032
<code>typed_mean</code>	<code>f0</code>	<code>a0h0</code>	92.50	12,358,656
<code>typed_mean</code>	<code>f1</code>	<code>a0h1</code>	57.84	11,612,160

Reading. Six of the eight cells reach 98.2% or better and three exceed 99%, so the candidate-action head is a viable actor on the corrected boundary. No factor separates. The apparent marginals, `mean` at 98.9% against `typed_mean` at 87.0% and `f0` at 97.3% against `f1` at 88.6%, are artifacts of two low cells; removing the single failure at 57.84% brings the pooling gap to about a point. At one seed the defensible statement is that the head works and that the three choices do not order themselves.

The descriptor is the one asymmetry worth recording: the best cell withholds it and both failing cells include it, against a prior that favours the more elaborate option. It is not redundant information, since every action still reaches the scorer with its own pooled context – agent 0’s 57 actions produce 57 distinct row-sets, because an action toggles one to three of a substation’s lines rather than a single busbar. Withholding it removes a static label and leaves the learned description intact.

6.8.1 Replication with Corrected Graph Inputs

At seed 0, the eight cells were repeated with corrected graph inputs (NLS): deterministic physical scaling and line-angle differences, without running standardization. The GNN remains shared, but each agent retains its own scoring MLP and optional idle head; this tests preprocessing, not scorer sharing.

Table 6.8: Full-test survival over 201 held-out bus14 chronics with original (NL) and corrected (NLS) graph inputs. The final column is NLS minus NL in percentage points. One seed per cell.

Pooling	Descriptor	Idle head	NL (%)	NLS (%)	Δ
mean	f1	a0h0	98.32	100.00	+1.68
typed_mean	f0	a0h0	92.50	99.84	+7.33
mean	f0	a0h1	99.52	99.46	-0.07
mean	f0	a0h0	99.18	99.43	+0.25
typed_mean	f1	a0h1	57.84	90.71	+32.87
mean	f1	a0h1	98.66	83.81	-14.85
typed_mean	f1	a0h0	99.50	71.67	-27.82
typed_mean	f0	a0h1	98.17	69.67	-28.50




Figure 6.6: The paired comparison of Table 6.8. NLS changes both power scaling and angle representation; the differences cannot be attributed to scaling alone.

Reading. Four NLS cells exceed 98%, but three remain below 91%, with changes above 25 points in either direction. At one seed, this indicates an interaction with the candidate head rather than a stable preprocessing effect. Only `mean/f0` is robust: both idle-head choices exceed 99.1% with NL and 99.4% with NLS. Chapter 8 separately studies transfer with a scorer shared across agents.

Chapter 7

What the Shared Encoder Reads

One encoder serves every agent, and the same encoder is what a transfer study moves between grids. Its inputs are raw physical quantities, so before either form of sharing can be interpreted the inputs themselves have to be measured. This chapter does that: it establishes what the encoder consumes, how far those distributions differ between two grids, which of the differences a network absorbs and which it does not, and what corrections follow.

Chapter 8 then applies the corrections and runs the transfer. The measurements here are the reason that study is interpretable at all: without them, a deficit under transfer could not be told apart from a representation reading miscalibrated inputs.

7.1 Why the input distribution matters

The graph encoder is grid-size independent and can therefore be transferred across grids. A conventional action head cannot: it has one output per action index, so its shape changes with the agent partitions and action sets. The candidate scorer of Section 5.5 escapes that by applying shared weights and action-conditioned pooling to a variable set of actions, which makes the action head transferable too. Only the flat centralized critic stays grid-specific, its input dimension being set by the grid.

The two input paths must be distinguished. The critic receives the normalized flat observation; the transferred actor encoder receives only graph node and edge features, for which physical scaling and running standardization are disabled here. A frozen `bus14` encoder therefore reads WCCI's raw graph quantities, and whether that matters depends on how closely the two distributions match.

7.2 Measurement protocol

Both environments were rolled out under a do-nothing policy for 6,000 environment steps, with a cap of 300 steps per chronic so the sample spreads over roughly twenty scenarios per grid rather than concentrating in one. At each step the per-agent local graphs were rebuilt and every visible node recorded, pooling agents together because a single encoder is shared across them. This yields 426,000 node observations for `bus14` and 1,236,000 for WCCI.

Only nodes are measured, because on this schema only nodes carry measurements: the nine edge channels are relation indicators, exactly one of which is set on any active edge, with no continuous content to compare. The transfer runs also set `gnn_include_neighbors = false`, so every node in a local graph is controlled and `domain_mask` is constant at one throughout.

Four statistics summarize each feature. Writing F_A and F_B for the two empirical distributions with

means μ_A, μ_B and standard deviations σ_A, σ_B , and $\sigma_p = \sqrt{(\sigma_A^2 + \sigma_B^2)/2}$ for their pooled spread:

$$\text{offset} = \frac{\mu_B - \mu_A}{\sigma_p}, \quad \text{scale ratio} = \frac{\sigma_B}{\sigma_A}, \quad (7.1)$$

$$\text{KS} = \sup_x |F_A(x) - F_B(x)|, \quad W_1/\sigma_p = \frac{1}{\sigma_p} \int_0^1 |F_A^{-1}(q) - F_B^{-1}(q)| dq. \quad (7.2)$$

Reading the equations. A is always **bus14**, the source grid, and B always **WCCI**, the target: an offset is therefore positive when **WCCI** sits higher, and a scale ratio above one means **WCCI** is the wider of the two. F_A and F_B are the two *empirical cumulative distribution functions*, one per grid, built from the samples of Section 7.2: $F_A(x)$ is the fraction of **bus14**’s sampled values that fall at or below x , so it rises from 0 to 1 as x sweeps the axis. Its inverse $F_A^{-1}(q)$ runs the other way, returning the value below which a fraction q of the sample lies — $F_A^{-1}(0.5)$ is the median. The two definitions are what let the last two statistics compare the same pair of samples from perpendicular directions: KS measures the gap between the curves vertically at fixed x , W_1 between their inverses horizontally at fixed q . Their names: KS is the Kolmogorov–Smirnov statistic, the largest vertical gap between the two empirical cumulative distribution functions. It lies in $[0, 1]$, reaching 0 when the distributions coincide and 1 when their supports are disjoint, and the last panel of Figure 7.1 draws the gap it measures. W_1 is the 1-Wasserstein distance, also called the earth mover’s distance: the average horizontal gap between the two quantile functions, or equivalently the least mass $\times$ distance needed to rearrange one distribution into the other. The subscript is the exponent applied to that displacement — W_2 would square it and so weigh far excursions more heavily — and this work uses $p = 1$ throughout. Because W_1 carries the units of the feature, it is reported in pooled standard deviations so that channels measured in MW and in degrees stay comparable.

The pairing is deliberate: the offset is blind to a change in spread and the scale ratio is blind to a displacement, so a channel is only unproblematic when both sit near their neutral values. KS measures disagreement in the bulk and is comparatively insensitive to tails, which W_1 does weigh. Section 7.5 shows why the two failures need separating.

7.3 What the encoder reads

The measurements are taken on the disaggregated line-node graph, the schema the transfer study uses. Its twelve node channels are specified in Table 5.6: four type indicators, the shared **p** and **theta** columns, four line-state columns, the substation cooldown, and **domain_mask**.

Every node type shares one feature vector, so a channel is filled only on the type that owns it and is zero-filled elsewhere. That makes the sparsity structural and severe: **p** and **theta** carry a value only on generator and load nodes, which are 24.0 % of rows on **bus14** and 28.7 % on **WCCI**, while **rho** and the other line-state channels are filled only on transmission-line nodes, 36.6 % and 36.4 % of rows. A statistic taken over all rows therefore mixes a measurement with the zeros of three other node types, which is why every channel below is reported twice.

Two channels are constant in this sample rather than sparse. The two cooldown columns are exactly zero only because the policy is do-nothing; under control a topology change sets the substation cooldown to 3 and a reconnection sets the line cooldown to 3. **timestep_overflow** is zero on all but a handful of rows. One channel is dimensionless by construction: **rho** divides each line’s flow by that line’s own thermal limit, while every other physical channel carries the units of the grid it came from.

7.4 Results

Table 7.1 reports the shift statistics twice: over all sampled rows, and over only the node type that carries the channel. The two readings differ because a shared column is a point mass at zero, contributed by the three types that do not own it, mixed with the measurement of the type that does.

Table 7.1: Distribution shift between the two grids on the disaggregated line-node graph, ordered by KS. Offset is in pooled standard deviations and the scale ratio is WCCI over **bus14**. The right-hand block restricts each measured channel to the node type that owns it: generator and load nodes for **p** and **theta**, transmission-line nodes for the line state. **The type indicators and domain_mask own no type, so the restriction does not apply; constant marks a channel with no variation in this sample, whose ratio is undefined.**

Channel	Zeros (%)		All sampled rows			Owning node type		
	bus14	WCCI	offset	ratio	KS	offset	ratio	KS
theta	77.5	73.8	+0.653	0.83	0.157	+1.864	1.43	0.678
p	77.8	76.7	-0.037	2.14	0.042	-0.158	2.54	0.201
rho	63.4	63.6	-0.088	0.92	0.053	-0.260	1.05	0.141
line_status	63.4	63.6	-0.005	1.00	0.002	constant at one		
timestep_overflow	100	100	-0.002	0.70	0.000	-0.004	0.70	0.000
time_before_cooldown_line	100	100	—	—	0.000	constant		
time_before_cooldown_sub	100	100	—	—	0.000	constant		
is_busbar	60.6	65.0	-0.093	0.98	0.045	—		
is_generator	91.5	89.3	+0.076	1.11	0.022	—		
is_load	84.5	82.0	+0.066	1.06	0.025	—		
is_transmission_line	63.4	63.6	-0.004	1.00	0.002	—		
domain_mask	0	0	—	—	0.000	constant at one		

Read on all rows the table looks harmless: no channel reaches $KS = 0.2$. That is the type mixing, not agreement. Restricted to the owning type, **theta** reaches 0.678 with +1.86 standard deviations of displacement, and **p** widens by a factor of 2.54. The four type indicators differ only slightly, which says the two grids have a similar mix of busbars, assets, and lines per local graph rather than saying anything about their physics.

The extremes make the same point concretely. **p** spans $[0, 97]$ MW on **bus14** and $[-281, 399]$ MW on WCCI, where 7.8% of asset rows are negative and **bus14** has none. **theta** spans $[-11.6, 1.5]$ degrees against $[-7.1, 16.8]$: not a wider version of the same distribution but a displaced one.




Figure 7.1: The measured channels on the node type that owns them. `rho` lands almost on top of itself, as expected from a quantity already normalized by each grid’s own thermal limits. `p` shares a centre but not a width, and WCCI alone produces negative values. `theta` is displaced rather than rescaled: `bus14` sits almost entirely on the negative side. The last panel reads that displacement as the KS statistic. Full ranges on log density are in Figure A.10.

7.5 Two failure modes

The left panel of Figure 7.2 separates the channels along the two axes; the right panel repeats the diagnostic after the corrections and is read in Section 7.6.5. The distinction is not cosmetic, because the architecture responds to the two axes differently.

Reading the axes. Each point is one channel, placed by the two statistics of Equation (7.1) taken over all sampled entries. The horizontal axis is the *magnitude* of the offset, $|\mu_B - \mu_A|/\sigma_p$, in pooled standard deviations: how far the two distributions have moved apart, with the sign discarded because only the size of a displacement matters here. The vertical axis is the scale ratio σ_B/σ_A , WCCI over `bus14`: how much wider or narrower WCCI’s spread is, with 1 meaning identical spread. It is drawn on a logarithmic scale so that halving and doubling sit the same distance from 1, which a linear axis would misrepresent. A channel needing no correction sits at the origin of this reading — near-zero offset and ratio near 1 — so distance from the bottom-left corner is distance from comparability, and the direction of that distance names which of the two failures it is.




Figure 7.2: Absolute offset against scale ratio, over all sampled entries. The shaded band and the dashed line mark a scale ratio within a third of unity; the dotted vertical line marks an offset of 0.3 pooled standard deviations. *Left*: as measured. Channels outside those bounds fail in one of two distinct ways — the power channels keep their centre and widen, the angles keep their width and move. *Right*: the same diagnostic after the two corrections of Section 7.6, discussed in Section 7.6.5. Points are placed on the owning node type. Every channel now sits inside the scale band and the worst displacement falls from 1.86 to 0.81, but *p* acquires an offset it did not have.

The actor’s node pre-encoder is a linear map followed by ReLU and LayerNorm. LayerNorm absorbs much of a uniform change of *scale*: multiplying every input channel by a constant leaves the normalized activations largely unchanged. It does not absorb an *offset*. Displacing a channel by a constant moves the pre-activations to a different region of the ReLU, changing which units fire at all, and normalization applied afterwards cannot undo that.

This makes the angle, not the power, the binding problem. *p* differs in width by a factor of 2.54, which the architecture partially self-corrects, though the negative WCCI region reaching -281 MW is genuinely unseen territory rather than a wider version of something familiar. *theta* differs by 1.86 pooled standard deviations of displacement, which it does not correct at all.

rho, meanwhile, needs no correction at all: scale ratio 1.050, offset -0.26 — the same values the busbar schema gives, since it is the same set of lines carrying the same loadings. The single most decision-relevant channel is already comparable across grids because it was defined in per-unit terms, an argument for expressing features in grid-relative units wherever the physics permits it.

7.6 Correcting the inputs

Two preprocessing mechanisms exist in the codebase. Deterministic physical scaling divides each channel by a per-environment physical constant. Running standardization maintains per-environment, per-node-type statistics over *p* and *theta*, excluding *rho*, which the measurements above confirm needs no help. On the diagnosis alone the second looks like the match, being the only one that can move a displaced channel. It is not.

7.6.1 Running standardization makes the mismatch worse

Standardizing each environment with its own statistics aligns the aggregate moments exactly: every channel comes out at offset 0.00 and scale ratio 1.00. It also moves both channels it touches an order of magnitude further apart (Table 7.2).

Table 7.2: KS between the two grids before and after per-environment standardization, over all sampled rows. The last column is the number a node that does *not* own the channel presents to the encoder once standardized. Raw, such a row reads 0 on both grids; standardizing sends it to $-\mu/\sigma$, a property of the grid. ρ is shown for contrast and is excluded from the running-statistics feature set.

Channel	KS raw	KS standardized	“Not this type” encoding bus14 vs WCCI
p	0.042	0.773	−0.371 vs −0.145
theta	0.157	0.673	+0.481 vs −0.143
rho (excluded)	0.053	0.634	−0.658 vs −0.622

The cause is the shared-column structure of Section 7.3. Each channel is a point mass at zero, contributed by the three node types that do not own it, plus the measurement of the type that does. Subtracting a per-grid mean moves that point mass to a grid-dependent location. Since three quarters of all rows are zeros for **p** and **theta**, the encoder’s single most frequent input stops meaning the same thing on the two grids, which is a larger mismatch than the displacement being corrected. On **theta** it is starker still: a row that does not own the channel reads +0.481 on bus14 and −0.143 on WCCI, so the same absence lands on opposite sides of zero.

Two further properties rule the mechanism out independently of these numbers. Standardization is scale-invariant, so it erases any scaling applied before it — dividing by a constant and then standardizing reproduces standardizing alone to four decimals — meaning the two mechanisms cannot be combined. And running statistics begin empty and converge during training, so a transferred encoder would spend early target training reading a distribution that is itself still moving, where a deterministic scaling is fixed at construction and identical on both grids from the first step.

7.6.2 Physical scaling for the power channels

Physical scaling, defined in Section 5.1, divides the power channels by each environment’s own maximum generator capacity. Two properties answer the failure diagnosed above: being multiplicative it leaves zeros at zero, so the point mass stays aligned on both grids, and being fixed at construction it does not drift the way estimated moments do. The divisor is a single generator’s nameplate rating, a crude proxy for grid scale, so it is worth checking against the alternatives (Table 7.3); it is the closest of the four to the divisor that would equalize the two spreads exactly.

Table 7.3: Candidate power normalizers and the p scale ratio each would produce on asset nodes. A ratio of 1.00 would be exact; the raw ratio is 2.54.

Divisor	bus14	WCCI	Resulting ratio
Exact	—	—	1.00
Maximum $P_{\max}$ (used)	140	400	0.89
Total $P_{\max}$	540	2,450	0.56
Generator count	6	22	0.69
Median $P_{\max}$	85	71	3.04

7.6.3 The angle displacement is a gauge artifact

Scaling cannot repair the angle displacement, and the implementation does not attempt it: angle channels are divided by a constant 180, identical on both grids, so both sides move together and the displacement is untouched. The measurement explains why. An absolute voltage angle is referenced to whichever bus is slack, and the two grids do not place theirs alike, so the displacement in Figure 7.1 — bus14 spanning $[-11.6, 1.5]$ degrees against WCCI’s $[-7.1, 16.8]$ — is largely gauge rather than physics. Nothing that rescales or recentres a reference-dependent quantity can fix that; the quantity has to be replaced.

The replacement is the angle drop along each line, defined in Section 5.1.2. Being a difference it

is gauge-invariant, and on this schema it lives on the transmission-line node, which is the most populous type in the graph, so it also trades a channel owned by a quarter of the rows for one owned by a third. This study therefore enables it, alongside physical scaling, in both source and target environments.

7.6.4 The action descriptor must not be normalized per agent

The same error appears once more, on an input this chapter has not otherwise considered, and with the agents rather than the grids as the populations. One input reaches the actor without passing through the encoder: the candidate-action head of Section 5.5 appends a static descriptor of each action to the pooled node context. Its count columns were divided by the largest value in that agent’s *own* action set — the construction Section 7.6.1 rejects for node features, applied to a different input, with the agents as the populations.

Table 7.4: Divisors applied to the action count descriptors under the previous per-agent rule, `bus14`.

Agent	Actions	Topology changes	Line endpoints
0	57	5	3
1	51	5	4
2	83	6	3

An action moving three elements therefore reached the shared scorer as 0.60 for agent 0 and 0.50 for agent 2, and every agent’s busiest action reached it as exactly 1.0 whether it moved five elements or six. On another grid the divisors change again. The columns are now divided by a constant, selected by `candidate_action_feature_scaling`; the busiest actions then read 1.25, 1.25 and 1.5. This is the only correction here that applies within a grid as well as across grids, and it qualifies Screen E of Chapter 6, whose runs predate it: their `f1` cells appended a descriptor inconsistent between the agents sharing one scorer.

7.6.5 The corrected distribution

Table 7.5 measures both changes together, under the same protocol as Section 7.2, and the right panel of Figure 7.2 places the result back on the offset–scale diagnostic. The angle channel lands in the band occupied by `rho`, and the power spread comes within 11 % of parity.

Table 7.5: Distribution shift with `gnn_angle_representation = edge_diff` and `gnn_physical_scaling = true`, against the uncorrected baseline, on the node type owning each channel. The offset column is included because scaling moves it: dividing each grid by its own constant is not offset-preserving when the two constants stand in a different relation to typical operation.

Channel	Uncorrected			Corrected		
	offset	ratio	KS	offset	ratio	KS
<code>theta</code>	+1.864	1.43	0.678	replaced		
<code>theta_diff</code>	—	—	—	−0.383	1.15	0.230
<code>p</code>	−0.158	2.54	0.201	−0.808	0.89	0.511
<code>rho</code>	−0.260	1.05	0.141	−0.260	1.05	0.141

The angle correction is unambiguous: replacing a reference-dependent quantity by a gauge-invariant one takes the worst channel in the chapter from 0.678 to 0.230, and its scale ratio stays near one throughout.

The power correction is a genuine exchange rather than a repair. Dividing by each grid’s largest generator rating brings the spread from 2.54 to 0.89, but the raw means, nearly equal at −0.158, are not preserved: `bus14` runs its generators at 39 % of its fleet maximum against WCCI’s 20 %, so a divisor taken from the largest unit lands the two grids at different normalized levels and the displacement grows to −0.808. KS rises with it, from 0.201 to 0.511. Section 7.5 argued that offsets, not scale errors, are what the architecture cannot absorb, so this moves `p` onto the axis that binds.

The exchange is still favourable on the evidence available: the worst displacement across all channels falls from 1.864 to 0.808, every scale ratio ends inside the band, and no channel remains outside both bounds at once. But it is an exchange, not an elimination. A divisor reflecting typical rather than peak generation would plausibly do better on both axes; Table 7.3 tested four candidates against the scale ratio alone, and selecting one against the offset as well is left open. The transfer results of Chapter 8 are the test of whether the residual displacement matters in practice, and the design of Section 8.3 runs both input conditions precisely so that question can be answered.




Figure 7.3: Corrected distributions, for comparison with Figure 7.1. The angle drop is centred on both grids, and the scaled power channels overlap over most of their range instead of differing by a factor of 2.5 in width.

Chapter 8

Transfer to the 36-Substation WCCI Grid

The controlled method comparisons so far were developed on `bus14`. That grid is small enough to iterate on but too small to claim anything about scale: fourteen substations, twenty lines, and three agents. The preliminary WCCI stress test of Section 4.7 reproduced several sparse-control mechanisms with an MLP actor on the larger grid, but even the matched baseline was weak. Because that pilot changed grid size, action space, and maintenance at once, it could not identify the source of failure.

This chapter turns that pilot into a controlled transfer study. It moves to the L2RPN WCCI 2020 grid, more than twice as many substations and three times as many lines, and asks the question that motivates a graph encoder in the first place. If the actor reads the grid as a graph rather than as a flat vector, its weights are independent of how many substations exist, and a representation learned on `bus14` should carry over.

Before interpreting that transfer, the two environments have to be made comparable and the encoder’s inputs have to be checked. This chapter does both. The comparison turns out to matter more than expected: the graph features the encoder consumes are raw physical quantities, and they do not follow the same distribution on the two grids.

8.1 The target environment

Table 8.1 contrasts the source and target grids. The agent partitions follow the same principle in both cases, contiguous electrical areas, but WCCI is split into four regions rather than three, and the regions are markedly unequal: one agent controls seventeen substations while another controls five.

Table 8.1: Source and target environments. Both use topology actions, decentralized agents and an 80/20 train/test chronic split with split seed 0.

	<code>bus14</code>	<code>bus36_wcci_nomaint</code>
Grid2Op environment	<code>12rpn_case14_sandbox</code>	<code>12rpn_wcci_2020</code>
Substations	14	36
Lines	20	59
Agents	3	4
Chronics	1,004	2,880
Episode length	up to 8,064 steps	up to 8,062 steps
Maintenance	none	disabled (Section 8.1.1)

8.1.1 Removing maintenance

Of the environments Grid2Op distributes, `12rpn_case14_sandbox` is the only one above the toy scale that ships without scheduled maintenance. Every larger environment, WCCI 2020 included, injects planned line outages. That is a problem for a transfer study: a policy trained without maintenance and evaluated with it faces a hazard it has never seen, and the two observation spaces differ as well, since the maintenance countdown features are only present when maintenance exists.

The two settings are therefore aligned by removing maintenance from WCCI rather than inventing a schedule for `bus14`, for which Grid2Op provides no specification.

The effect is measured on the complete held-out test split using the same 576 chronics and a do-nothing policy in both settings.

Table 8.2: Do-nothing performance on the full WCCI test split. The chronics and policy are identical; only scheduled maintenance differs.

Setting	Mean survival	Median survival	Complete episodes
Maintenance enabled	27.3%	12.1%	39/576 (6.8%)
Maintenance disabled	57.3%	100.0%	303/576 (52.6%)
Difference	+30.0 pp	+87.9 pp	+264 (+45.8 pp)

Removing maintenance increases mean survival by 30.0 percentage points and episode completion by 45.8 points, confirming that maintenance substantially changes the baseline difficulty of the target grid.

8.1.2 Action space

The unreduced WCCI action space is not trainable. Enumerating topology actions over each agent’s region gives the sizes in the first column of Table 8.3: nearly 67,000 actions in total, of which agent 1 alone accounts for 65,642. All experiments therefore use the reduced action space obtained by the brute-force outcome procedure, at $k = 256$ per agent.

Table 8.3: Per-agent action-space sizes on `bus36_wcci_nomaint`. Agents 0 and 2 are already exhaustive at $k = 256$, so the reduction only binds on agents 1 and 3.

Agent	Substations	Full	Reduced ($k = 256$)
<code>agent_0</code>	5	77	77
<code>agent_1</code>	5	65,642	256
<code>agent_2</code>	9	127	127
<code>agent_3</code>	17	1,119	256
Total	36	66,965	716

The choice of k was made by evaluating a one-step simulator-greedy policy over each candidate reduction on the same fifty chronics (Table 8.4). Performance saturates at $k = 256$: the difference between $k = 256$ and $k = 1024$ is four chronics won against three lost, which is noise, while both clearly beat $k \leq 128$. Since $k = 256$ is the smallest reduction in the saturated group, it gives the smallest action heads for no measurable loss in what the action set can achieve.

Table 8.4: One-step simulator-greedy policy over each reduced action space, on the same fifty held-out chronics of `bus36_wcci_nomaint`. The do-nothing row is the common reference; no reduction was evaluated against a different chronic set.

Policy	Mean survival	Full-episode survival	Episodes worse than do-nothing
Do nothing	0.560	52%	—
Greedy, $k = 32$	0.679	60%	5
Greedy, $k = 64$	0.819	64%	3
Greedy, $k = 128$	0.847	76%	3
Greedy, $k = 256$	0.976	92%	0
Greedy, $k = 512$	0.883	82%	2
Greedy, $k = 1024$	0.953	90%	0

These two rows bracket the transfer experiments. A learned policy that lands near 0.56 mean survival has learned nothing beyond doing nothing; one approaching 0.9 is doing real work. The greedy row is not a strict upper bound, since it queries the simulator at every decision step and is myopic, but it is a strong reference.

8.2 Corrected inputs

Chapter 7 measures the encoder’s inputs on both grids and settles two corrections, both applied in every run below. Power channels are divided by each grid’s own maximum generator rating, which leaves the structural zeros in place and is fixed at construction rather than estimated online (Section 7.6.2). The absolute voltage angle is replaced by the drop along each line, which removes a displacement that is a property of the slack reference rather than of the grid state (Section 7.6.3). Running standardization is left off: Section 7.6.1 measures it degrading three of the four channels it touches by a factor of five or more. Together these bring the angle channel into the band occupied by `rho` and the power spreads to within 15 % of parity (Table 7.5).

8.3 Transfer experiment design

The transferred model is the shared candidate-scoring actor of Section 5.5 on the disaggregated line-node graph, the configuration Screen E selected on `bus14`. Four of its settings are varied factorially: the input condition (raw features and absolute angles, or physical scaling with the angle drop), the touched-node pooling (joint or type-separated mean), the static action descriptor (off or on), and the idle action (shared or dedicated scorer). Each of the sixteen cells runs in two arms. In the *fine-tuned* arm the encoder and the candidate scorer are both loaded from the paired `bus14` run and both keep training; in the *scratch* arm the same configuration starts from random initialization. Every other setting is copied from the corresponding `bus14` configuration, which is what makes the weights loadable at all.

Transferring the scorer as well as the encoder is what the candidate head makes possible. Its parameters do not depend on the action count, so one agent’s scorer, `agent_0` here, can initialize every agent on the target grid.

Loading needs care in one place. The encoder registers its edge index and node identifiers as persistent buffers, so they appear in the saved state dictionary at source-grid shape. They describe the source graph and are supplied per agent from the target environment’s own specification at call time, so they are dropped during the load, while any other shape disagreement is treated as an architecture mismatch and raises rather than leaving part of the encoder randomly initialized.

Table 8.5: Transfer study design. The two arms of a pair differ only in the two transfer settings.

Component	Setting
Environment	bus36_wcci_nomaint , difficulty 0, topology actions, 80/20 split with split seed 0
Action space	Reduced at $k = 256$, giving 77 / 256 / 127 / 256 actions per agent
Actor	Shared candidate-scoring head on the disaggregated line-node graph, share_actor_gnn
Factors	Input condition $\times$ touched-node pooling $\times$ descriptor $\times$ idle scorer, $2^4 = 16$ cells
Arms	Fine-tuned encoder and scorer from the paired bus14 run, both trainable; identical configuration from scratch
Seeds	0
Budget	15M steps, matching the bus14 source runs
Inputs	Both conditions are run: raw features with absolute angles, and physical scaling with edge_diff

Running both input conditions rather than only the corrected one is deliberate. Comparing the scratch arms with and without the correction isolates whether it helps or harms on the target grid at all, independently of any transfer effect and on otherwise identical settings; comparing each scratch arm against its fine-tuned partner isolates the transfer.

Appendix A

Supplementary Results

This appendix collects detailed result tables and diagnostic figures that support Chapters 4 and 6. The main chapters retain the compact comparisons needed for the discussion.

A.1 Sparse Intervention Control

The joint survival/action-0 comparison is retained in Figure 4.1; the tables below provide the exact aggregate and per-agent values.

The following tables report average full-test episodic survival and action-0 fractions over the three evaluated seeds. The all-agent value is the arithmetic mean of the three per-agent fractions. The do-nothing row is a reference policy that always selects action 0.

Table A.1: Full-test summary for evaluation-time rho heuristics.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Global rho heuristic, $\rho_{safe} = 0.90$	98.36	0.936	0.933	0.936	0.940
Local rho heuristic, $\rho_{safe} = 0.90$	97.07	0.968	0.992	0.985	0.928

Table A.2: Full-test summary for intervention-gated actors.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Intervention gate, final-action MAP	89.82	0.802	0.826	0.889	0.690
Intervention gate, hierarchical greedy	93.66	0.395	0.066	0.844	0.276

Table A.3: Full-test summary for flat-policy fixed-penalty runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Flat policy, $\lambda_{fixed} = 0.000$	98.23	0.494	0.524	0.463	0.496
Flat policy, $\lambda_{fixed} = 0.001$	96.75	0.486	0.642	0.401	0.416
Flat policy, $\lambda_{fixed} = 0.003$	93.97	0.575	0.691	0.346	0.688
Flat policy, $\lambda_{fixed} = 0.010$	97.44	0.905	0.907	0.954	0.853

Table A.4: Full-test summary for gated-policy fixed-penalty runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Gated policy, $\lambda_{fixed} = 0.000$	78.51	0.787	0.747	0.949	0.665
Gated policy, $\lambda_{fixed} = 0.001$	73.82	0.862	0.956	0.832	0.799
Gated policy, $\lambda_{fixed} = 0.003$	94.59	0.887	0.934	0.929	0.799
Gated policy, $\lambda_{fixed} = 0.010$	85.69	0.974	0.990	0.984	0.947

Table A.5: Full-test summary for adaptive intervention budget runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Flat local-safe AIB, $d = 0.20$	98.31	0.978	0.980	0.989	0.966
Flat local-safe AIB, $d = 0.10$	98.39	0.979	0.992	0.995	0.949
Flat local-safe AIB, $d = 0.35$	96.61	0.934	0.959	0.947	0.896
Gated local-safe AIB, $d = 0.20$, hierarchical greedy	33.87	0.999	1.000	1.000	0.998
Flat non-idle AIB, $d = 0.20$	95.08	0.986	0.997	0.992	0.970

Table A.6: Full-test summary for the teacher-student distillation runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Teacher: MAPPO with local rho heuristic	97.07	0.968	0.992	0.985	0.928
Student BC, balanced non-idle 0.50, weight 5, aux 0.50	92.14	0.939	0.974	0.925	0.919
Student BC, balanced non-idle 0.20, weight 3, aux 0.25	90.76	0.957	0.982	0.948	0.940

A.2 Screen A: Encoder Depth and Width

Table A.7 gives the complete endpoint grid, while Table A.8 summarizes each factor across the other. Figures A.1 and A.2 provide the corresponding training-curve and parameter-count diagnostics.

Table A.7: Screen A. Full-test survival of the best checkpoint, one seed per cell, by message-passing depth and encoder width. Screen B reuses mp2_h128 as its all-visible/mean reference.

Layers K	Encoder width d_h			
	16	32	64	128
1	84.02	92.72	65.52	92.96
2	92.12	91.37	53.38	94.72
3	89.11	97.80	65.57	86.48

Table A.8: Screen A. Marginal effect of each factor, averaged over the other. Four runs per depth, three per width.

Level	Mean (%)	Min (%)	Max (%)
<i>Message-passing depth</i>			
$K = 1$	83.80	65.52	92.96
$K = 2$	82.90	53.38	94.72
$K = 3$	84.74	65.57	97.80
<i>Encoder width</i>			
$d_h = 16$	88.42	84.02	92.12
$d_h = 32$	93.97	91.37	97.80
$d_h = 64$	61.49	53.38	65.57
$d_h = 128$	91.39	86.48	94.72




Figure A.1: Screen A. Test episodic survival during training, one panel per message-passing depth, with the four widths compared inside each panel. Restarted runs use the canonical history in which the continuation replaces the overlapping post-checkpoint branch.




Figure A.2: Screen A. Full-test survival against actor parameter count on a logarithmic axis, joined within each depth. The grid spans 83.6k to 244.9k actor parameters, a factor of three.

A.3 Screen B: Graph Readout Scope and Reducer

Figures A.3 and A.4 complement the endpoint comparison in Table 6.5 with the eight training trajectories and the full-test interaction between readout scope and reducer.




Figure A.3: Screen B. Smoothed training episodic survival curves by readout scope and reducer (seed 0).




Figure A.4: Screen B. Full-test survival of each best checkpoint. Connecting the two reducers within a scope makes the scope-reducer interaction visible; energized-only readout leads under both reductions.

A.4 Screen C: Structural Augmentations of the Busbar Graph

Tables A.9 and A.10 and Figures A.5 and A.6 complement the factorial cube retained in Section 6.6 with the complete results.

Table A.9: Screen C. Cells ranked by full-test survival, with the difference from the unaugmented e0n0v0 baseline.

Cell	Mean (%)	Std (pp)	Min (%)	Max (%)	Δ baseline (pp)
e0n0v0	97.31	1.22	96.55	98.72	0.00
e1n0v0	96.20	3.94	91.66	98.68	-1.11
e0n0v1	95.63	6.19	88.52	99.82	-1.68
e0n1v0	93.51	6.39	86.72	99.40	-3.80
e1n1v0	90.01	4.92	85.42	95.19	-7.30
e1n0v1	89.92	9.42	80.48	99.33	-7.39
e0n1v1	68.91	5.77	62.59	73.91	-28.40
e1n1v1	66.11	33.95	27.86	92.68	-31.20

Table A.10: Screen C. Main effect of each factor (12 runs per level) and the four single-factor flips of the cube, each one enabling that factor with the other two held fixed. **Spread** is the range of those four flips.

Factor	Off (%)	On (%)	Effect (pp)	Flips (pp)	Mean flip (pp)	Spread (pp)
<i>e</i> (substation edges)	88.84	85.56	-3.28	-1.1, -5.7, -3.5, -2.8	-3.28	4.60
<i>v</i> (virtual readout)	94.26	80.14	-14.11	-1.7, -24.6, -6.3, -23.9	-14.11	22.92
<i>n</i> (summary nodes)	94.77	79.63	-15.13	-3.8, -26.7, -6.2, -23.8	-15.13	22.92




Figure A.5: Screen C. Test episodic survival during training, one panel per augmented cell against the unaugmented **e0n0v0** baseline in grey. Thin lines are the three seeds, the thick line is their mean, and the band is one standard deviation.



Figure A.6: Screen C. The same eight cells as a heatmap of the full-test survival of each run's best checkpoint, with the three individual seed values printed under each cell mean.

A.5 Screen D: Generator and Load Message Direction

Figure A.7 shows the training trajectories behind the endpoint heatmap retained in Section 6.7, while Table A.11 gives the complete ranking.




Figure A.7: Screen D. Test episodic survival during training, one panel per direction pair against the bidirectional baseline in grey. Thin lines are the three seeds, the thick line is their mean, and the band is one standard deviation.

Table A.11: Screen D. Direction pairs ranked by full-test survival, with the difference from the bi/bi baseline.

Generator / load	Mean (%)	Std (pp)	Min (%)	Max (%)	Δ baseline (pp)
b2a / b2a	97.94	2.20	95.40	99.25	+5.30
b2a / a2b	97.43	1.42	96.32	99.03	+4.79
a2b / b2a	94.47	7.53	85.78	99.03	+1.83
bi / bi	92.65	8.24	83.21	98.40	0.00
bi / a2b	91.35	5.54	86.52	97.40	-1.30
bi / b2a	87.76	11.04	77.64	99.53	-4.89
b2a / bi	84.93	13.01	70.11	94.46	-7.72
a2b / bi	83.16	12.75	72.27	97.19	-9.49
a2b / a2b	56.41	22.12	30.88	69.60	-36.24

A.6 Screen E: Candidate-Action Scoring

Figures A.8 and A.9 complement the full-test table retained in Section 6.8 with the training trajectories and endpoint ordering.




Figure A.8: Screen E. Test episodic survival during training, smoothed over nine evaluations. Six of the eight cells converge to the top of the range and stay there; `typed_mean/f1/a0h1` never does.




Figure A.9: Screen E. Full-test survival of each best checkpoint. The dashed line marks 98%.

A.7 Encoder Input Distributions

Chapter 7 retains the compact shift table (Table 7.1) and the consolidated distribution panel (Figure 7.1). This section provides the raw moments behind those statistics and the tail behaviour that linear axes compress. All values are measured on the disaggregated line-node graph under the protocol of Section 7.2.

Table A.12: Raw moments behind Table 7.1, with the Wasserstein distance. “Owning” restricts each channel to the node type that carries it. W_1/σ_p tracks the offset almost exactly on the displaced channel, which is why the main table omits it: for distributions that differ mainly by a shift the two statistics carry the same information.

Channel	Rows	bus14		WCCI		offset	ratio	W_1/σ_p
		mean	sd	mean	sd			
p	all	6.538	17.615	5.461	37.718	−0.037	2.141	0.204
	owning	27.306	26.998	19.067	68.613	−0.158	2.541	0.485
theta	all	−1.236	2.570	0.305	2.127	+0.653	0.827	0.653
	owning	−5.162	2.707	1.064	3.871	+1.864	1.430	1.864
rho	all	0.133	0.203	0.116	0.187	−0.088	0.922	0.088
	owning	0.364	0.168	0.319	0.176	−0.260	1.050	0.261




Figure A.10: Full range of the two asset channels on log density. WCCI carries a generator pinned near its 399 MW rating and a load region reaching −281 MW, neither of which occurs anywhere on **bus14**: a transferred encoder has no calibration for either. The angle panel shows the displacement rather than a difference in width.

A.8 Deterministic Feature Scales

Section 5.1 gives the two references that carry physical content, the power base and the angular scale. The complete list of divisors, including the discrete counters, is below.

Table A.13: Deterministic physical scales used for continuous graph features. The power base is the largest installed generator rating on the grid; the remaining discrete channels use the corresponding Grid2Op parameter bound.

Features	Scale s_f	Source
<code>gen_p</code> , <code>load_p</code> , <code>p</code>	Power base in MW	Largest installed generator rating.
<code>gen_theta</code> , <code>load_theta</code> , <code>theta</code> , <code>theta_diff</code>	180 degrees	Fixed angular scale.
<code>time_before_cooldown_sub</code>	Maximum substation cooldown	Grid2Op environment parameter.
<code>time_before_cooldown_line</code>	Maximum line cooldown	Grid2Op environment parameter.
<code>timestep_overflow</code>	Allowed overflow duration	Grid2Op environment parameter.
Maintenance time and duration	Episode horizon	Grid2Op chronic horizon.

Appendix B

GINE Graph-Actor Architecture

This appendix specifies the complete edge-aware actor used in the graph experiments: the shared input and encoder pipeline, the GINE message-passing update, the graph readout, and the agent-specific action head.

B.1 Shared input and encoder pipeline

For agent a , the encoder receives the local graph $\mathcal{G}_t^{(a)}$. Candidate edges with `edge_mask`= 0 are removed before message passing, so the layer sees only the active directed set $\mathcal{E}_t^{(a)}$. The source of $u \rightarrow v$ sends a message and v aggregates it. Node and edge masks do not become learned features.

Each raw node vector $\mathbf{x}_v$ first passes through

$$\mathbf{h}_v^{(0)} = \text{LayerNorm}(\text{ReLU}(\mathbf{W}_x \mathbf{x}_v + \mathbf{b}_x)), \quad (\text{B.1})$$

which produces a 128-dimensional node state. Learned global node-identifier embeddings and the optional external edge pre-encoder are disabled. Every GINE layer nevertheless contains its own trainable projection of the edge vector to the dimension required by that layer.

Two message-passing layers are used. After each layer, the implementation applies ReLU followed by LayerNorm:

$$\mathbf{h}_v^{(k+1)} = \text{LayerNorm}\left(\text{ReLU}(\tilde{\mathbf{h}}_v^{(k+1)})\right). \quad (\text{B.2})$$

The same encoder parameters are shared by all agents. Each agent nevertheless passes its own local graph, with its own node set, active edges, and one-hop context, through that shared encoder.

B.2 GINE message passing

For an active edge $u \rightarrow v$, layer k projects the complete edge feature vector and adds it to the sender state:

$$\tilde{\mathbf{e}}_{uv}^{(k)} = \mathbf{W}_e^{(k)} \mathbf{e}_{uv} + \mathbf{b}_e^{(k)}, \quad (\text{B.3})$$

$$\mathbf{m}_{u \rightarrow v}^{(k)} = \text{ReLU}\left(\mathbf{h}_u^{(k)} + \tilde{\mathbf{e}}_{uv}^{(k)}\right), \quad (\text{B.4})$$

$$\tilde{\mathbf{h}}_v^{(k+1)} = \text{MLP}^{(k)}\left((1 + \epsilon)\mathbf{h}_v^{(k)} + \sum_{u:(u,v) \in \mathcal{E}_t} \mathbf{m}_{u \rightarrow v}^{(k)}\right). \quad (\text{B.5})$$

The implementation uses the PyTorch Geometric default $\epsilon = 0$ without a trainable epsilon parameter. Each internal GINE MLP is Linear–ReLU–Linear and outputs 128 features. Equation B.5 shows why edge attributes are part of the message content. Two identical node states can send different messages when one edge is an overloaded physical line and the other is an asset-attachment relation. Likewise,

reversing a typed disaggregated edge activates a different relation-one-hot column and can produce a different message.

GINE uses an unweighted sum over incoming active edges. It does not learn a softmax competition between neighbors. The edge embedding changes what is sent; the number of incoming messages and their vector sum determine the aggregate.

B.3 Pooling and graph output

With mean pooling, the graph representation after the second message-passing layer is

$$\bar{\mathbf{h}}_G = \frac{\sum_v m_v \mathbf{h}_v^{(K)}}{\max(1, \sum_v m_v)}, \quad (\text{B.6})$$

where m_v is the dynamic visibility mask. This is one mean over all visible node types, not a separate mean per type. Therefore:

- the busbar graph pools controlled busbars and only the far-end busbars currently attached by live boundary lines;
- the disaggregated graph additionally pools controlled generators and loads, but never neighboring private assets;
- the disaggregated line-node graph additionally pools transmission-line nodes and contains no neighboring equipment.

An inactive contextual candidate contributes neither to the numerator nor to the denominator. The ordinary GINE mean readout includes visible context; `controlled_mean` replaces m_v by the product of the visibility and controlled-node masks. Visible context can still influence controlled nodes through message passing, but it is not aggregated directly. Sum and maximum pooling, including their `controlled_*` variants, change only the readout and do not change message passing or the information boundary.

The pooled 128-dimensional vector then passes through the shared graph output projection

$$\mathbf{z}_a = \text{ReLU}(\mathbf{W}_o \bar{\mathbf{h}}_G + \mathbf{b}_o), \quad \mathbf{z}_a \in \mathbb{R}^{128}. \quad (\text{B.7})$$

When virtual-node readout is selected, the final virtual-node state replaces $\bar{\mathbf{h}}_G$; the output projection and its dimension stay the same.

B.4 Agent-specific action head

The flat observation is not concatenated to $\mathbf{z}_a$. Each agent has its own two-layer MLP action head:

$$\mathbf{q}_a^{(1)} = \text{ReLU}(\mathbf{W}_a^{(1)} \mathbf{z}_a + \mathbf{b}_a^{(1)}), \quad (\text{B.8})$$

$$\mathbf{q}_a^{(2)} = \text{ReLU}(\mathbf{W}_a^{(2)} \mathbf{q}_a^{(1)} + \mathbf{b}_a^{(2)}), \quad (\text{B.9})$$

$$\boldsymbol{\ell}_a = \mathbf{W}_a^{(3)} \mathbf{q}_a^{(2)} + \mathbf{b}_a^{(3)}, \quad \boldsymbol{\ell}_a \in \mathbb{R}^{|\mathcal{A}_a|}, \quad (\text{B.10})$$

$$\pi_a(\cdot \mid \mathcal{G}_t^{(a)}) = \text{Categorical}(\text{logits} = \boldsymbol{\ell}_a). \quad (\text{B.11})$$

Both hidden layers contain 128 units. During training, an action is sampled from this categorical distribution. Deterministic evaluation selects $\arg \max_j \ell_{a,j}$. Action zero is the local do-nothing action.

Only the GNN encoder and graph output projection are shared across agents. The action heads are agent-specific because agents can have different local action sets and therefore different output widths $|\mathcal{A}_a|$. The critic is a separate MLP over the centralized flat state and does not reuse the actor graph embedding.

Equation B.11 is the default head. The optional candidate-action head of Section 5.5 keeps this pipeline up to $\mathbf{z}_a$ and additionally consumes the post-message-passing node states $\mathbf{h}_v^{(K)}$, which are the same states that Equation B.6 averages.

B.5 Exact node inputs by representation

All node types of one representation receive the same feature columns. Columns not listed for a node type are zero. Type identity reaches GINE through the `is_*` columns. The inventory below describes visible nodes. For an inactive contextual candidate, the complete row—including type and role columns—is zero, all incident edges are inactive, and the node is excluded from readout.

Table B.1: Node types and nonzero input features received by GINE, under the default `gnn_angle_representation = node`. Under `edge_diff` the `gen_theta`, `load_theta` and `theta` entries leave the node rows in every representation; the busbar and disaggregated graphs move the angle drop onto the physical line edge, while the disaggregated line-node graph writes it on the transmission-line row as `theta_diff` (Section 5.1.2).

Representation	Node type	Populated input features
bus	Busbar	<code>gen_p</code> , <code>gen_theta</code> , <code>load_p</code> , <code>load_theta</code> , <code>time_before_cooldown_sub</code> , <code>domain_mask</code>
disaggregated	Busbar	<code>is_busbar</code> , <code>time_before_cooldown_sub</code> , <code>domain_mask</code>
	Generator	<code>is_generator</code> , <code>p=gen_p</code> , <code>theta=gen_theta</code> , <code>domain_mask</code>
	Load	<code>is_load</code> , <code>p=load_p</code> , <code>theta=load_theta</code> , <code>domain_mask</code>
disaggregated line-node	Busbar	<code>is_busbar</code> , <code>time_before_cooldown_sub</code> , <code>domain_mask</code>
	Generator	<code>is_generator</code> , <code>p=gen_p</code> , <code>theta=gen_theta</code> , <code>domain_mask</code>
	Load	<code>is_load</code> , <code>p=load_p</code> , <code>theta=load_theta</code> , <code>domain_mask</code>
	Transmission line	<code>is_transmission_line</code> , <code>line_status</code> , <code>rho</code> , <code>timestep_overflow</code> , <code>time_before_cooldown_line</code> , optional maintenance times, <code>domain_mask</code>
Any augmented schema	Substation summary	Learned projection of the substation structural descriptor; no raw physical feature row
	Virtual node	One learned vector shared across graphs; no raw physical feature row

B.6 Exact edge inputs by representation

The edge vector contains all edge information consumed by GINE. Inactive candidates are filtered out before the vectors in Table B.2 reach a convolution.

Table B.2: Directed edge relations and edge-feature vectors received by GINE, under the default `gnn_angle_representation = node`. “Relation” denotes one active relation-one-hot column. Under `edge_diff` the physical-line block of the busbar and disaggregated graphs gains a `theta_diff` column; the line-node rows are unchanged, because there the drop is a line-node feature and the attachment edges stay free of physical channels.

Representation	Directed relation	Edge features
bus	Physical busbar $\leftrightarrow$ busbar	<code>line_status</code> , <code>rho</code> , <code>timestep_overflow</code> , <code>time_before_cooldown_line</code> , optional maintenance times; no relation one-hot
	Optional self or same-substation	Same width; all physical values zero; no relation one-hot
disaggregated	Physical busbar $\leftrightarrow$ busbar	Complete measured physical-line block plus <code>relation_physical_line</code>
	Generator $\leftrightarrow$ busbar	Neutral physical block plus direction-specific generator relation
	Load $\leftrightarrow$ busbar	Neutral physical block plus direction-specific load relation
	Optional self or same-substation	Neutral physical block plus the corresponding relation
disaggregated line-node	Busbar $\leftrightarrow$ line node	One origin/extremity-and-direction relation; no line-state attributes
	Generator $\leftrightarrow$ busbar	One direction-specific generator relation
	Load $\leftrightarrow$ busbar	One direction-specific load relation
	Optional self or same-substation	One corresponding relation
Any augmented schema	Busbar $\rightarrow$ summary/virtual, with optional reverse	Zero edge vector; the relation carries no flow

The complete relation names and activation rules are given in Tables 5.5 and 5.7. Those tables distinguish, for example, generator-to-busbar from busbar-to-generator and origin from extremity line attachments. These distinctions matter directly to an edge-aware model such as GINE, which folds the edge vector into every message.

Appendix C

Graph Construction Details

This appendix collects the construction and sizing details of the three graph schemas of Section 5.2: how large each candidate graph is, how inactive candidates are neutralized, and how a local actor graph is cut out of the full grid. None of it changes the representations themselves.

C.1 Candidate graph sizes

Let S be the number of included substations, B the busbars per substation, L the included physical lines, and N_g, N_d the generators and loads. Write $d_g, d_d, d_\ell \in \{1, 2\}$ for the number of retained directions of the generator, load, and line-node attachments: 2 for **bidirectional** and 1 for either one-way mode. Let P_ℓ be the number of represented line endpoints: an internal line contributes two and a boundary line in a local disaggregated line-node graph contributes only its controlled endpoint. Thus $L \leq P_\ell \leq 2L$. Table C.1 gives the size of the fixed candidate graph before any mask is applied.

Table C.1: Size of the candidate graph of each schema. Enabling self edges adds $|\mathcal{V}|$ directed edges and enabling same-substation edges adds $SB(B - 1)$, in every schema.

Schema	Candidate nodes	Candidate directed edges
bus	SB	$2B^2L$
disaggregated	$SB + N_g + N_d$	$2B^2L + B(d_gN_g + d_dN_d)$
disaggregated line-node	$SB + N_g + N_d + L$	$Bd_\ell P_\ell + B(d_gN_g + d_dN_d)$

Three consequences are worth naming. A busbar graph with both optional relations enabled therefore has $2B^2L + SB + SB(B - 1)$ candidate directed edges. The disaggregated schema adds $N_g + N_d$ nodes and between $B(N_g + N_d)$ and $2B(N_g + N_d)$ asset candidates to the busbar graph, and buys individual asset states with them. For a complete state graph, $P_\ell = 2L$, so the disaggregated line-node schema replaces the $2B^2L$ busbar-to-busbar term by $2Bd_\ell L$: identical at $B = 2$ and cheaper for $B > 2$. A local boundary line has candidates only at its controlled endpoint, which reduces P_ℓ and prevents construction of a far-end attachment. The schema adds L line nodes and one message-passing hop for internal endpoint exchange.

The hierarchical augmentations of Section 5.3 add on top of any of these: one substation-summary node per substation, with SB directed edges for **toward_summary** or $2SB$ for **bidirectional**; and one virtual node, with N_{bus} or $2N_{\text{bus}}$ edges, or S or $2S$ when summary nodes are also enabled.

C.2 Indexing and masking conventions

Grid2Op numbers busbars from one in its topology vector; those identifiers are converted to the zero-based index b used in the node identifier $i(s, b)$.

An inactive candidate edge keeps its row in the edge-feature tensor, so the tensor shape is fixed for an episode, but the row is zero-filled and its **edge_mask** entry is 0. Masked edges are removed before the

encoder is called, so those zero rows never generate a message. This is what allows a topology change to alter connectivity without rebuilding the graph. An inactive contextual node likewise retains its candidate row, but its complete feature vector is zero and its `node_mask` entry is 0. Every incident physical or structural edge is deactivated. The row therefore participates in neither message passing nor graph readout; a mean readout excludes it from both numerator and denominator.

C.3 Local actor graphs

A local actor graph always contains the controlled substations and every line touching them. Additional membership is schema-specific. The `bus` graph contains only the far-end busbar currently attached by each live boundary line, and that busbar exposes aggregated neighboring power and angle. The direct disaggregated graph contains the same far-end busbar, but never its generator or load nodes. The disaggregated line-node graph represents the shared line itself and therefore constructs no neighboring busbar or private asset and no far-end attachment edge. The centralized state graph is intentionally unrestricted for centralized training. Ordinary pooling aggregates all visible nodes, whereas `controlled_*` pooling restricts readout to action-domain nodes; neither choice changes the construction boundary.

C.4 Boundary verification

The regression suite checks that contextual visibility follows a live busbar attachment, moves when the attachment is switched, and disappears after a line outage; controlled nodes remain present; structural relations cannot make a hidden candidate visible; inactive rows are zero; neighboring private assets are absent from disaggregated graphs; and all neighboring equipment is absent from disaggregated line-node local graphs. Candidate-action tests additionally verify that a boundary line remains addressable through its line node without adding far-end equipment. Controlled-readout tests verify both that the pooled set is fixed and that gradients can still reach visible context through message passing.

Appendix D

Experiment Configurations

This appendix provides the complete run configurations for the GINE encoder and heavy GINE ablations.

Table D.1: GINE rerun configuration summary. The reference run is Heavy GINE Baseline.

Run	Actor/critic encoder	GNN hidden/layers	Actor/critic heads	Concat flat	Edge pre-enc.	Entropy	Init action-0	Opt. critic	Difference from Heavy GINE Baseline
Heavy GINE Baseline	<code>gnn / gnn</code>	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.7	false	Baseline: heavy concat-flat GINE actor, GNN critic, and legacy critic updates
Light GINE (Concat, MLP Critic)	<code>gnn / mlp</code>	64 / 1	$2 \times 128 / 2 \times 128$	true	false	0.02	0.7	true	Light GINE actor, MLP critic, smaller heads, no edge pre-encoder, optimized critic updates
Heavy GINE (No Bias)	<code>gnn / gnn</code>	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.0	false	Bias ablation: same as baseline, but the initial do-nothing prior is 0.0
No Bias, Opt Critic	<code>gnn / gnn</code>	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.0	true	Bias 0.0 plus optimized critic updates; otherwise heavy concat-flat GINE with GNN critic
Optimized Heavy GINE	<code>gnn / gnn</code>	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.01	0.0	true	Optimized heavy GINE: bias 0.0, lower entropy 0.01, optimized critic updates
Optimized Light GINE	<code>gnn / mlp</code>	64 / 1	$2 \times 128 / 2 \times 128$	false	false	0.01	0.0	true	Light optimized GINE: MLP critic, no concat-flat, lower entropy 0.01, bias 0.0, optimized critic updates

Table D.2: Heavy-GINE parameter comparison from configs/gine_s0_s1_s2.

Run	Main change from baseline	GNN	GNN size	Critic	Actor/critic heads	Concat	Shared	Edge pre.	Node ID	Entropy	Init action-0	Opt. critic
No Flat Concat	Removes flat-observation concatenation	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	no	yes	yes	yes	0.02–0.02	0.7	no
Heavy GINE Baseline	Reference: heavy concat-flat GINE with GNN critic	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
Opt. Critic Updates	Enables optimized critic updates	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	yes
MLP Critic	Replaces the GNN critic with an MLP critic	gine	2×128	mlp	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
Unshared Actor	Uses non-shared actor GNN weights	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	no	yes	yes	0.02–0.02	0.7	no
Light GINE Encoder	Uses a lighter GINE encoder	gine	1×64	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
Zero Entropy Decay	Decays entropy from 0.02 to 0.0	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.00	0.7	no
No Node-ID Embs	Removes learned node-ID embeddings	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	no	0.02–0.02	0.7	no
GAT Message Passing	Replaces GINE with GAT message passing	gat	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
Weighted GCN	Replaces GINE with weighted GCN message passing	gcn	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
Light GINE (No Concat, MLP Critic)	Light no-concat GINE with MLP critic and optimized critic updates	gine	1×64	mlp	$2 \times 128 / 2 \times 128$	no	yes	no	yes	0.02–0.02	0.7	yes
Light GINE (Concat, MLP Critic)	Light concat-flat GINE with MLP critic and optimized critic updates	gine	1×64	mlp	$2 \times 128 / 2 \times 128$	yes	yes	no	yes	0.02–0.02	0.7	yes
No Initial Bias	Removes the initial do-nothing bias	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.0	no

Appendix E

Full-Test Checkpoint Registry

Table [E.1](#) reports the checkpoint step used for each full-test result table. The values are the `checkpoint_global_step` fields stored in the corresponding JSON files under `Topology_Task/outputs/full_test_eval`. The do-nothing policy has no learned checkpoint. For the behavioral cloning students, the stored step is the student checkpoint’s optimizer-step count rather than a MAPPO environment step.

Table E.1: Checkpoint steps used for the full-test evaluations.

Run type	Seed 0 checkpoint step	Seed 1 checkpoint step	Seed 2 checkpoint step
Do-nothing policy	—	—	—
Baseline flat policy	13,600,000	13,040,000	13,760,000
Global rho heuristic, $\rho_{safe} = 0.90$	13,920,000	14,320,000	14,800,000
Local rho heuristic, $\rho_{safe} = 0.90$	14,160,000	14,400,000	14,160,000
Intervention gate, final-action MAP	12,000,000	12,720,000	13,760,000
Intervention gate, hierarchical greedy	7,120,000	14,720,000	13,760,000
Flat Sparse16, $\lambda_{fixed} = 0.000$	9,600,000	12,160,000	7,920,000
Flat Sparse16, $\lambda_{fixed} = 0.001$	10,640,000	11,520,000	11,600,000
Flat Sparse16, $\lambda_{fixed} = 0.003$	12,000,000	12,960,000	12,480,000
Flat Sparse16, $\lambda_{fixed} = 0.010$	8,960,000	5,760,000	12,560,000
Gated Sparse16, $\lambda_{fixed} = 0.000$	8,480,000	4,320,000	6,000,000
Gated Sparse16, $\lambda_{fixed} = 0.001$	5,760,000	7,600,000	3,280,000
Gated Sparse16, $\lambda_{fixed} = 0.003$	8,960,000	7,120,000	7,600,000
Gated Sparse16, $\lambda_{fixed} = 0.010$	12,000,000	7,840,000	7,920,000
Flat local-safe AIB, $d = 0.20$	13,120,000	9,360,000	12,160,000
Flat local-safe AIB, $d = 0.10$	12,880,000	13,200,000	12,720,000
Flat local-safe AIB, $d = 0.35$	12,800,000	5,360,000	11,520,000
Gated local-safe AIB, $d = 0.20$, hierarchical greedy	14,960,000	5,280,000	16,800,000
Flat non-idle AIB, $d = 0.20$	12,480,000	14,880,000	11,920,000
Teacher: MAPPO with local rho heuristic	14,160,000	14,400,000	14,160,000
Student BC, balanced non-idle 0.50, weight 5, aux 0.50	12,375	14,994	13,617
Student BC, balanced non-idle 0.20, weight 3, aux 0.25	12,375	14,994	13,617

Bibliography

- [1] E. B. Fisher, R. P. O'Neill, and M. C. Ferris, "Optimal transmission switching," *IEEE Transactions on Power Systems*, vol. 23, no. 3, pp. 1346–1355, 2008.
- [2] A. Marot et al., "Learning to run a power network challenge for training topology controllers," *Electric Power Systems Research*, 2021.
- [3] J. Viebahn, M. Naglic, A. Marot, B. Donnot, and S. H. Tindemans, "Potential and challenges of AI-powered decision support for short-term system operations," in *CIGRE Paris Session*, 2022.
- [4] B. Donnot, "Grid2op: A testbed platform to model sequential decision making in power systems," *arXiv preprint arXiv:2011.07929*, 2020.
- [5] M. Tan, "Multi-agent reinforcement learning: Independent vs. cooperative agents," in *International Conference on Machine Learning*, 1993.
- [6] R. Lowe, Y. Wu, A. Tamar, J. Harb, P. Abbeel, and I. Mordatch, "Multi-agent actor-critic for mixed cooperative-competitive environments," in *Advances in Neural Information Processing Systems*, 2017.
- [7] C. Yu et al., "The surprising effectiveness of ppo in cooperative multi-agent games," *Advances in Neural Information Processing Systems*, 2022.
- [8] M. Hassouna, C. Holzhüter, P. Lytaev, J. Thomas, B. Sick, and C. Scholz, *Graph reinforcement learning for power grids: A comprehensive survey*, 2024. arXiv: [2407.04522](https://arxiv.org/abs/2407.04522). [Online]. Available: <https://arxiv.org/abs/2407.04522>
- [9] M. Subramanian, J. Viebahn, S. H. Tindemans, B. Donnot, and A. Marot, "Exploring grid topology reconfiguration using a simple deep reinforcement learning approach," in *IEEE Madrid PowerTech*, 2021.
- [10] D. Yoon, S. Hong, B.-J. Lee, and K.-E. Kim, "Winning the L2RPN challenge: Power grid management via semi-Markov afterstate actor-critic," in *International Conference on Learning Representations*, 2021.
- [11] A. Chauhan, M. Baranwal, and A. Basumatary, "PowRL: A reinforcement learning framework for robust management of power networks," in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 37, 2023, pp. 14 757–14 764.
- [12] M. Lehna, J. Viebahn, A. Marot, S. Tomforde, and C. Scholz, "Managing power grids through topology actions: A comparative study between advanced rule-based and reinforcement learning agents," *Energy and AI*, vol. 14, p. 100 276, 2023.
- [13] B. Manczak, J. Viebahn, and H. van Hoof, *Hierarchical reinforcement learning for power network topology control*, 2023. arXiv: [2311.02129](https://arxiv.org/abs/2311.02129). [Online]. Available: <https://arxiv.org/abs/2311.02129>
- [14] E. van der Sar, A. Zocca, and S. Bhulai, *Multi-agent reinforcement learning for power grid topology optimization*, 2023. arXiv: [2310.02605](https://arxiv.org/abs/2310.02605). [Online]. Available: <https://arxiv.org/abs/2310.02605>
- [15] B. de Mol, D. Barbieri, J. Viebahn, and D. Grossi, "Centrally coordinated multi-agent reinforcement learning for power grid topology control," in *Proceedings of the 16th ACM International Conference on Future and Sustainable Energy Systems (e-Energy)*, 2025.
- [16] E. Marchesini et al., "MARL2Grid-TR: A multi-agent RL benchmark in power grid operations," in *International Conference on Learning Representations*, 2026.
- [17] T. N. Kipf and M. Welling, "Semi-supervised classification with graph convolutional networks," *International Conference on Learning Representations*, 2017.

- [18] P. Veličković, G. Cucurull, A. Casanova, A. Romero, P. Liò, and Y. Bengio, “Graph attention networks,” in *International Conference on Learning Representations*, 2018.
- [19] K. Xu, W. Hu, J. Leskovec, and S. Jegelka, “How powerful are graph neural networks?” *International Conference on Learning Representations*, 2019.
- [20] M. de Jong, J. Viebahn, and Y. Shapovalova, *Generalizable graph neural networks for robust power grid topology control*, 2025. arXiv: [2501.07186](https://arxiv.org/abs/2501.07186). [Online]. Available: <https://arxiv.org/abs/2501.07186>
- [21] E. Marchesini et al., *RL2Grid: Benchmarking reinforcement learning in power grid operations*, 2025. arXiv: [2503.23101](https://arxiv.org/abs/2503.23101). [Online]. Available: <https://arxiv.org/abs/2503.23101>
- [22] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [23] J. Achiam, D. Held, A. Tamar, and P. Abbeel, “Constrained policy optimization,” *International Conference on Machine Learning*, 2017. [Online]. Available: <https://arxiv.org/abs/1705.10528>
- [24] A. Stooke, J. Achiam, and P. Abbeel, “Responsive safety in reinforcement learning by PID lagrangian methods,” *International Conference on Machine Learning*, 2020. [Online]. Available: <https://arxiv.org/abs/2007.03964>
- [25] N. Carrara, E. Leurent, R. Laroche, T. Urvoy, and O.-A. Maillard, “Budgeted reinforcement learning in continuous state space,” *Advances in Neural Information Processing Systems*, 2019. [Online]. Available: <https://arxiv.org/abs/1903.01004>
- [26] L. Lautenbacher et al., *Multi-objective reinforcement learning for power grid topology control*, 2025. arXiv: [2502.00040](https://arxiv.org/abs/2502.00040). [Online]. Available: <https://arxiv.org/abs/2502.00040>
- [27] W. Hu et al., “Strategies for pre-training graph neural networks,” in *International Conference on Learning Representations*, 2020.