

Thesis Plan

Decentralized Multi-Agent Reinforcement Learning for Power Grid Topology Control

Corentin Plumet

July 22, 2026

Contents

1	Thesis Plan	1
1.1	Plan Inspiration	1
1.2	Working Title	1
1.3	Central Thesis Question	1
1.4	Research Questions	1
1.5	Expected Contributions	1
1.6	Chapter Outline	2
1.7	Appendix Plan	4
1.8	Immediate Writing Order	4
1.9	Open Decisions	4
2	Methods	5
2.1	Graph Representation	5
2.1.1	Busbar nodes	5
2.1.2	Candidate line edges and dynamic topology	6
2.1.3	Physical scaling and running normalization	7
2.1.4	Additional relation edges	7
2.1.5	Local actor graphs and the global state graph	8
2.1.6	Boundary between representation and encoder	9
2.2	Heterogeneous Graph Representation	9
2.2.1	Typed nodes and features	9
2.2.2	Typed relations	10
2.2.3	Tensor representation and encoder compatibility	11
2.2.4	Local heterogeneous graphs	11
2.3	Heterogeneous Graph with Transmission-Line Nodes	11
2.3.1	Node allocation and features	12
2.3.2	Endpoint-specific typed relations	13
2.3.3	Tensor representation and local graphs	13
2.4	Substation Summary Nodes	13
2.5	Virtual-Node Graph Readout	14
2.5.1	Augmented graph	15
2.5.2	Learned summary instead of terminal pooling	15
2.5.3	Relation attributes and encoder compatibility	16
2.6	Staged Graph-Architecture Screening Protocol	16
2.6.1	Common screening controls	16
2.6.2	Stage 1: representation and input preprocessing	16
2.6.3	Stage 2: structural graph factors	17
2.6.4	Stage 3: encoder comparison	17
2.6.5	Promotion, confirmation, and reproducibility	18
3	Experimental Results	19
3.1	How to Read This Chapter	19
3.2	Evaluation Protocol	19
3.3	Experiment Index	19
3.4	IEEE 14-Bus Experiments	20
3.4.1	A0 Core Reruns	20
3.4.2	Initial Do-Nothing Bias	22

3.4.3	GINE Encoder Reruns	23
3.4.4	Heavy GINE	26
3.5	WCCI Action-Space Reduction	30
3.6	Decentralized Sparse Intervention Control	30
3.6.1	Evaluation-Time Rho Heuristics	31
3.6.2	Intervention-Gated Actor	31
3.6.3	Fixed Sparse-Action Penalty	32
3.6.4	Adaptive Intervention Budget	32
3.7	Current Takeaway	34
A	GINE Architecture	36
B	Full-Test Checkpoint Registry	37

Chapter 1

Thesis Plan

1.1 Plan Inspiration

This plan follows the structure of the reference thesis in `mAI2024deMo1B.pdf`: start from the power-grid motivation and research questions, introduce the scientific background, describe the environment before the algorithms, then separate methods from results and close by answering the research questions. The adapted structure below is tuned to this repository: Grid2Op topology control, decentralized MAPPO, graph encoders, rerun ablations, and sparse intervention mechanisms.

1.2 Working Title

Decentralized Multi-Agent Reinforcement Learning for Sparse Power Grid Topology Control

1.3 Central Thesis Question

How can decentralized multi-agent reinforcement learning agents learn reliable power-grid topology control policies that preserve survival performance while scaling to combinatorial action spaces and avoiding unnecessary interventions?

1.4 Research Questions

1. How should the Grid2Op topology-control problem be formulated as a decentralized multi-agent decision problem?
2. Which architectural choices improve survival and generalization: flat MLP policies, graph neural network encoders, shared actor weights, critic design, or optimized critic updates?
3. How do initialization and training choices, especially the initial do-nothing bias and reward normalization, affect exploration and learned behavior?
4. Can sparse-control mechanisms, such as rho-based overrides, intervention gates, fixed penalties, or adaptive intervention budgets, reduce unnecessary topology actions without sacrificing episodic survival?

1.5 Expected Contributions

- A clear decentralized formulation of topology optimization in Grid2Op.
- A controlled comparison of MAPPO baselines, GNN encoders, critic variants, and initialization choices.
- An ablation study of sparse intervention mechanisms for more operator-like control.
- A reproducible experiment and analysis pipeline based on W&B histories, cached metrics, notebooks, and final evaluation scripts.

1.6 Chapter Outline

Chapter 1: Introduction

Goal. Motivate topology control as a low-cost but difficult operational tool for power-grid reliability.

Planned content.

- Energy-transition context: higher demand, renewable uncertainty, and stressed transmission networks.
- Topology control as an alternative or complement to redispatching and grid expansion.
- Why the problem is hard: sequential control, physical constraints, rare failures, and combinatorial topology actions.
- Why decentralized MARL is a natural fit: local substations, regional decision making, reduced per-agent action spaces, and scalability.
- Research questions, contributions, and thesis outline.

Chapter 2: Background and Related Work

Goal. Give the reader enough power-systems and RL background to understand the design choices.

Planned content.

- Power-grid operation: substations, busbars, line loading, overloads, contingencies, and remedial actions.
- Grid2Op and the Learning to Run a Power Network setting.
- Reinforcement learning basics: MDPs, policy gradients, PPO, and MAPPO.
- Multi-agent reinforcement learning: decentralized execution, centralized training, credit assignment, and non-stationarity.
- Graph neural networks for grid-structured observations.
- Safe, sparse, and constrained RL ideas relevant to intervention budgets.

Chapter 3: Environment and Problem Formulation

Goal. Define exactly what the agents observe, choose, optimize, and are evaluated on.

Planned content.

- Grid2Op environment used in the repository and the selected chronics.
- Multi-agent decomposition: agent-region or agent-substation assignment.
- Observation spaces: local grid state, line loadings, topology state, and optional graph observations.
- Action spaces: local topology actions, action 0 as do nothing, legal action constraints, and action-space reduction.
- Reward and evaluation metrics: episodic survival, illegal action rate, action-0 fraction, non-idle fraction, and train/test chronic split.
- Reproducibility setup: seeds, W&B logging, cached histories, and deterministic evaluation.

Chapter 4: Methods

Goal. Present the implemented learning systems and the ablations.

Planned content.

- MAPPO baseline: actor, critic, rollout collection, PPO objective, and centralized training assumptions.
- Baseline MLP agents and the A0 family of reruns.
- Graph-encoder agents: bus graph construction, GINE/GAT/GCN variants, node and edge pre-encoders, graph readout, and flat-feature concatenation.
- Critic variants: GNN critic versus MLP critic, legacy versus optimized critic updates.

- Initialization and exploration choices: entropy coefficient, entropy decay, and initial do-nothing probability.
- Sparse-control mechanisms: evaluation-time rho heuristic, intervention gate, fixed sparse-action penalties, and adaptive intervention budgets.
- Experimental protocol for ablations: controlled seeds, fixed evaluation cadence, and fair baseline comparisons.

Chapter 5: Experiments and Results

Goal. Show which design choices matter and what behaviors the agents learn.

Planned content.

- Main baseline performance: training stability and test survival.
- GNN architecture ablations: concat-flat, actor sharing, GINE versus GAT or weighted GCN, light versus full encoders, and critic design.
- A0 rerun ablations: optimized critic updates, smaller MLP, reward normalization, and do-nothing initialization bias.
- Sparse-control ablations: fixed penalty, global-safe weighting, local-safe adaptive budget, rho-threshold sensitivity, and budget target sensitivity.
- Behavioral analysis: how often agents act, whether interventions are local, and whether high survival comes with excessive topology churn.
- Generalization: seen versus unseen chronics and robustness across seeds.

Figures and tables to prepare.

- Table of run families and what each ablation changes.
- Survival curves for baseline, GNN, A0, and sparse-control comparisons.
- Action-frequency plots: action 0 fraction, any non-idle action, and multi-agent simultaneous interventions.
- Summary table with final survival, peak survival, and action sparsity.

Chapter 6: Discussion

Goal. Interpret the results rather than just restating them.

Planned content.

- Answer each research question directly.
- Explain which components appear essential, useful, or unnecessary.
- Discuss the tradeoff between survival and sparse interventions.
- Relate learned behavior to operator-like control: act rarely, act locally, and act when the grid is unsafe.
- Discuss limits: simulator scope, incomplete operational constraints, seed variance, training cost, and dependency on chosen chronics.

Chapter 7: Conclusion and Future Work

Goal. Close with the main thesis claim and next research directions.

Planned content.

- Short summary of the problem, approach, and core findings.
- Final conclusion about decentralized MARL for topology control.
- Future work: richer graph representations, larger grids, outage contingencies, constrained RL tuning, better coordination mechanisms, and operational validation.

1.7 Appendix Plan

- Appendix A: Environment details and action-space construction.
- Appendix B: Hyperparameters and training configuration tables.
- Appendix C: Full ablation run registry.
- Appendix D: Additional survival and behavior plots.
- Appendix E: Reproducibility notes for notebooks, cached W&B histories, and evaluation scripts.

1.8 Immediate Writing Order

1. Write Chapter 3 first, because it fixes notation and the experiment vocabulary.
2. Write Chapter 4 next, using the existing configuration READMEs as source notes for method variants.
3. Draft the Chapter 5 table of ablations before prose, so the results chapter has a stable spine.
4. Write Chapter 1 once the main conclusion is clearer.
5. Finish Chapters 6 and 7 after the final plots and summary metrics are frozen.

1.9 Open Decisions

- Confirm the final thesis title and whether to emphasize GNNs, sparse-control, or decentralized MARL most strongly.
- Decide which experiment families are final and which are only exploratory.
- Decide whether the main claim is performance improvement, improved sparsity at equal survival, or reproducible ablation insight.
- Choose the final bibliography style required by the university.

Chapter 2

Methods

2.1 Graph Representation

The current implementation represents a Grid2Op observation as a homogeneous, directed graph whose nodes correspond to busbars and whose edges describe possible transmission-line connections. Grid2Op exposes the instantaneous electrical state and topology through a structured observation [1]. The graph construction maps this observation to fixed-size tensors suitable for batched graph neural network training. Its main design choice is to keep the set of candidate nodes and edges fixed throughout an episode while using a dynamic edge mask to express the topology that is electrically active at time step t .

For agent a , the graph observation can be written as

$$\mathcal{G}_t^{(a)} = \left(\mathcal{V}^{(a)}, \mathcal{E}_{\text{cand}}^{(a)}, \mathbf{X}_t^{(a)}, \mathbf{Z}_t^{(a)}, \mathbf{m}_t^{(a)} \right), \quad (2.1)$$

where $\mathcal{V}^{(a)}$ and $\mathcal{E}_{\text{cand}}^{(a)}$ are static, $\mathbf{X}_t^{(a)}$ contains node features, $\mathbf{Z}_t^{(a)}$ contains edge features, and $\mathbf{m}_t^{(a)}$ indicates which candidate edges are active. This separates the physical equipment layout, which does not change during an episode, from the busbar assignment and line status, which can change after every topology action.

Figure 2.1 illustrates this baseline schema: busbars are the only nodes, while active transmission lines form the graph backbone.




Figure 2.1: Busbar graph representation. Blue nodes denote busbars and black connections denote active transmission-line relations. Generator and load quantities are aggregated into the features of their associated busbars.

2.1.1 Busbar nodes

Let S denote the number of substations and let B denote the maximum number of busbars per substation. A node is created for every substation–busbar pair, including busbars to which no asset is connected at

the current time. A global node identifier is assigned as

$$i(s, b) = sB + b, \quad s \in \{0, \dots, S-1\}, \quad b \in \{0, \dots, B-1\}. \quad (2.2)$$

The complete graph therefore contains SB nodes. Grid2Op uses one-based busbar identifiers in its topology vector; these identifiers are converted to the zero-based index b in Equation 2.2.

Table 2.1 lists the six node features used in the current representation. Generator and load quantities are assigned to the busbar specified by the current topology vector. Active powers are summed when several assets occupy the same busbar, whereas the corresponding angles are averaged. Disconnected assets do not contribute to any node. The substation cooldown is repeated on every busbar node of that substation.

Table 2.1: Features attached to each busbar node.

Feature	Definition
<code>gen_p</code>	Sum of active generator power connected to the busbar.
<code>gen_theta</code>	Mean voltage angle of generators connected to the busbar.
<code>load_p</code>	Sum of active load power connected to the busbar.
<code>load_theta</code>	Mean voltage angle of loads connected to the busbar.
<code>time_before_cooldown_sub</code>	Remaining substation cooldown.
<code>domain_mask</code>	One for a substation controlled by the agent and zero for a contextual neighbor.

This aggregation produces a compact representation but deliberately removes the identities and counts of individual generators and loads. Consequently, two collections of assets with identical aggregated values have the same node representation.

2.1.2 Candidate line edges and dynamic topology

Consider a physical line ℓ whose origin and extremity belong to substations $o(\ell)$ and $e(\ell)$. For every pair of possible endpoint busbars (b_o, b_e) , the graph reserves the two directed candidate edges

$$i(o(\ell), b_o) \rightarrow i(e(\ell), b_e) \quad \text{and} \quad i(e(\ell), b_e) \rightarrow i(o(\ell), b_o). \quad (2.3)$$

Thus, each physical line contributes $2B^2$ candidate directed edges. The two directions share the same line attributes. For the usual case $B = 2$, one line produces eight candidate edges even though at most two—one per direction—are electrically active at a given time.

The mask for a candidate endpoint assignment is

$$m_{\ell, b_o, b_e, t} = \mathbf{1}[\text{line } \ell \text{ is connected at } t] \mathbf{1}[b_{\ell, t}^{\text{or}} = b_o] \mathbf{1}[b_{\ell, t}^{\text{ex}} = b_e]. \quad (2.4)$$

If a line is disconnected, all of its candidate edges are masked. Otherwise, only the pair corresponding to its current origin and extremity busbars remains active. Inactive edge features are set to zero, and inactive edges are removed before message passing. Topology changes can therefore modify graph connectivity without changing tensor shapes or rebuilding the static graph specification.

The dynamic attributes copied onto active physical-line edges are given in Table 2.2. The two maintenance features are present only for environments with scheduled maintenance.

Table 2.2: Features attached to physical-line edges.

Feature	Definition
<code>line_status</code>	Binary connection status of the physical line.
<code>rho</code>	Line loading relative to its thermal limit.
<code>timestep_overflow</code>	Number of consecutive time steps in overload.
<code>time_before_cooldown_line</code>	Remaining cooldown before another line action is legal.
<code>time_next_maintenance</code>	Time until scheduled maintenance, when available.
<code>duration_next_maintenance</code>	Duration of scheduled maintenance, when available.

2.1.3 Physical scaling and running normalization

Two optional preprocessing stages are implemented for every graph schema. They are controlled independently with

`gnn_physical_scaling = true,      gnn_running_norm = true.`

Both options default to false so that historical experiments and checkpoints retain their original inputs. When enabled together, deterministic physical scaling is applied first and running standardization is applied second. The GNN input is therefore

$$\hat{x}_{i,f,t} = \text{clip} \left(\frac{x_{i,f,t}/s_f - \mu_{f,t}}{\sqrt{\sigma_{f,t}^2 + \epsilon}}, -c, c \right), \quad (2.5)$$

for a continuous feature f with physical scale s_f , running mean $\mu_{f,t}$, running variance $\sigma_{f,t}^2$, and clipping threshold c . The default is $c = 10$, configurable through `gnn_norm_clip`. A non-positive value disables clipping.

The first stage converts physical quantities to dimensionless ranges while preserving a physical zero as zero. Table 2.3 summarizes the implemented factors. The power base can be specified explicitly with `gnn_power_scale_mw`. When its value is non-positive, the largest installed generator rating is used, with a 100 MW fallback when the environment does not expose generator ratings.

Table 2.3: Deterministic physical scales used for continuous graph features.

Features	Scale s_f	Source
<code>gen_p, load_p, p</code>	Power base in MW	Configured value or maximum generator rating.
<code>gen_theta, load_theta, theta</code>	180 degrees	Fixed angular scale.
<code>time_before_cooldown_sub</code>	Maximum substation cooldown	Grid2Op environment parameter.
<code>time_before_cooldown_line</code>	Maximum line cooldown	Grid2Op environment parameter.
<code>timestep_overflow</code>	Allowed overflow duration	Grid2Op environment parameter.
Maintenance time and duration	Episode horizon	Grid2Op chronic horizon.

The loading ratio `rho` is already dimensionless and is not rescaled or standardized. Preserving it is also necessary when GCN uses `rho` directly as a non-negative message weight. Binary connection states, node and edge masks, domain masks, type indicators, and relation one-hot attributes are also left unchanged.

The second stage maintains Welford running statistics over the physically scaled values. One shared normalizer serves every agent. At each environment step, its statistics are updated once from the full state graph rather than from every local graph; duplicated regional observations therefore do not receive extra weight. Heterogeneous graphs maintain a separate node statistic for each node type, preventing the structurally zero entries of busbars, generators, loads, and transmission-line nodes from contaminating one another. Edge statistics include only active physical-line relations. Inactive candidate edges and self, asset-attachment, same-substation, and virtual relations do not contribute.

Each parallel environment worker maintains local online statistics. Before evaluation and checkpointing, the workers are merged using the parallel Welford equations. The merged state is broadcast to evaluation environments, saved in the checkpoint, restored when training resumes, and frozen during evaluation. If running normalization is requested for standalone evaluation but the checkpoint does not contain statistics, the existing normalization policy either raises an error or disables the running stage according to the selected evaluation mode.

This input normalization is distinct from the LayerNorm operations inside the encoder. Equation 2.5 gives a stable meaning to each physical input channel, whereas LayerNorm acts on the learned hidden channels of each individual node or edge after projection. Both stages can therefore remain enabled simultaneously.

2.1.4 Additional relation edges

The graph schema distinguishes three directed edge relations: *self*, *physical line*, and *same-substation busbar*. The same-substation relation can be enabled for every encoder and graph representation with

`gnn_add_substation_edges = true.`

For each substation it adds a directed edge for every ordered pair of distinct busbars. These edges remain active throughout the episode, carry zero-valued physical-line features except for the neutral `rho` message weight, and are distinguished by a categorical edge type. When the general option is left unset, the existing behavior is preserved: same-substation and self relations are enabled by default for the sparse graph transformer and disabled by default for the other encoders.

Same-substation edges are computational attention relations; they do not imply that the two busbars are electrically connected. They allow information to pass between alternative busbar states within a substation, while edge-type embeddings and relation-specific attention biases allow the encoder to treat them differently from transmission lines. In the heterogeneous schemas, GAT and GINE also receive the relation one-hot attribute. GCN and GraphSAGE use the additional connectivity and identify the endpoints through their node features.




Figure 2.2: Direct same-substation busbar relation. The highlighted orange edge allows two busbars belonging to the same substation to exchange messages; it is a computational relation and not a physical transmission line.

For a full graph containing S substations and L physical lines, the number of candidate directed edges is

$$|\mathcal{E}_{\text{cand}}| = 2B^2L \quad (2.6)$$

for the standard representation, and

$$|\mathcal{E}_{\text{cand}}| = 2B^2L + SB + SB(B - 1) \quad (2.7)$$

when self and same-substation relations are both enabled. As an illustration, a 14-substation, 20-line system with two busbars has 28 nodes and 160 physical candidate edges. Adding both relation types gives 216 candidate edges in total.

2.1.5 Local actor graphs and the global state graph

Each agent controls a set of substations $\mathcal{C}_a$. Under decentralized observability, its local graph can be constructed in two ways. Without neighbor context, the node set contains only substations in $\mathcal{C}_a$, and the line set contains only lines whose two endpoints lie inside this region. With one-hop context enabled, the graph additionally contains every line touching $\mathcal{C}_a$ and the substations at the opposite endpoints of those lines. The `domain_mask` node feature and a separate controlled-node mask distinguish the controlled region from contextual neighbors. This second mask also permits graph readout to pool only the nodes on which the agent can act.

Alongside every local graph, the environment constructs a full-grid *state graph*. A graph-based actor consumes its agent-specific local graph, whereas a graph-based centralized critic consumes the state graph.

This follows the centralized-training, decentralized-execution structure of MAPPO [2]: actors can be restricted to regional information while the critic uses a global representation during training.

The observation delivered to each agent has the conceptual structure

$$\{\text{flat}, \text{graph}, \text{state_graph}\}. \quad (2.8)$$

The local and state graphs each contain node features, edge features, node and edge masks, edge types, and a controlled-node mask. Their static specifications, including node identifiers and candidate edge indices, are built once when the environment is initialized.

2.1.6 Boundary between representation and encoder

The busbar graph is shared by several encoder architectures. GAT and GINE use the edge attributes directly; GCN can optionally use `rho` as an edge weight; GraphSAGE uses the active connectivity without edge attributes; and the sparse graph transformer can use both continuous edge attributes and categorical edge relations. Optional learned substation and busbar embeddings may be concatenated to the node features. After message passing, the node embeddings are pooled into one fixed-dimensional graph embedding. This embedding can be used alone or concatenated with the original flat observation before it enters the actor or critic head.

These choices change how the graph is encoded, but not which physical entities the graph represents. The busbar schema does not create separate generator, load, line, or substation nodes, and it does not encode reactive power, voltage magnitudes, or directional endpoint flows. Empty busbars remain in the node set, and the node mask is currently one for every allocated node. Moreover, graph features are taken from the raw Grid2Op observation by default. The optional physical scaling and shared running normalization from Section 2.1.3 can be enabled without changing the graph schema. These properties define the busbar baseline against which the heterogeneous representation introduced next can be compared.

2.2 Heterogeneous Graph Representation

The second representation makes the equipment categories explicit instead of aggregating generators and loads into busbar features. It follows the structure illustrated conceptually by a network in which busbars form the blue backbone, generators are red terminal nodes, and loads are green terminal nodes. The graph contains three node types,

$$\mathcal{V} = \mathcal{V}_{\text{bus}} \cup \mathcal{V}_{\text{gen}} \cup \mathcal{V}_{\text{load}}, \quad (2.9)$$

and typed directed relations between them. This preserves the identity and state of every generator and load while retaining the busbar-level description of the controllable topology.

The implementation is selected with `gnn_graph_type = heterogeneous`. It uses the same local actor and global state-graph interface as the busbar representation, so the graph schema can be changed without modifying MAPPO rollout collection or the actor and critic APIs.

2.2.1 Typed nodes and features

For a graph with S included substations, B busbars per substation, N_g included generators, and N_d included loads, the number of nodes is

$$|\mathcal{V}| = SB + N_g + N_d. \quad (2.10)$$

Busbar nodes are created exactly as in Section 2.1.1. In addition, each generator and load whose substation is present in the graph receives its own node. Asset nodes remain allocated when an asset is disconnected, thereby preserving fixed tensor shapes throughout an episode.

All node types are stored in one feature matrix. Type-specific values occupy a shared feature schema, with unused entries set to zero. Table 2.4 lists this schema.




Figure 2.3: Heterogeneous equipment representation. Busbars are blue, generators are red, and loads are green. Physical transmission lines remain direct relations between busbar nodes.

Table 2.4: Unified node features in the heterogeneous representation.

Feature	Definition
<code>is_busbar</code>	One for busbar nodes and zero otherwise.
<code>is_generator</code>	One for generator nodes and zero otherwise.
<code>is_load</code>	One for load nodes and zero otherwise.
<code>p</code>	Generator active production or load active consumption.
<code>theta</code>	Generator or load voltage angle.
<code>time_before_cooldown_sub</code>	Substation cooldown on busbar nodes.
<code>connected</code>	One for busbars and for currently connected assets; zero for a disconnected asset.
<code>domain_mask</code>	One when the node belongs to a substation controlled by the agent.

When an asset is disconnected, its `connected` feature becomes zero and its dynamic `p` and `theta` entries are zeroed. Its type and substation identity remain available. Unlike the busbar representation, generator and load values are not summed or averaged, so two assets connected to the same busbar can retain different states.

2.2.2 Typed relations

The heterogeneous graph uses seven possible directed edge relations:

1. *self*;
2. *physical line*, between two busbars;
3. *same-substation busbar*;
4. *generator-to-busbar*;
5. *busbar-to-generator*;
6. *load-to-busbar*;
7. *busbar-to-load*.

Direction-specific asset relations allow an encoder to learn different transformations for information entering and leaving a generator or load. Self and same-substation relations remain optional and are enabled by default for the sparse graph transformer, as in the busbar representation.

Physical transmission lines use the candidate endpoint construction from Equation 2.3. Asset attachments follow the same fixed-shape principle. For generator g located at substation $s(g)$, the candidate set contains

$$g \rightarrow i(s(g), b) \quad \text{and} \quad i(s(g), b) \rightarrow g, \quad b \in \{0, \dots, B-1\}. \quad (2.11)$$

Only the two directed edges associated with the generator’s current busbar are active. Loads are handled identically. If an asset is disconnected in the topology vector, all of its attachment edges are masked. A topology action can therefore move an asset from one busbar to another by changing masks, while the candidate edge index remains static.

Without optional self and same-substation relations, the candidate edge count is

$$|\mathcal{E}_{\text{cand}}| = 2B^2L + 2B(N_g + N_d), \quad (2.12)$$

where L is the number of included physical lines. Enabling both optional relations adds $|\mathcal{V}|$ self-edges and $SB(B-1)$ directed same-substation busbar edges.

2.2.3 Tensor representation and encoder compatibility

Semantically, the graph is heterogeneous because its nodes and relations have different physical meanings. Operationally, it is represented using unified tensors rather than separate PyTorch Geometric node stores. The observation adds a categorical `node_type` tensor and retains the categorical `edge_type` tensor. The three node-type indicators are also included in the node feature matrix. Similarly, a one-hot relation block is appended to the continuous edge attributes. This dual representation makes the schema usable by the existing encoders:

- GAT and GINE receive the relation one-hot values as edge attributes;
- the sparse graph transformer receives the integer relation identifiers through its relation embeddings and attention biases;
- GCN and GraphSAGE infer equipment roles from the node-type features, even though their standard layers do not directly consume edge types.

The physical-line part of the edge feature vector retains the attributes from Table 2.2. Structural relations use a neutral unit value in the `rho` position so they remain effective when GCN is configured to use `rho` as a message weight; this value is a computational weight rather than a physical loading ratio. Inactive candidate relations are removed before message passing, exactly as for physical lines.

2.2.4 Local heterogeneous graphs

The decentralized domain rules remain unchanged. A local graph first selects its controlled substations and, optionally, their one-hop neighboring substations. It then includes the busbars, generators, and loads belonging to those substations, together with the selected physical lines. Consequently, neighbor assets are available as context when one-hop expansion is enabled, but the controlled-node mask still distinguishes them from equipment inside the agent’s action domain. The centralized state graph contains all busbars, generators, loads, physical lines, and active attachment relations in the grid.

The principal trade-off is therefore explicitness against size. The busbar representation is smaller but aggregates equipment values. The heterogeneous representation introduces $N_g + N_d$ nodes and $2B(N_g + N_d)$ candidate attachment edges, but preserves individual asset states and gives the encoder typed paths that mirror the intended physical structure.

2.3 Heterogeneous Graph with Transmission-Line Nodes

The third representation extends Section 2.2 by representing every transmission line as a node rather than as a direct edge between two busbars. In the conceptual diagram, busbars are blue, generators are red, loads are green, and transmission lines are yellow. The line nodes form the intermediate elements through which information passes between the busbars at the two ends of each physical line. The resulting node set is

$$\mathcal{V} = \mathcal{V}_{\text{bus}} \cup \mathcal{V}_{\text{gen}} \cup \mathcal{V}_{\text{load}} \cup \mathcal{V}_{\text{line}}. \quad (2.13)$$

This schema is selected with `gnn_graph_type = heterogeneous_line`. It implements the same fixed tensor interface used by the two preceding representations and can therefore be used by the existing GNN actor and critic encoders without changing the MAPPO training loop.




Figure 2.4: Heterogeneous representation with explicit transmission-line nodes. Yellow nodes represent physical lines and connect the blue busbars at their endpoints; generators and loads remain explicit terminal nodes.

2.3.1 Node allocation and features

For S included substations, B busbars per substation, N_g generators, N_d loads, and L included transmission lines, the graph contains

$$|\mathcal{V}| = SB + N_g + N_d + L \quad (2.14)$$

nodes. Busbar, generator, and load nodes retain the features described in Table 2.4. A fourth type indicator, `is_transmission_line`, identifies line nodes. The unified node-feature matrix additionally provides the line-specific state shown in Table 2.5; entries that do not apply to a given node type are set to zero.

Table 2.5: Features carried by an explicit transmission-line node.

Feature	Definition
<code>line_status</code>	One when the physical line is connected and zero otherwise.
<code>rho</code>	Line loading ratio.
<code>timestep_overflow</code>	Number of consecutive overflow time steps.
<code>time_before_cooldown_line</code>	Remaining line-reconnection cooldown.
<code>time_next_maintenance</code>	Time until planned maintenance, when maintenance features are enabled.
<code>duration_next_maintenance</code>	Planned maintenance duration, when maintenance features are enabled.
<code>connected</code>	One for an online line and zero for a disconnected line.
<code>domain_mask</code>	One when either endpoint belongs to a substation controlled by the agent.

A disconnected line remains allocated in the graph, so neither the node count nor the candidate edge index changes during an episode. Its `line_status` and `connected` entries become zero and all of its busbar attachments are masked. Other state, such as its cooldown and maintenance timing, remains visible to the encoder. This distinction lets the policy reason about a temporarily unavailable physical element without mistaking it for an element that does not exist.

2.3.2 Endpoint-specific typed relations

Let $o(\ell)$ and $e(\ell)$ denote the origin and extremity substations of line ℓ . For every busbar index b , the fixed candidate graph contains

$$i(o(\ell), b) \leftrightarrow \ell \quad \text{and} \quad i(e(\ell), b) \leftrightarrow \ell, \quad b \in \{0, \dots, B-1\}. \quad (2.15)$$

The bidirectional symbol denotes two directed relations. At a particular time, the topology vector activates only the origin and extremity busbars to which the line is currently attached. Thus, an online line has four active directed attachment edges: busbar-to-line and line-to-busbar at each endpoint. A topology action only changes the edge mask; it does not rebuild the candidate graph.

The representation distinguishes ten relation types:

1. self;
2. same-substation busbar;
3. busbar-to-line at the origin;
4. line-to-busbar at the origin;
5. busbar-to-line at the extremity;
6. line-to-busbar at the extremity;
7. generator-to-busbar;
8. busbar-to-generator;
9. load-to-busbar;
10. busbar-to-load.

Separating direction and line endpoint allows a relation-aware encoder to learn different transformations for the two terminals of a transmission line. The generator and load attachment rules are unchanged from Equation 2.11. Self and same-substation busbar relations remain optional.

Without the optional relations, the number of candidate directed edges is

$$|\mathcal{E}_{\text{cand}}| = 4BL + 2B(N_g + N_d). \quad (2.16)$$

For the common case $B = 2$, the line-attachment term has the same size as the $2B^2L$ direct busbar-to-busbar candidate set in Equation 2.12. For $B > 2$, explicit line nodes make this term grow linearly rather than quadratically with the number of busbars.

2.3.3 Tensor representation and local graphs

As in the preceding heterogeneous representation, all node types share one node-feature tensor. The categorical `node_type` tensor has four values, and the type indicators are also present in the node features. The categorical `edge_type` tensor identifies the ten relations, while a corresponding one-hot relation block is appended to the edge features for encoders such as GAT and GINE. The line loading ratio is copied to active line-attachment edges so it can also be used as a GCN message weight; non-line structural relations receive a neutral unit weight. The complete physical line state remains available principally on the line node.

Local actor graphs use the same domain and one-hop-neighbor selection as the other schemas. An included line receives one line node and its candidate attachments to the included busbars. A line node is marked as controlled when at least one of its endpoints belongs to the agent’s controlled domain. The centralized state graph contains all busbar, generator, load, and transmission line nodes. Compared with the direct-edge heterogeneous graph, this representation adds L nodes and an extra message-passing step between the two endpoint busbars, but makes each transmission line an identifiable entity whose state can be updated and embedded independently.

2.4 Substation Summary Nodes

An optional hierarchical augmentation introduces one learned node for every substation included in the graph. It is enabled with

`gnn_add_substation_nodes = true.`

For substation s , let u_s denote its summary node and let $\mathcal{B}(s)$ be its included busbar nodes. The augmented graph contains the bidirectional structural relations

$$\mathcal{E}_{\text{sub}} = \{(i, u_s), (u_s, i) \mid i \in \mathcal{B}(s)\}. \quad (2.17)$$

Each substation node therefore gathers information from every busbar belonging to that substation and can redistribute its learned summary to those busbars in later message-passing layers. It is not connected directly to generators, loads, or transmission-line nodes; those elements influence it through their busbar attachments.




Figure 2.5: Learned substation-summary augmentation. The orange substation node receives messages from each of its blue busbar nodes and can redistribute the resulting substation-level representation.

The initial substation-node states are learned embeddings indexed by the global substation identifiers. For S included substations and B busbars per substation, the augmentation adds S nodes and $2SB$ directed edges. The nodes are constructed inside the encoder, leaving the environment observation and rollout-storage shapes unchanged. Standard mean, sum, maximum, or attention readout includes the new summary nodes. Controlled readout marks a substation node as controlled whenever at least one of its busbars is controlled. The sparse graph transformer assigns the attachments a dedicated learned relation.

Substation nodes can also be combined with the virtual-node readout introduced next. In that case, information follows a hierarchy from physical equipment to busbars, from busbars to substation nodes, and from substation nodes to one graph-level virtual node. The node augmentation is distinct from the direct same-substation edges in Section 2.1.4: the former introduces a learned hub u_s , whereas the latter directly connects pairs of busbars. Either mechanism, or both, may be enabled.

2.5 Virtual-Node Graph Readout

The preceding graph representations normally produce a fixed-dimensional graph embedding by applying mean, sum, maximum, or attention pooling after the final message-passing layer. An alternative readout is implemented using one *virtual node*. This node acts as a learned summary of all substations in the current graph and replaces the terminal pooling operation. The option is selected with

`gnn_readout_aggr = virtual_node.`

It is independent of the graph representation and can therefore be combined with the busbar, heterogeneous, or heterogeneous-line schema.

2.5.1 Augmented graph

Let $v_\star$ denote the virtual node. Define the summary set $\mathcal{V}_{\text{sum}}$ as the included substation nodes when the augmentation from Section 2.4 is enabled, and as the included busbar nodes otherwise. At encoder runtime, the graph is augmented as

$$\mathcal{V}^+ = \mathcal{V} \cup \{v_\star\}, \quad (2.18)$$

$$\mathcal{E}^+ = \mathcal{E} \cup \{(i, v_\star), (v_\star, i) \mid i \in \mathcal{V}_{\text{sum}}\}. \quad (2.19)$$

Without substation summary nodes, the virtual node is connected bidirectionally to every included busbar, as illustrated by the purple links in Figure 2.6. With the hierarchical augmentation enabled, the virtual node instead connects to every included substation node. It is never connected directly to generator, load, or transmission-line nodes. Those entities nevertheless influence it through their ordinary paths to the busbar backbone. Masked busbar nodes do not receive hierarchy connections.




Figure 2.6: Virtual-node graph readout. The purple node exchanges messages with every included busbar and its final embedding replaces terminal node pooling. When substation-summary nodes are enabled, the virtual node connects to those summaries instead.

The virtual node is appended inside the encoder rather than the environment. Consequently, the Grid2Op observation tensors, static graph specifications, and rollout storage retain their original dimensions. For a minibatch of M graphs, the encoder creates M independent copies of the virtual node and does not introduce edges between different graphs.

2.5.2 Learned summary instead of terminal pooling

The initial virtual representation $\mathbf{h}_\star^{(0)}$ is a learned parameter shared across input graphs. Each graph receives a copy of this parameter. Because the virtual edges are present during every layer, a generic message-passing update jointly evolves the physical and virtual nodes:

$$\left(\{\mathbf{h}_i^{(k+1)}\}_{i \in \mathcal{V}}, \mathbf{h}_\star^{(k+1)}\right) = \text{GNN}^{(k)}\left(\{\mathbf{h}_i^{(k)}\}_{i \in \mathcal{V}}, \mathbf{h}_\star^{(k)}, \mathcal{E}^+\right). \quad (2.20)$$

Incoming busbar or substation messages let $v_\star$ accumulate a graph-level summary, while the reverse edges let its current state influence the hierarchy in subsequent layers. With K layers, the graph embedding is obtained directly from the final virtual state,

$$\mathbf{z}_G = \text{ReLU}\left(\mathbf{W}_o \mathbf{h}_\star^{(K)} + \mathbf{b}_o\right), \quad (2.21)$$

rather than from $\text{Pool}(\{\mathbf{h}_i^{(K)}\}_{i \in \mathcal{V}})$. The output dimension and the actor–critic interface are therefore unchanged.

2.5.3 Relation attributes and encoder compatibility

The readout is supported by GCN, GAT, GINE, GraphSAGE, and the sparse graph transformer. Virtual connections are structural rather than physical: their continuous edge attributes are zero, except for a neutral unit value in the `rho` position when that feature exists. The unit value ensures that a GCN configured to use `rho` as an edge weight does not suppress virtual messages. The sparse graph transformer allocates dedicated learned relations for substation attachments and virtual-node attachments, allowing its attention layers to distinguish hierarchy links from all physical and equipment relations.

Without substation nodes, a graph containing N_{bus} included busbars receives one virtual node and $2N_{\text{bus}}$ directed edges. With the hierarchy enabled, the virtual-node stage instead adds $2S$ edges on top of the substation augmentation. The method avoids a fixed, parameter-free reduction such as mean or sum pooling, but its summary quality depends on message-passing depth: after the first layer it has aggregated its immediate busbar or substation neighbors, and later layers permit repeated interaction between the summary and the physical graph. In a decentralized actor, the virtual node summarizes the controlled domain together with any included one-hop context; in the centralized critic, it summarizes the full state graph.

2.6 Staged Graph-Architecture Screening Protocol

The representation, preprocessing, hierarchy, readout, and encoder choices defined above form a large factorial design. Even when graph readout is limited to mean pooling and the virtual node, a simultaneous sweep over three graph schemas, four input-preprocessing modes, three binary structural factors, and five encoders would require

$$3 \times 4 \times 2^3 \times 5 = 480 \quad (2.22)$$

conditions before adding random seeds. The implemented screening therefore uses sequential promotion: each stage varies one group of factors and inherits the selected setting from the preceding stage. This reduces the initial screen to 25 runs while retaining an explicit test of every proposed mechanism.

2.6.1 Common screening controls

All screening runs first use the `bus14` environment, the same action space, the same train–test chronic split, training seed zero, 72 parallel environments, and rollouts of 576 steps. A run receives 8 million environment steps, slightly more than half of the 15-million-step `bus14` confirmation budget. Evaluation is deterministic and uses the same held-out chronic set in every condition. The smaller network makes the 25-run screen affordable before the selected design is scaled to WCCI.

The actor uses a 128-dimensional, two-layer graph encoder followed by the same action head in every condition. Two message-passing layers are necessary for a fair evaluation of the explicit-line schema: information can travel from a busbar to a line node and then to the opposite busbar. The critic is a fixed MLP, so the screen measures changes to the actor graph encoder rather than changes to the centralized value function. The critic is used only during centralized training and is not part of decentralized actor execution. Replacing it simultaneously with a GNN would introduce a second representation change, increase computation, and make an improvement difficult to attribute. Its flat input remains normalized through `norm_obs`; this is independent of graph input normalization. A GNN critic is reserved for a subsequent ablation using the winning actor configuration. It is particularly relevant if the complete actor–critic training system, rather than only the shared actor encoder, must support different grid sizes.

To preserve the intended grid-size generalization, one actor GNN is shared by all agents, learned global node-identifier embeddings are disabled, and the flat observation is not concatenated to the actor embedding. One-hop neighboring context is enabled. LayerNorm remains enabled inside every GNN layer and is not treated as a screening factor. Thus the external preprocessing factors described in Section 2.1.3 are varied without also changing hidden-state normalization.

2.6.2 Stage 1: representation and input preprocessing

Stage 1 fixes the encoder to GINE, uses mean graph pooling, and disables the three structural additions. GINE is used because it consumes both node and physical-line edge attributes. The three graph schemas

are crossed with the four preprocessing conditions in Table 2.6, giving 12 paired runs.

Table 2.6: Stage-1 graph input preprocessing modes applied to each of the busbar, heterogeneous-equipment, and heterogeneous-line schemas.

Code	Physical scaling	Running normalization	Interpretation
n0_none	Off	Off	Raw graph input baseline
n1_physical	On	Off	Dimensionless physical ranges
n2_running	Off	On	Online standardized input
n3_both	On	On	Physical scaling, then standardization

The physical scale and running moments are expected to have partially redundant effects after the moments converge: standardization cancels a positive constant scale apart from initialization, numerical epsilon, and clipping. The fourth condition is nevertheless retained to test whether physically bounded early inputs improve optimization before the running estimates stabilize.

The primary stage-1 selection criterion is held-out survival and chronic completion. Ties are resolved using return, overload frequency and maximum ρ , topology-action count, learning stability, wall-clock time, and memory consumption. The selected graph schema and preprocessing pair are then frozen for Stage 2.

2.6.3 Stage 2: structural graph factors

Stage 2 inherits the Stage-1 winner and evaluates all eight binary combinations shown in Table 2.7. Here E denotes direct same-substation busbar edges, N denotes one learned substation summary node per included substation, and V replaces terminal mean pooling with virtual-node readout. When $V = 0$, mean pooling is used. When $V = 1$, the readout is `virtual_node`.

Table 2.7: The complete Stage-2 structural screen.

Code	E	N	V	Readout	Added mechanism
e0n0v0	0	0	0	Mean	Continuity rerun of Stage-1 winner
e0n0v1	0	0	1	Virtual	Global learned summary
e0n1v0	0	1	0	Mean	Substation hierarchy
e0n1v1	0	1	1	Virtual	Hierarchy and learned summary
e1n0v0	1	0	0	Mean	Direct same-substation communication
e1n0v1	1	0	1	Virtual	Direct edges and learned summary
e1n1v0	1	1	0	Mean	Direct edges and hierarchy
e1n1v1	1	1	1	Virtual	Complete structural augmentation

The e0n0v0 condition deliberately repeats the promoted condition under a new run identifier and provides a stage-boundary consistency check. The current summary-node initialization contains a learned embedding indexed by substation identifier. Results involving $N = 1$ therefore measure in-grid performance, but are not used alone as evidence of zero-shot transfer to an unseen grid. Direct substation edges and the single shared virtual node do not require a table of grid-specific identifiers.

2.6.4 Stage 3: encoder comparison

Stage 3 freezes the selected representation, preprocessing, structural factors, hidden dimension, depth, actor head, and critic. It compares GCN, GAT, GINE, GraphSAGE, and the relation-aware sparse graph transformer. Four attention heads are configured; this option is ignored by encoders that do not use multi-head attention. The GINE condition repeats the Stage-2 winner as a consistency run.

Only mean and virtual-node readout are eligible for promotion because both are implemented by every encoder. Attention and controlled-attention pooling are valid sparse-transformer-specific studies, but including them here would confound the encoder comparison with a change in readout. The configuration launcher now rejects those readouts early when they are paired with an ordinary GNN instead of allowing a later model-construction failure.

2.6.5 Promotion, confirmation, and reproducibility

The committed configuration directory is `configs/gnn_graph_screening`. Stage 1 contains the 12 executable TOML files, a manifest of the varied factors, and a batch launcher. Because the later configurations must contain the actual selected winner, they are generated only after promotion. In the commands below, `screen.py` abbreviates `tools/gnn_graph_screening.py`. From the `Topology_Task` directory, the procedure is:

1. execute `bash launch_all.sh` in the Stage-1 directory;
2. generate Stage 2 with `screen.py promote-structure WINNER.toml`;
3. generate Stage 3 with `screen.py promote-encoder WINNER.toml`;
4. validate generated files with `screen.py validate DIRECTORY`;
5. promote selected configurations with `screen.py confirm SELECTED.toml`.

Each promotion writes runnable TOML files, `manifest.csv`, and `launch_all.sh` into the corresponding stage directory. Existing changed files are not overwritten unless the explicit `-force` option is supplied.

The confirmation command creates fresh 15-million-step bus14 runs for seeds 0, 1, and 2. Up to four screened configurations are retained. Final claims use the mean and dispersion across these seeds and deterministic evaluation over the same complete held-out chronic set. In addition to survival and return, the report records overload statistics, topology interventions, runtime, memory, and transfer performance on a larger grid. The one-seed stages are therefore a resource-efficient ranking procedure, whereas the three-seed stage supplies the evidence used for the final bus14 comparison. The selected shared actor encoder is then retrained on WCCI with new agent-specific action heads, because the WCCI action spaces do not match those of bus14.

Chapter 3

Experimental Results

3.1 How to Read This Chapter

This chapter is organized as a sequence of experiment blocks. The goal is to make the thesis easy to navigate while many tests are still being added. Each block follows the same structure:

1. **Question.** What uncertainty or design choice does this test address?
2. **Baseline and change.** Which run is the control, and what is changed in the comparison runs?
3. **Protocol.** Which data split, checkpoint, seeds, and metric are used?
4. **Implementation detail.** Which algorithmic, architectural, or mathematical change is being tested, when this is relevant?
5. **Evidence.** Which table and figure summarize the result?
6. **Interpretation.** What can be concluded, and what still needs to be verified?

Using the same block structure for every new experiment should make it possible to add future MSc-thesis tests without rewriting the whole results chapter.

3.2 Evaluation Protocol

The experiments use a train/test chronic split. The training set contains 80% of the chronics and is used to update the policy. The test set contains the remaining 20% and is used to estimate generalization during training and for final evaluation.

During training, evaluation rollouts are periodically run on both the train and test splits using 10 episodes each time. In the current rerun notebook, the main reported metric is **test episodic survival**, expressed as a percentage. Higher survival means that the agent keeps the grid alive for a larger fraction of the evaluation episodes. The final plots show seed-aggregated mean curves, faint individual seed curves, and the variability band.

The final evaluation currently uses the saved checkpoint with the best combined train/test survival. This selection rule should be reported explicitly because the test split participates in checkpoint selection; the values in this chapter are therefore best interpreted as controlled rerun comparisons, not as a final untouched test-set estimate.

3.3 Experiment Index

Table 3.1 gives a compact map of the experiment blocks included so far. New result sections should add one row to this index before adding detailed prose.

Table 3.1: Current result blocks and their role in the thesis.

Block	Main question	Primary evidence	Status
A0 core reruns	Which baseline training changes stabilize MAPPO?	Table 3.3, Figures 3.1 and 3.2	Drafted
Initial do-nothing bias	Does an initial action-0 prior help exploration?	Table 3.4, Figure 3.3	Drafted
GINE reruns	Do graph encoders improve survival over the A0 baseline?	Tables 3.5 and 3.6, Figure 3.4	Drafted
Heavy GINE	Does a larger GINE architecture help or overcomplicate optimization?	Table 3.7, Figure 3.5	Drafted
Sparse intervention control	Can decentralized agents intervene rarely, locally, and only when useful?	Tables 3.9 and 3.10	Design drafted

3.4 IEEE 14-Bus Experiments

The first result group focuses on the IEEE 14-bus topology-control task. The initial PPO/MAPPO baseline from the MARL2Grid framework was not satisfactory: the agents did not learn a robust policy and the survival rate remained low. The following experiments therefore test training and architecture changes that could make the baseline more stable.

3.4.1 A0 Core Reruns

Question. Which changes are needed to obtain a stable MLP MAPPO baseline on the 14-bus task?

Baseline and changes. The A0 family starts from the flat MLP MAPPO baseline and tests three main stabilization choices:

- increasing actor and critic capacity from two hidden layers to three hidden layers;
- using reward normalization;
- changing the critic update so that the critic is updated once for the whole batch instead of once per agent.

Implementation detail. There is no new mathematical objective in this ablation. The implementation changes are practical training choices: network depth, reward normalization, critic-update scheduling, and the initialization of the action-0 probability. Table 3.2 records these settings so that the survival plots can be interpreted without repeatedly restating the run names.

Evidence. Table 3.3 reports final and peak survival for the core A0 ablations. Figure 3.1 shows each comparison against the `rerun_a0_opt` baseline, and Figure 3.2 overlays all core A0 curves.

Interpretation. The optimized A0 baseline, the non-optimized critic-update variant, and the smaller MLP variant all finish around 93% survival. This suggests that the main A0 recipe is robust once the rerun protocol is fixed. Disabling reward normalization still reaches high peak survival, but it learns more slowly and finishes lower. The old no-validation, non-optimized baseline is clearly not competitive and should not be used as the main reference for later ablations.

Table 3.2: A0-family run configurations. The reference run is `rerun_a0_opt`.

Run family	Actor/critic hidden layers	Reward norm.	Initial action-0 prob.	Optimized critic update	Difference from <code>rerun_a0_opt</code>
<code>rerun_a0_opt</code>	$3 \times 256 / 3 \times 256$	true	0.0	true	Reference baseline
<code>rerun_a0_nono</code>	$3 \times 256 / 3 \times 256$	true	0.0	false	Disables optimized critic updates
<code>rerun_a01</code>	$2 \times 256 / 2 \times 256$	true	0.0	true	Uses a smaller MLP actor and critic
<code>rerun_a02</code>	$3 \times 256 / 3 \times 256$	false	0.0	true	Disables reward normalization
<code>old_a0_nono</code>	$3 \times 256 / 3 \times 256$	true	0.3	false	Uses action-0 bias 0.3 and non-optimized critic updates
<code>rerun_bias_05</code>	$3 \times 256 / 3 \times 256$	true	0.5	true	Uses action-0 bias 0.5
<code>rerun_bias_07</code>	$3 \times 256 / 3 \times 256$	true	0.7	true	Uses action-0 bias 0.7

A0 core reruns: `rerun_a0_opt` baseline pairwise comparisons




Figure 3.1: Pairwise A0 rerun comparisons against the `rerun_a0_opt` baseline.

A0 core reruns: rerun_a0_opt vs rerun_a0_nonopt / rerun_a01 / rerun_a02 / old noval20 nonopt




Figure 3.2: All A0 core rerun survival curves shown on one axis.

Table 3.3: A0 core rerun survival summary.

Run	Intended change	Final survival (%)	Peak survival (%)
rerun_a0_opt	A0 optimized baseline	93.4	98.6
rerun_a0_nonopt	Disable optimized critic update	93.1	95.8
rerun_a01	Smaller MLP actor/critic	93.7	100.0
rerun_a02	Disable reward normalization	81.8	97.1
old_a0_nonopt_noval20	Older non-optimized protocol	22.3	68.4

3.4.2 Initial Do-Nothing Bias

Question. Does biasing the initial policy toward action 0, the do-nothing action, help or hurt learning?

Baseline and changes. The A0 baseline uses no initial do-nothing bias in this rerun set. The bias controls keep the same general protocol but change the initial probability mass assigned to action 0.

Implementation detail. The bias is applied only at initialization: it changes the starting action distribution before learning, without changing the reward, architecture, or PPO objective. This makes the ablation a test of exploration bias rather than a test of representational capacity.

Evidence. Table 3.4 and Figure 3.3 compare the zero-bias A0 baseline with rerun_bias_05 and rerun_bias_07.

Interpretation. The effect of the action-0 bias is large and non-monotonic. The 0.5-bias curve remains unstable and substantially below the zero-bias baseline, whereas the 0.7-bias curve eventually reaches the highest final survival in this comparison. The current conclusion is therefore not that larger bias is always better. Instead, the initialization prior appears to strongly shape exploration and should be treated as a first-order ablation variable.

Table 3.4: Initial do-nothing bias survival summary.

Run	Final survival (%)	Peak survival (%)
<code>rerun_a0_opt</code>	93.4	98.6
<code>rerun_bias_05</code>	50.8	70.2
<code>rerun_bias_07</code>	99.3	100.0

`rerun_a0_opt` baseline vs `rerun_bias_05` / `rerun_bias_07`




Figure 3.3: A0 baseline compared with initial do-nothing bias controls `rerun_bias_05` and `rerun_bias_07`.

3.4.3 GINE Encoder Reruns

Question. Can graph encoders improve topology-control performance, and which GINE design choices matter most?

Baseline and changes. The comparison starts from the historical `best_00` GINE baseline. The reruns then test lighter GINE encoders, MLP versus GNN critics, optimized critic updates, flat-feature concatenation, A0-style entropy, and different do-nothing-bias settings.

Protocol. All GINE runs use `gnn_type = gine`, reward normalization, 15,000,000 total timesteps, evaluation every 80,000 timesteps, 2,000 rollout steps, deterministic evaluation, and the same 80/20 chronic split used by the A0 experiments. The reference run is `best_00`.

Implementation detail. The full GINE model architecture is described separately in Appendix A. This block separates representational changes from optimization changes. The GINE encoder changes how grid observations are embedded, while the critic type, flat-feature concatenation, entropy setting, and critic-update schedule affect how the PPO objective is optimized.

Evidence. Table 3.5 records the run configurations, Table 3.6 gives the selected survival values, and Figure 3.4 shows the pairwise survival curves against `best_00`.

Interpretation. The strongest GINE result is `best_14`, which reaches about 95.7% final survival. This suggests that the graph encoder can be competitive when it is paired with a simpler MLP critic and stable A0-style training choices. The heavier historical GINE baseline has high peak survival but much weaker final survival, which points to optimization stability rather than representational capacity as the main issue.

Table 3.5: GINE rerun configuration summary. The reference run is **best_00**.

Run	Actor/critic encoder	GNN hidden/layers	Actor/critic heads	Concat flat	Edge pre-enc.	Entropy	Init action-0	Opt. critic	Difference from best_00
best_00	gnn / gnn	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.7	false	Baseline: heavy concat-flat GINE actor, GNN critic, and legacy critic updates
best_10	gnn / mlp	64 / 1	$2 \times 128 / 2 \times 128$	true	false	0.02	0.7	true	Light GINE actor, MLP critic, smaller heads, no edge pre-encoder, optimized critic updates
best_11	gnn / gnn	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.0	false	Bias ablation: same as baseline, but the initial do-nothing prior is 0.0
best_12	gnn / gnn	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.02	0.0	true	Bias 0.0 plus optimized critic updates; otherwise heavy concat-flat GINE with GNN critic
best_13	gnn / gnn	128 / 2	$3 \times 256 / 3 \times 256$	true	true	0.01	0.0	true	A0-style heavy GINE: bias 0.0, lower entropy 0.01, optimized critic updates
best_14	gnn / mlp	64 / 1	$2 \times 128 / 2 \times 128$	false	false	0.01	0.0	true	Light A0-style GINE: MLP critic, no concat-flat, lower entropy 0.01, bias 0.0, optimized critic updates

Table 3.6: GINE rerun survival summary.

Run	Intended change	Final survival (%)	Peak survival (%)
best_00	Historical GINE baseline	59.2	93.9
best_10	Light GINE with MLP critic	78.9	90.4
best_12	Optimized critic and bias 0.0	82.5	85.7
best_14	Light A0-style GINE with MLP critic	95.7	97.4

configs/gine_s0_s1_s2: best_00 vs best_10-14




Figure 3.4: Seed-aggregated test episodic survival for GINE best_00 against best_10–best_14.

3.4.4 Heavy GINE

Question. Does the heavier concat-flat GINE baseline fail because it lacks capacity, or because specific architectural and optimization choices make it unstable?

Baseline and changes. The reference run is best_00, the historical shared-actor, concat-flat GINE with a GNN critic and legacy critic updates. The best-search 2 family compares this baseline against controlled alternatives that change the flat-feature concatenation, critic type, actor sharing, GNN size, entropy schedule, node identifiers, message-passing operator, and initialization bias.

Protocol. The configurations are read from configs/gine_s0_s1_s2. Each family is run with seeds s0, s1, and s2. Unless changed by the family, the runs use reward normalization, deterministic evaluation, 5 evaluation episodes, 15,000,000 total timesteps, evaluation every 80,000 timesteps, and the same 80/20 chronic split.

Implementation detail. Table 3.7 gives the parameter comparison used for this ablation. The table is based on the `s0` TOML file for each family; the corresponding `s1` and `s2` files keep the same configuration except for the training seed and run name.

Evidence. Figure 3.5 compares the `best_00` baseline against each heavy-GINE alternative. Each panel shows the seed-averaged test-survival curve for `best_00` and one comparison family, with faint seed curves and variability bands.

Interpretation. The broad search suggests that the heavy baseline is not limited by capacity alone. Several simpler or more constrained variants improve final survival over `best_00`, especially removing node-ID embeddings, removing flat concatenation, using an MLP critic, and combining a light GINE actor with an MLP critic. In contrast, simply switching message passing to GAT or weighted GCN, or removing the initial do-nothing bias while keeping the heavy GINE/GNN-critic structure, does not clearly solve the stability problem.

Table 3.7: Heavy-GINE best-search 2 parameter comparison from configs/gine_s0_s1_s2.

Run	Main change from best_00	GNN	GNN size	Critic	Actor/critic heads	Concat	Shared	Edge pre.	Node ID	Entropy	Init action-0	Opt. critic
best_000	Removes flat-observation concatenation	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	no	yes	yes	yes	0.02–0.02	0.7	no
best_00	Reference: heavy concat-flat GINE with GNN critic	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
best_01	Enables optimized critic updates	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	yes
best_02	Replaces the GNN critic with an MLP critic	gine	2×128	mlp	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
best_03	Uses non-shared actor GNN weights	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	no	yes	yes	0.02–0.02	0.7	no
best_04	Uses a lighter GINE encoder	gine	1×64	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
best_05	Decays entropy from 0.02 to 0.0	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.00	0.7	no
best_06	Removes learned node-ID embeddings	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	no	0.02–0.02	0.7	no
best_07	Replaces GINE with GAT message passing	gat	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
best_08	Replaces GINE with weighted GCN message passing	gcn	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.7	no
best_09	Light no-concat GINE with MLP critic and optimized critic updates	gine	1×64	mlp	$2 \times 128 / 2 \times 128$	no	yes	no	yes	0.02–0.02	0.7	yes
best_10	Light concat-flat GINE with MLP critic and optimized critic updates	gine	1×64	mlp	$2 \times 128 / 2 \times 128$	yes	yes	no	yes	0.02–0.02	0.7	yes
best_11	Removes the initial do-nothing bias	gine	2×128	gnn	$3 \times 256 / 3 \times 256$	yes	yes	yes	yes	0.02–0.02	0.0	no




Figure 3.5: Heavy-GINE best-search 2 survival comparisons. Each subplot compares the **best_00** heavy concat-flat GINE baseline against one controlled alternative from `configs/gine_s0_s1_s2`.

3.5 WCCI Action-Space Reduction

The WCCI bus36 setting has a much larger discrete topology-action space than bus14. Before training MAPPO on this environment, the action space is reduced offline from simulated action outcomes. During rollouts, states are collected when the grid is stressed,

$$\rho_{\max} > 0.90.$$

For each collected state s , candidate actions are simulated and compared against do nothing:

$$\Delta(s, a) = \rho_{\text{after}}(s, a) - \rho_{\text{after}}(s, 0).$$

An action is counted as useful when it lowers the post-action loading relative to do nothing, i.e. $\Delta(s, a) < -\epsilon$. Each action is then scored by its empirical improvement rate,

$$\text{score}(a) = \frac{\#\{\text{tested states where } a \text{ improves the grid}\}}{\#\{\text{tested states where } a \text{ is evaluated}\}},$$

with a minimum-count filter to avoid selecting rarely sampled actions. The final reduced action space keeps the best actions per agent, while always preserving action 0. The candidate sets are then sanity-checked with a one-step simulator-greedy controller: at each risky state, the controller simulates the available reduced actions and executes the best improving one.

Table 3.8: One-step greedy evaluation of WCCI reduced action spaces on 50 random test chronics. All rows use the same chronic sample (`chronic_sample_seed=0`); the do-nothing mean survival on this sample is 28.30%.

Reduced action space	Kept actions per agent (a_0, a_1, a_2, a_3)	Greedy survival (%)	Do-nothing survival (%)	Survival delta (%)	Greedy full survival (%)
Do-nothing reference	1, 1, 1, 1	–	28.30	–	–
WCCI $h = 1$, mk64	64, 64, 64, 64	55.24	28.30	26.94	22.00
WCCI $h = 1$, mk128	77, 128, 127, 128	60.65	28.30	32.35	22.00
WCCI $h = 1$, mk256	77, 256, 127, 256	68.61	28.30	40.31	36.00
WCCI $h = 1$, mk512	77, 512, 127, 512	71.21	28.30	42.90	40.00
WCCI $h = 1$, mk1024	77, 1024, 127, 1024	78.18	28.30	49.87	54.00

3.6 Decentralized Sparse Intervention Control

Question. Can the agents keep high survival while behaving more like operators: acting rarely, acting locally, acting mainly when the grid state makes an intervention useful, and keeping the final policy decentralized?

Baseline and change. This block starts from the flat decentralized MAPPO topology-control actor. For agent i , the baseline policy is one categorical distribution over the full local action space:

$$\pi_i(a_i \mid o_i), \quad a_i \in \{0, 1, \dots, |\mathcal{A}_i| - 1\}.$$

Action 0 is the no-op action. During training, actions are sampled from the policy, while deterministic evaluation uses the greedy action by default:

$$a_{i,t}^{\text{eval}} = \arg \max_a \pi_i(a \mid o_{i,t}).$$

The sparse-control roadmap tests four ways to reduce unnecessary topology changes: evaluation-time rho overrides, an intervention-gated actor, fixed non-idle action penalties, and adaptive Lagrangian intervention budgets.

Motivation. The intervention gate is motivated by hierarchical topology-control work, where the action can be decomposed into “do nothing versus intervene”, then where to act, then which local topology configuration to apply [3], [4]. In this thesis only the first level is adapted: each regional agent independently decides whether to do nothing or to execute one of its existing local topology actions. The fixed and adaptive sparse-control objectives are closer to constrained and budgeted RL: non-idle topology interventions are treated as a limited resource, following the general Lagrangian view of constrained policy optimization and budgeted control [5], [6], [7]. This is also operationally meaningful because switching frequency and topology deviation are real power-grid objectives, not just cosmetic metrics [8].

3.6.1 Evaluation-Time Rho Heuristics

Implementation detail. The rho heuristic is an evaluation-only diagnostic. It does not change the training objective and it does not prove that the policy learned sparse behavior. The policy first proposes the deterministic action $\bar{a}_{i,t}$, then the evaluator may replace it with action 0.

The global version uses the pre-action maximum line loading

$$\rho_t^{global} = \max_{\ell \in \mathcal{L}} \rho_{\ell,t},$$

and executes

$$a_{i,t}^{exec} = \begin{cases} 0, & \rho_t^{global} < \rho_{safe}, \\ \bar{a}_{i,t}, & \text{otherwise,} \end{cases} \quad \rho_{safe} = 0.90.$$

The local version is decentralized. Agent i computes the maximum rho on the lines touching its regional substations,

$$\rho_{i,t}^{local} = \max_{\ell \in \mathcal{L}_i} \rho_{\ell,t},$$

and only that agent is forced to no-op when $\rho_{i,t}^{local} < \rho_{safe}$. A boundary line can appear in the local neighborhood of both adjacent agents, but no communication is required: both agents can observe the line through their own physical neighborhood.

Interpretation. The heuristic is useful to ask what happens if safe-state interventions are blocked, but it is not a learned method. Its action-0 rates must therefore be reported as final executed evaluation actions, not as evidence that the policy itself prefers action 0.

3.6.2 Intervention-Gated Actor

Implementation detail. The gated actor changes the actor factorization. Instead of one categorical head over all actions, each agent has a gate head and a non-idle action head:

$$\pi_i^{gate}(g_i | o_i), \quad g_i \in \{0, 1\},$$

where 0 means do nothing and 1 means intervene, and

$$\pi_i^{nonidle}(b_i | o_i), \quad b_i \in \{1, \dots, |\mathcal{A}_i| - 1\}.$$

During training, the gate is sampled first. If $g_{i,t} = 0$, then $a_{i,t} = 0$. If $g_{i,t} = 1$, the final action is sampled from the non-idle head. The PPO log-probability of the executed action is therefore

$$\log P(a_i = 0) = \log P(g_i = 0),$$

and, for $a_i > 0$,

$$\log P(a_i) = \log P(g_i = 1) + \log P(b_i = a_i | g_i = 1).$$

Two deterministic decoders are used. The `final_action_map` decoder selects the most likely final executed action:

$$P(a_i = 0) = P(g_i = 0), \quad P(a_i = a > 0) = P(g_i = 1)P(b_i = a | g_i = 1).$$

The `hierarchical_greedy` decoder first chooses the most likely gate decision, then chooses the most likely non-idle action only if the gate selects intervention.

The original coupled entropy term is

$$H_i = H(g_i) + P(g_i = 1)H(b_i).$$

This can accidentally reward intervention because larger $P(g_i = 1)$ unlocks more non-idle entropy. The separated entropy variant is

$$H_i = \alpha_{gate}H(g_i) + \alpha_{nonidle}H(b_i),$$

and is used in the gate-plus-budget diagnostic runs.

Interpretation. The gate is learned and decentralized, but it can restrict exploration more strongly than a reward penalty. If it collapses early to no-op, the non-idle head is rarely executed and receives little learning signal. For this reason, the gate is treated as an important diagnostic rather than the main sparse control contribution.

3.6.3 Fixed Sparse-Action Penalty

Implementation detail. The Sparse16 sweep adds a fixed reward penalty to the final executed action of each agent:

$$r'_{i,t} = r_{i,t} - \lambda_{fixed} \mathbf{1}[a_{i,t} \neq 0],$$

with

$$\lambda_{fixed} \in \{0.0, 0.001, 0.003, 0.01\}.$$

The grid is

$$\text{actor} \in \{\text{flat}, \text{gated}\}, \quad \lambda_{fixed} \in \{0.0, 0.001, 0.003, 0.01\}, \quad \text{seed} \in \{0, 1, 2\},$$

which gives $2 \times 4 \times 3 = 24$ runs. All Sparse16 runs set `safe_intervention_penalty = 0.0`, so the penalty is not rho-weighted.

Interpretation. Sparse16 changes only the training reward. It does not override actions during evaluation and it does not hard-mask exploration during training. The policy can still sample non-idle actions, but each such action must be useful enough to pay its fixed cost. The key comparison is `sparse16_flat_p010` versus `sparse16_gated_p010`, which asks whether the gate adds value once a simple sparse objective already exists.

3.6.4 Adaptive Intervention Budget

Implementation detail. The adaptive intervention budget turns sparse topology control into a local constraint. For each agent,

$$c_i(s_t, a_{i,t}) = \mathbf{1}[a_{i,t} \neq 0] w_i(s_t),$$

and the desired constraint is

$$\mathbb{E}_{\pi_i} [c_i(s_t, a_{i,t})] \leq d.$$

The implemented rollout reward removes the constant term that does not affect the PPO action gradient:

$$r'_{i,t} = r_{i,t} - \lambda_i c_i(s_t, a_{i,t}).$$

After each rollout, the local multiplier is updated by

$$\lambda_i \leftarrow \text{clip} \left(\lambda_i + \eta (\hat{C}_i - d), 0, \lambda_{\max} \right), \quad \hat{C}_i = \frac{1}{T} \sum_{t=1}^T c_i(s_t, a_{i,t}).$$

If the observed cost $\hat{C}_i$ is above the budget d , future interventions become more expensive; if it is below the budget, the penalty relaxes. The default settings are `intervention_budget_lr = 0.02`, `intervention_budget_init_lambda = 0.0`, and `intervention_budget_max_lambda = 10.0`.

Three cost modes are used:

- **Non-idle cost:** $w_i(s_t) = 1$, so $c_i(s_t, a_{i,t}) = \mathbf{1}[a_{i,t} \neq 0]$. This is an adaptive version of a plain sparse-action penalty.
- **Global-safe cost:** $w_i^{global}(s_t) = \sigma(k(\rho_{safe} - \rho_t^{global}))$. This penalizes non-idle actions mostly when the whole grid is safe.
- **Local-safe cost:** $w_i^{local}(s_t) = \sigma(k(\rho_{safe} - \rho_{i,t}^{local}))$. This is the decentralized version: safe local states make interventions expensive, while stressed local states make them cheap.

In the AIB setting, $\rho_{safe} = 0.90$ is not a hard action override. It is the center of a smooth sigmoid safety weight.

AIB experiment grid. The first AIB folder, `configs/adaptive_intervention_budget_7`, contains the main local-budget ablations:

- `aib_00_flat_local_t020`: flat actor, local-safe cost, $d = 0.20$;
- `aib_01_flat_local_t010`: stricter local-safe budget, $d = 0.10$;
- `aib_02_flat_local_t035`: looser local-safe budget, $d = 0.35$;
- `aib_03_gate_hgreedy_sep_local_t020`: gated actor, hierarchical-greedy evaluation, separated entropy, local-safe cost, $d = 0.20$;
- `aib_04_flat_nonidle_t020`: flat actor with adaptive plain non-idle cost, $d = 0.20$.

The comparison `aib_00_flat_local_t020` versus `aib_04_flat_nonidle_t020` isolates the value of local rho weighting.

The mechanism-ablation folder `configs/adaptive_intervention_budget_mechanism_15` contains 15 additional flat-actor runs. It separates fixed versus adaptive coefficients, plain non-idle versus state-dependent costs, global versus local rho weighting, and rho-threshold sensitivity.

Table 3.9: Sparse-control mechanisms and where they affect the policy.

Method	Training effect	Evaluation effect	Main risk
Baseline	Samples from the flat policy	Greedy argmax by default	Standard entropy-controlled exploration
Rho heuristic	No training change	Hard no-op override in safe states	Evaluation is manually changed
Intervention gate	Changes action factorization and sampling	Final-action MAP or hierarchical greedy	Gate collapse can starve non-idle exploration
Sparse16	Fixed reward penalty on non-idle actions	No action override	Penalty coefficient must be tuned
AIB non-idle	Adaptive reward penalty on non-idle actions	No action override	Multiplier can grow if target is too strict
AIB local-safe	Adaptive state-dependent reward penalty	No action override	Depends on the local rho safety-weight design

Table 3.10: Main sparse-control comparisons to report.

Question	Comparison
Does a fixed sparse penalty already work?	Baseline versus <code>sparse16_flat_p010</code>
Does the gate help beyond a sparse objective?	<code>sparse16_flat_p010</code> versus <code>sparse16_gated_p010</code>
Fixed versus adaptive coefficient?	<code>sparse16_flat_p010</code> or <code>aibm_00_fixed_nonidle_p006</code> versus <code>aib_04_flat_nonidle_t020</code>
Does local-safe weighting matter?	<code>aib_00_flat_local_t020</code> versus <code>aib_04_flat_nonidle_t020</code>
Does global-safe weighting matter?	<code>aibm_02_adaptive_global_t020</code> versus <code>aib_00_flat_local_t020</code>
Is $\rho_{safe} = 0.90$ special?	<code>aibm_03</code> with 0.85, <code>aib_00</code> with 0.90, and <code>aibm_04</code> with 0.95
Is the heuristic only an evaluation trick?	Rho-heuristic runs versus learned sparse/AIB runs

Full-test checkpoint summaries. The following tables report the average full-test episodic survival over the three evaluated seeds. The survival values are computed directly from the JSON files in `Topology_Task/outputs/full_test_eval`. These JSON files also record where the compact action-distribution artifacts were written. The action-0 fractions are computed from the corresponding local `Topology_Task/outputs/full_test_eval_actions/*/action_summary.json` files and averaged over

seeds. The all-agent action-0 column is the arithmetic mean of the three per-agent action-0 fractions. The do-nothing row is a reference policy that always selects action 0; its survival is computed from `Topology_Task/outputs/do_nothing_eval/bus14_test_20`. Remaining “–” entries indicate that the compact action summaries for those full-test jobs are not available locally.

Table 3.11: Full-test summary for evaluation-time rho heuristics.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Global rho heuristic, $\rho_{safe} = 0.90$	98.36	0.936	0.933	0.936	0.940
Local rho heuristic, $\rho_{safe} = 0.90$	97.07	0.968	0.992	0.985	0.928

Table 3.12: Full-test summary for intervention-gated actors.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Intervention gate, final-action MAP	89.82	0.802	0.826	0.889	0.690
Intervention gate, hierarchical greedy	93.66	0.395	0.066	0.844	0.276

Table 3.13: Full-test summary for Sparse16 flat-policy runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Flat Sparse16, $\lambda_{fixed} = 0.000$	98.23	0.494	0.524	0.463	0.496
Flat Sparse16, $\lambda_{fixed} = 0.001$	96.75	0.486	0.642	0.401	0.416
Flat Sparse16, $\lambda_{fixed} = 0.003$	93.97	0.575	0.691	0.346	0.688
Flat Sparse16, $\lambda_{fixed} = 0.010$	97.44	0.905	0.907	0.954	0.853

Metrics to report. The main performance metric remains episodic survival. It should be paired with sparsity and physical-diagnostic metrics: per-agent action-0 fractions, per-agent non-idle rates, the distribution of how many agents act simultaneously, intervention-budget multipliers, budget costs and violations, the rate of interventions in costly safe states, pre- and post-action max rho, and topology-distance changes.

Interpretation. The current roadmap interpretation is that the gate is not the strongest contribution by itself. Sparse reward shaping already improves action-0 usage, the adaptive budget makes the sparse penalty self-tuning, and the local-safe cost makes the penalty more physically meaningful. In particular, `sparse16_flat_p010` is a strong fixed-penalty baseline, while the most principled learned sparse-control candidate is `aib_00_flat_local_t020`: it keeps the original decentralized actor, does not rely on an evaluation-time override, learns the intervention penalty adaptively, and uses a local state-dependent safety weight.

3.7 Current Takeaway

The current reruns indicate that the strongest direction is not a single isolated modification. Stable performance comes from a combination of A0-style training choices, reward normalization, careful initialization, and architectures that keep the critic and graph encoder simple enough to optimize reliably. Optimized critic updates and lighter GINE/MLP-critic variants can improve the final GINE rerun performance, but the A0 baseline remains the clearest stable reference point for future ablations. The sparse-control roadmap adds a second axis to the thesis: after survival is stabilized, the next question is whether the same decentralized agents can achieve high survival with fewer unnecessary topology interventions.

Table 3.14: Full-test summary for Sparse16 gated-policy runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Gated Sparse16, $\lambda_{fixed} = 0.000$	78.51	0.787	0.747	0.949	0.665
Gated Sparse16, $\lambda_{fixed} = 0.001$	73.82	0.862	0.956	0.832	0.799
Gated Sparse16, $\lambda_{fixed} = 0.003$	94.59	0.887	0.934	0.929	0.799
Gated Sparse16, $\lambda_{fixed} = 0.010$	85.69	0.974	0.990	0.984	0.947

Table 3.15: Full-test summary for adaptive intervention budget runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Baseline flat policy	93.67	0.434	0.554	0.385	0.363
Flat local-safe AIB, $d = 0.20$	98.31	0.978	0.980	0.989	0.966
Flat local-safe AIB, $d = 0.10$	98.39	0.979	0.992	0.995	0.949
Flat local-safe AIB, $d = 0.35$	96.61	0.934	0.959	0.947	0.896
Gated local-safe AIB, $d = 0.20$, hierarchical greedy	33.87	0.999	1.000	1.000	0.998
Flat non-idle AIB, $d = 0.20$	95.08	0.986	0.997	0.992	0.970

Table 3.16: Full-test summary for the teacher-student distillation runs.

Run type	Avg. episodic survival (%)	Avg. action-0 all agents	Avg. action-0 agent 0	Avg. action-0 agent 1	Avg. action-0 agent 2
Do-nothing policy	20.25	1.000	1.000	1.000	1.000
Teacher: MAPPO with local rho heuristic	97.07	0.968	0.992	0.985	0.928
Student BC, balanced non-idle 0.50, weight 5, aux 0.50	92.14	0.939	0.974	0.925	0.919
Student BC, balanced non-idle 0.20, weight 3, aux 0.25	90.76	0.957	0.982	0.948	0.940

Appendix A

GINE Architecture

This appendix will describe the GINE encoder used in the graph-based topology control experiments. Building on the graph definition in Section [2.1](#), the final version should specify the message-passing layers, pooling/readout operation, and how the actor and critic heads consume the encoded graph representation.

Appendix B

Full-Test Checkpoint Registry

Table [B.1](#) reports the checkpoint step used for each full-test result table. The values are the `checkpoint_global_step` fields stored in the corresponding JSON files under `Topology_Task/outputs/full_test_eval`. The do-nothing policy has no learned checkpoint. For the behavioral cloning students, the stored step is the student checkpoint’s optimizer-step count rather than a MAPPO environment step.

Table B.1: Checkpoint steps used for the full-test evaluations.

Run type	Seed 0 checkpoint step	Seed 1 checkpoint step	Seed 2 checkpoint step
Do-nothing policy	—	—	—
Baseline flat policy	13,600,000	13,040,000	13,760,000
Global rho heuristic, $\rho_{safe} = 0.90$	13,920,000	14,320,000	14,800,000
Local rho heuristic, $\rho_{safe} = 0.90$	14,160,000	14,400,000	14,160,000
Intervention gate, final-action MAP	12,000,000	12,720,000	13,760,000
Intervention gate, hierarchical greedy	7,120,000	14,720,000	13,760,000
Flat Sparse16, $\lambda_{fixed} = 0.000$	9,600,000	12,160,000	7,920,000
Flat Sparse16, $\lambda_{fixed} = 0.001$	10,640,000	11,520,000	11,600,000
Flat Sparse16, $\lambda_{fixed} = 0.003$	12,000,000	12,960,000	12,480,000
Flat Sparse16, $\lambda_{fixed} = 0.010$	8,960,000	5,760,000	12,560,000
Gated Sparse16, $\lambda_{fixed} = 0.000$	8,480,000	4,320,000	6,000,000
Gated Sparse16, $\lambda_{fixed} = 0.001$	5,760,000	7,600,000	3,280,000
Gated Sparse16, $\lambda_{fixed} = 0.003$	8,960,000	7,120,000	7,600,000
Gated Sparse16, $\lambda_{fixed} = 0.010$	12,000,000	7,840,000	7,920,000
Flat local-safe AIB, $d = 0.20$	13,120,000	9,360,000	12,160,000
Flat local-safe AIB, $d = 0.10$	12,880,000	13,200,000	12,720,000
Flat local-safe AIB, $d = 0.35$	12,800,000	5,360,000	11,520,000
Gated local-safe AIB, $d = 0.20$, hierarchical greedy	14,960,000	5,280,000	16,800,000
Flat non-idle AIB, $d = 0.20$	12,480,000	14,880,000	11,920,000
Teacher: MAPPO with local rho heuristic	14,160,000	14,400,000	14,160,000
Student BC, balanced non-idle 0.50, weight 5, aux 0.50	12,375	14,994	13,617
Student BC, balanced non-idle 0.20, weight 3, aux 0.25	12,375	14,994	13,617

Bibliography

- [1] B. Donnot, “Grid2op: A testbed platform to model sequential decision making in power systems,” *arXiv preprint arXiv:2011.07929*, 2020.
- [2] C. Yu et al., “The surprising effectiveness of ppo in cooperative multi-agent games,” *Advances in Neural Information Processing Systems*, 2022.
- [3] J. Manczak et al., *Hierarchical reinforcement learning for power network topology control*, 2023. arXiv: [2311.02129](https://arxiv.org/abs/2311.02129). [Online]. Available: <https://arxiv.org/abs/2311.02129>
- [4] E. van der Sar et al., *Hierarchical multi-agent reinforcement learning for power grid topology optimization*, 2023. arXiv: [2310.02605](https://arxiv.org/abs/2310.02605). [Online]. Available: <https://arxiv.org/abs/2310.02605>
- [5] J. Achiam, D. Held, A. Tamar, and P. Abbeel, “Constrained policy optimization,” *International Conference on Machine Learning*, 2017. [Online]. Available: <https://arxiv.org/abs/1705.10528>
- [6] A. Stooke, J. Achiam, and P. Abbeel, “Responsive safety in reinforcement learning by PID lagrangian methods,” *International Conference on Machine Learning*, 2020. [Online]. Available: <https://arxiv.org/abs/2007.03964>
- [7] N. Carrara, E. Leurent, R. Laroche, T. Urvoy, and O.-A. Maillard, “Budgeted reinforcement learning in continuous state space,” *Advances in Neural Information Processing Systems*, 2019. [Online]. Available: <https://arxiv.org/abs/1903.01004>
- [8] L. Lautenbacher et al., *Multi-objective reinforcement learning for power grid topology control*, 2025. arXiv: [2502.00040](https://arxiv.org/abs/2502.00040). [Online]. Available: <https://arxiv.org/abs/2502.00040>
- [9] J. Schulman, F. Wolski, P. Dhariwal, A. Radford, and O. Klimov, “Proximal policy optimization algorithms,” *arXiv preprint arXiv:1707.06347*, 2017.
- [10] A. Marot et al., “Learning to run a power network challenge for training topology controllers,” *Electric Power Systems Research*, 2021.
- [11] T. N. Kipf and M. Welling, “Semi-supervised classification with graph convolutional networks,” *International Conference on Learning Representations*, 2017.